

Course Description: 6.852, Distributed Algorithms

1 People and places

Instructor: Prof. Nancy Lynch, 32-G668, MIT extension 3-7225, lynch@csail.mit.edu. Available by appointment; Monday-Thursday afternoons are generally good.

Teaching assistants: Cameron Musco, 32-G670, cnmusco@mit.edu
Office hours: Tuesdays 2:00-3:30pm or by appointment. In 32-G670.

Katerina Sotiraki, 32-G670, katesot@mit.edu
Office hours: Wednesdays 4:30-6:00pm or by appointment. In 32-G670.

Course secretary: Joanne Talbot Hanley, 32-G672A, MIT extension 3-6054.

Class meetings: Tuesdays and Thursdays, 11:00AM - 12:30PM in Room 32-144.

Web site and mailing list: Through Stellar, registered students will automatically be set up on a mailing list. However, if you are not registered and would like to be on the list, please email one of the TAs.

Website: <https://learning-modules.mit.edu/class/index.html?uuid=/course/6/fa15/6.852#info>

In addition to all the usual course-related material, this site will contain downloadable copies of course handouts and related research papers.

2 What is this course about?

Distributed Algorithms are algorithms that are designed to run on many processors, without centralized control. In general, they are much harder to design, and much harder to understand, than single-processor sequential algorithms. Distributed algorithms are important for most practical systems, including large-scale and local-area computer networks, data management systems, and multiprocessor shared-memory systems. Distributed algorithms also have a rich theory, which forms the subject matter for this course. A very recent Turing Award winner, Leslie Lamport, is a pioneer of this theory, as was one of the earliest Turing Award winners, Edsger Dijkstra.

Theoretical results about distributed algorithms appear in research conferences such as PODC (Principles Of Distributed Computing), DISC (International Symposium on DIStributed Computing), OPODIS (International Conference on Principles of Distributed Systems), and SPAA (ACM Symposium on Parallelism in Algorithms and Architectures). They also appear in general theoretical computer science conferences such as FOCS (Foundations of Computer Science) and STOC (Symposium on Theory of Computing), and in broad conferences involving distributed computing, such as ICDCS (International Conference on Distributed Computing Systems).

What do distributed algorithms researchers do? They (we) define various kinds of distributed system models, and identify important problems to be solved using those models. They design new algorithms to solve the problems, and analyze their behavior—including correctness, performance, and fault-tolerance. They also

prove lower bounds and other impossibility results, which explain why certain tasks cannot be carried out in certain kinds of distributed settings, or cannot be carried out within certain cost constraints. Researchers typically study problems that arise in practical distributed systems, including problems of communication, computation, building various graph structures, data management, resource management, synchronization, and consensus. Some of the results have impact on practical distributed system design.

This year, the course will focus mainly on basic distributed algorithms and impossibility results, as covered in Lynch's book *Distributed Algorithms*. We will supplement this with some related material on distributed graph network algorithms, wait-free computability, and failure detectors. At the end, we will introduce some current topics, such as self-stabilization, dynamic networks, and/or social insect colony algorithms.

6.852 is intended to do two things: provide an introduction to the most important basic results in the area of distributed algorithms, and prepare interested students to begin independent research or take a more advanced course in this area.

3 Prerequisites

To take 6.852, you should have:

- “Mathematical maturity”. In particular, you should be really good at reading and writing mathematical proofs.
- General knowledge about some distributed systems. For instance, MIT's undergraduate course 6.033, Computer Systems Engineering, would be good background.
- Experience with sequential algorithms and their analysis. MIT's undergraduate course 6.046 would be good background.
- (Desirable, but not essential) Experience with formal models of computation. MIT's course 6.045 or 6.840 would be good.

4 Source material

The primary source will be the book *Distributed Algorithms* by Nancy Lynch, which can be purchased at the MIT Coop. This book has gone through many printings, but we have made no changes since the fourth printing, so fourth and later are just fine. Known errata are collected in an errata list, which is accessible from the course Web page.

The book refers to many papers from the research literature on distributed algorithms; you might want to track down some of these and read them.

Other books that you should find useful are:

- Hagit Attiya and Jennifer Welch, *Distributed Computing: Fundamentals, Simulations, and Advanced Topics*. John Wiley and Sons, Inc., 2004. Second Edition.
This is another textbook on distributed algorithms. It was originally published soon after the Lynch book, but this one now has a second edition. The material covered overlaps quite a lot with the Lynch book, though Attiya and Welch do cover some additional topics, like clock synchronization.
- Shlomi Dolev. *Self-Stabilization*, MIT Press, 2000. This gives a good description of self-stabilizing distributed algorithms. Self-stabilization is a strong kind of fault-tolerance.
- Dilsun Kaynar, Nancy Lynch, Roberto Segala, and Frits Vaandrager, *The Theory of Timed I/O Automata, Second Edition*, Morgan Claypool, 2010. This monograph presents a basic modeling framework, Timed I/O Automata (TIOA), for describing and analyzing distributed algorithms, especially those that depend on timing.

- David Peleg, *Distributed Computing: A Locality-Sensitive Approach*, SIAM, 2000. This presents distributed graph network algorithms, focusing on topics that illustrate the role of locality in designing such algorithms.

These books are all available through MIT's library. All but *Self-Stabilization* are available as PDFs online if you search their titles at: <http://library.mit.edu>.

In addition, some research papers whose contents are not covered in the textbook will be covered in class and on problem sets. These papers will be listed on the course web site.

5 Course requirements

5.1 Readings

We will announce readings that cover the topics for each class beforehand. Most of the readings will be sections from the main textbook. Some will be sections of the other books and research papers listed on the web site. We expect you to read the assigned material ahead of time, and be prepared to discuss it in class.

5.2 Problem sets

These are intended to help you to understand the material covered in class. Most problems will be about results already covered; some will be designed to get you started thinking about ideas that will arise later. Specifically, we will assign approximately 4-5 problems every Thursday. The problems will be batched and due every two weeks, at class on alternate Thursdays. (See the course schedule, Handout 2, for exact dates.) There will be a total of six (two-part) problem sets. **No late homeworks will be accepted.** If you haven't finished, just hand in what you have completed. Of course, in case of a serious emergency, please talk to one of the course staff.

When grading homework problems, we will try to give full credit to solutions that include all the important logical steps and ideas. Your writeups should not be overly long and detailed. (However, we will sometimes ask specifically for details, for instance, when we are studying formal correctness proofs for algorithms.)

We will hand out solutions to homework problems, usually the best student solutions. If you would like us to use your writeups, you can help us by writing elegant and concise solutions and formatting them using L^AT_EX (and the L^AT_EX style files we will provide). When you submit your homework, keep the .tex file since you (or we) may need to edit it. Problem sets will be graded by teams of students in the class, led by the course staff.

Policy on homework collaboration: You are encouraged to discuss possible solutions with other class members. Many students in past incarnations of this course have formed homework discussion groups. However, *you must always write up the solutions entirely on your own.*

5.3 Problem set grading

For each problem set, a small group of students will be responsible for working with the course staff to grade the solutions. We would like the grading to be completed by the Monday afternoon after the homeworks are handed in, so we can record the grades and hand the graded homeworks back on Tuesday. The number of times you will have to grade over the course of the semester should be one or two, depending on the size of the class. Part of your grade will depend on the quality and promptness of your work on problem set grading.

5.4 Term projects

This year, students are required to carry out a term project related to distributed algorithms, interpreted broadly. Distributed algorithms that arise in computer systems or applications would be great to consider. We will provide a handout with more details about this in a couple of weeks, but here is a summary.

Students can work alone, or in teams of two or three. Three kinds of projects are acceptable:

- *Reading project:* Choose a topic of current interest in the distributed algorithms research community (we will make some suggestions). Read the most important papers on the topic and write an expository report explaining the key theoretical ideas.
- *Theoretical research project:* Find a research problem you are interested in, and devise your own new algorithm(s) and/or prove your own new lower bound result(s). Write a theoretical research report about it. For this, it's better to consider current research topics, so you will need to do some background reading for this type of project as well.
- *Experimental research project:* You can simulate known theoretical distributed algorithms and perform experiments to determine how they behave under various assumptions. Write an experimental research report about it. This kind of project will need a clear question and clear conclusions.

A project proposal will be due in class on Thursday, October 29, which will give us a chance to approve the project and make suggestions. The final report will be due at the last regular class, on Thursday, December 10. We will hold an extra meeting, on Friday, December 11, for students to present summaries of their projects. Attendance at this extra meeting is optional, and if you cannot attend it will not affect your grade. The main purpose of the presentations is to give you a chance to tell everyone what you worked on. We will provide lunch.

5.5 Exams

There will be no exams. However, we are not quite done after the last class—we will have one optional meeting on Friday, December 11, for project presentations. See above.

5.6 Course grade calculation

Your course grade will be based on problem set grades, term project grades, and grades for your class participation and problem set grading. Here is how we will calculate it:

- Problem sets: 60%
- Term projects: 25%
- Class participation and grading: 15%