

6.852: Distributed Algorithms


Fall, 2015

Lecture 15

Today's plan


- Asynchronous shared-memory system model
- The Mutual Exclusion problem
- Dijkstra's algorithm
- Peterson's algorithms
- Lamport's Bakery algorithm
- **Reading:** Chapter 9, Sections 10.1-10.5, 10.7
- Next:
 - More Mutual Exclusion algorithms
 - Lower bound on the number of shared variables
 - Resource allocation
- Reading: Sections 10.6,10.8, Chapter 11 (skim)

Asynchronous Shared-Memory Systems



Asynchronous Shared-Memory Systems

- We've covered basics of non-fault-tolerant asynchronous network algorithms:
 - How to model them.
 - Basic asynchronous network protocols---broadcast, spanning trees, leader election,...
 - General methods for designing asynchronous network algorithms:
 - Synchronizers
 - Logical time
 - Global snapshots
- Now consider asynchronous shared-memory systems:
- Processes, interacting via shared objects, possibly subject to some access constraints.
- Shared objects have types, e.g.:
 - Read/write (weak)
 - Read-modify-write, compare-and-swap (strong)
 - Queues, stacks, others (in between)



Asynch Shared-Memory systems




- Theory of ASM systems has much in common with theory of asynchronous networks:
 - Similar algorithms and impossibility results.
 - Even with failures.
 - Transformations from ASM model to asynch network model allow ASM algorithms to run in asynchronous networks.
 - “Distributed shared memory”.
- Historically, theory for ASM began first.
- Arose long ago, in study of early operating systems, in which several processes run on a single processor, sharing memory, with possibly-arbitrary interleavings of steps.
- Currently, ASM models are used to describe multiprocessor shared-memory systems, in which several processes run on separate processors and share memory.

Topics

- Define the basic system model, without failures.
- Basic problems:
 - Mutual Exclusion.
 - Other resource-allocation problems (briefly).
- Then, introduce process failures into the model.
- More basic problems:
 - Distributed consensus
 - Implementing atomic objects:
 - Atomic snapshot objects
 - Atomic read/write registers
- Wait-free and fault-tolerant computability theory


Basic ASM Model, Version 1

- Processes + objects, modeled as automata.
- Arrows:
 - Represent invocations and responses for operations on the objects.
 - Modeled as input and output actions.
- Fine-granularity model, can describe:
 - Delay between invocation and response.
 - Concurrent (overlapping) operations:
 - Object could reorder operations.
 - Could allow them to run concurrently, interfering with each other.
- We'll begin with a simpler, coarser model:
 - Object runs ops in invocation order, one at a time.
 - In fact, collapse each operation into a single step.
- Return to the finer model next week.



Basic ASM Model, Version 2


- One big shared memory system automaton A.
- External actions at process “ports”.
- Each process i has:
 - A set **states_i** of states.
 - A subset **start_i** of start states.
- Each variable x has:
 - A set **values_x** of values it can take on.
 - A subset **initial_x** of initial values.
- Automaton A:
 - **States**: State for each process, a value for each variable.
 - **Start**: Start states, initial values.
 - **Actions**: Each action associated with one process, and some also with a single shared variable.
 - **Input/output actions**: At the external boundary.
 - **Transitions**: Correspond to local process steps and variable accesses.
 - Action enabling, which variable is accessed, depend only on process state.
 - Changes to variable and process state depend also on variable value.
 - Must respect the type of the variable.
 - **Tasks**: One or more per process (threads).



Basic ASM Model

- **Execution of A:**

- As specified by general definitions of executions, fair executions for I/O automata.
- By fairness definition, each task gets infinitely many chances to take steps.
- Model environment as a separate automaton, to express restrictions on environment behavior.



- **Commonly-used variable types:**

- Read/write registers: Most basic object type.
 - Allows access using separate read and write operations.
- Read-modify-write: Most powerful object type:
 - Atomically, read variable, do local computation, write to variable.
- Compare-and-swap, fetch-and-add, queues, stacks, etc.

- Different computability and complexity results hold for different variable types.

The Mutual Exclusion Problem


The Mutual Exclusion Problem

- Share one resource among n user processes, U_1, U_2, \dots, U_n .
 - E.g., printer, portion of a database.
- U_i has four “regions”.
 - Subsets of its states, described by portions of its code.
 - C critical; R remainder; T trying; E exit

Protocols for obtaining and relinquishing the resource



- Cycle: $R \longrightarrow T \longrightarrow C \longrightarrow E$

- Architecture:
 - U_i s and A are IOAs, compose.




The Mutual Exclusion Problem

- **Actions at user interface:**
 - Connect U_i to P_i
 - p_i is U_i 's "agent"
- **Correctness conditions:**
 - **Well-formedness (Safety):**
 - System also obeys cyclic discipline.
 - E.g., doesn't grant resource when it wasn't requested.
 - **Mutual exclusion (Safety):**
 - System never grants to > 1 user simultaneously.
 - Trace safety property.
 - Or, there's no reachable system state in which >1 user is in C at once.
 - **Progress (Liveness):**
 - From any point in a fair execution:
 - If some user is in T and no user is in C then at some later point, some user enters C.
 - If some user is in E then at some later point, some user enters R.



The Mutual Exclusion Problem


- Well-formedness (Safety):
 - System obeys cyclic discipline.
- Mutual exclusion (Safety):
 - System never grants to > 1 user.
- Progress (Liveness):
 - From any point in a fair execution:
 - If some user is in T and no user is in C then at some later point, some user enters C.
 - If some user is in E then at some later point, some user enters R.



- Conditions constrain only system automaton A , not users.
 - System determines if/when users enter C and R.
 - Users determine if/when users enter T and E.
 - We don't state any requirements on the users, except that users respect well-formedness.

The Mutual Exclusion Problem

- Well-formedness (Safety):
- Mutual exclusion (Safety):
- Progress (Liveness):
 - From any point in a fair execution:
 - If some user is in T and no user is in C then at some later point, some user enters C.
 - If some user is in E then at some later point, some user enters R.



- **Fairness assumption:**
 - Progress condition requires fairness assumption (all process tasks continue to get turns to take steps).
 - Needed to guarantee that some user enters C or R.
 - In general, in the asynchronous model, liveness properties require fairness assumptions.
 - Contrast: Well-formedness and mutual exclusion are safety properties, don't depend on fairness.

One more assumption...

- No permanently active processes.
 - Locally-controlled actions enabled only when user is in T or E.
 - No always-awake, dedicated processes.
 - Motivation:
 - Multiprocessor settings, where users can run processes when needed, but are otherwise not involved in the protocol.
 - Avoid “wasting a processor”.

Dijkstra's Mutual Exclusion Algorithm

[Dijkstra 65]



Mutual Exclusion algorithm

- Based on Dekker's 2-process solution.
- Pseudocode, [p. 265-266](#)
 - Written in traditional sequential style, then translated into more detailed state/transition description.
- Shared variables: Read/write registers.
 - **turn**, in $\{1, 2, \dots, n\}$, multi-writer, multi-reader (mWmR), initially any
 - for each process i :
 - **flag(i)**, in $\{0, 1, 2\}$, single-writer, multi-reader (1WmR), initially 0
 - Written by i , read by everyone.
- Process i 's Stage 1:
 - Set **flag** := 1, test to see if **turn** = i .
 - If not, and **turn**'s current owner is seen to be inactive, then set **turn** := i .
 - Otherwise go back to testing...
 - When you see **turn** = i , move to Stage 2.

Dijkstra's algorithm

- Stage 2:
 - Set `flag(i)` := 2.
 - Check (one at a time, any order) that no other process has `flag` = 2.
 - If check completes successfully, go to C.
 - If not, go back to the beginning of Stage 1.
- Exit protocol:
 - Set `flag(i)` := 0.
- Problem with the sequential code style:
 - Unclear what constitutes an atomic step.
 - E.g., need three separate steps to test `turn`, test `flag(turn)`, and set `turn`.
 - Must rewrite to make this clear:
 - E.g., precondition/effect code (p. 268-269)
 - E.g., sequential-style code with explicit reads and writes, one per line.

Dijkstra's algorithm, pre/eff code

- One transition definition for each kind of atomic step.
- Explicit program counter, *pc*.
- Transitions:
 - *set-flag-1_i*: Sets *flag* to 1 and prepares to test *turn*.
 - *test-turn_i*: Tests *turn*, and either moves to Stage 2 or prepares to test the current owner's *flag*.
 - *test-flag(j)_i*: Tests *j*'s *flag*, and either goes on to set *turn* or goes back to test *turn* again.
 - ...
 - *set-flag-2_i*: Sets *flag* to 2 and initializes set *S*, preparing to check all other processes' *flags*.
 - *check(j)_i*: If *flag(j)* = 2, go back to beginning.
 - ...
- *S* keeps track of which processes have been successfully checked in Stage 2.

Precondition/effect code

Shared variables:

$\text{turn} \in \{1, \dots, n\}$, initially arbitrary

for every i :

$\text{flag}(i) \in \{0, 1, 2\}$, initially 0

Actions of process i :

Input: try_i , exit_i

Output: crit_i , rem_i

Internal: set-flag-1_i , test-turn_i , $\text{test-flag}(j)_i$, set-turn_i , set-flag-2_i ,
 $\text{check}(j)_i$, reset_i

Precondition/effect code, Dijkstra process i

try_i:

Eff: $pc := \text{set-flag-1}$

set-flag-1_i:

Pre: $pc = \text{set-flag-1}$

Eff: $\text{flag}(i) := 1$

$pc := \text{test-turn}$

test-turn_i:

Pre: $pc = \text{test-turn}$

Eff: if $\text{turn} = i$ then $pc := \text{set-flag-2}$
else $pc := \text{test-flag}(\text{turn})$

test-flag(j)_i:

Pre: $pc = \text{test-flag}(j)$

Eff: if $\text{flag}(j) = 0$ then $pc := \text{set-turn}$
else $pc := \text{test-turn}$

set-turn_i:

Pre: $pc = \text{set-turn}$

Eff: $\text{turn} := i$

$pc := \text{set-flag-2}$

set-flag-2_i:

Pre: $pc = \text{set-flag-2}$

Eff: $\text{flag}(i) := 2$

$S := \{i\}$

$pc := \text{check}$

More code, Dijkstra process i

check(j)_i :

Pre: **pc** = check

$j \notin S$

Eff: if **flag(j)** = 2 then

S := \emptyset

pc := set-flag-1

else

S := **S** \cup {j}

if |**S**| = n then **pc** := leave-try

crit_i :

Pre: **pc** = leave-try

Eff: **pc** := crit

exit_i

Eff: **pc** := reset

reset_i :

Pre: **pc** = reset

Eff: **flag(i)** := 0

S := \emptyset

pc := leave-exit

rem_i :

Pre: **pc** = leave-exit



Eff: **pc** := rem

Note on code style

- Explicit `pc` makes atomicity clear, but may look awkward.
- `pc` is often needed in invariants.
- Alternative idea:
 - Use sequential style, with explicit reads or writes (or other operations), one per line.
 - Need line numbers:
 - Play same role as `pc`.
 - Used in invariants: “If process *i* is at line 7 then...”


Correctness

- **Well-formedness:** Obvious.
- **Mutual exclusion:**
 - Argue using event order in executions, instead of invariants as usual.
 - By contradiction: Assume U_i, U_j find themselves in C at the same time.
 - Both must set-flag-2 before entering C ; consider the last time they do this.
 - WLOG, suppose set-flag-2 _{i} comes first.
 - Then $\text{flag}(i) = 2$ from that point onward (until the assumed point when they are both in C).
 - However, j must see $\text{flag}(i) \neq 2$, in order to enter C .
 - Impossible.

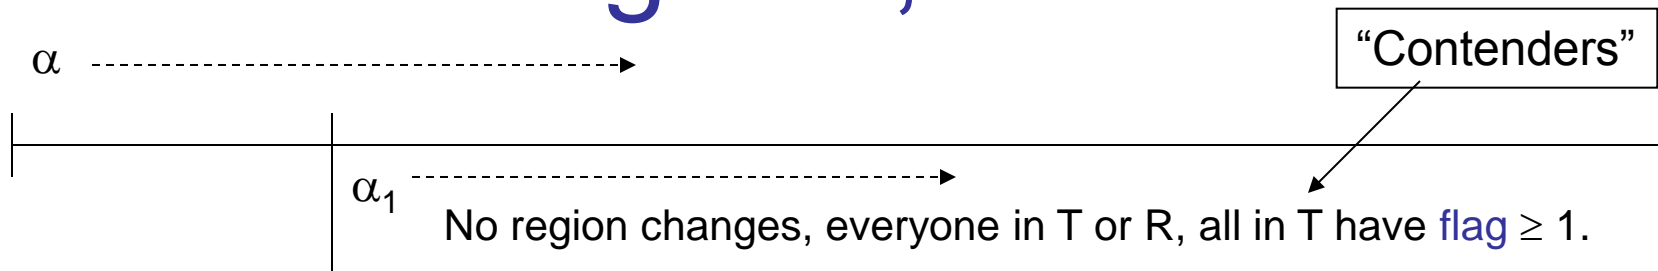


Progress

- Interesting case: Trying region.
- Proof by contradiction:
 - Suppose α is a fair execution, reaches a point where some process is in T, no process is in C, and thereafter, no process ever enters C.
 - Now start removing complications...
 - Eventually, all region changes stop and all in T keep their **flags** ≥ 1 .
 - Then it must be that everyone is in T and R, and all in T have **flag** ≥ 1 .



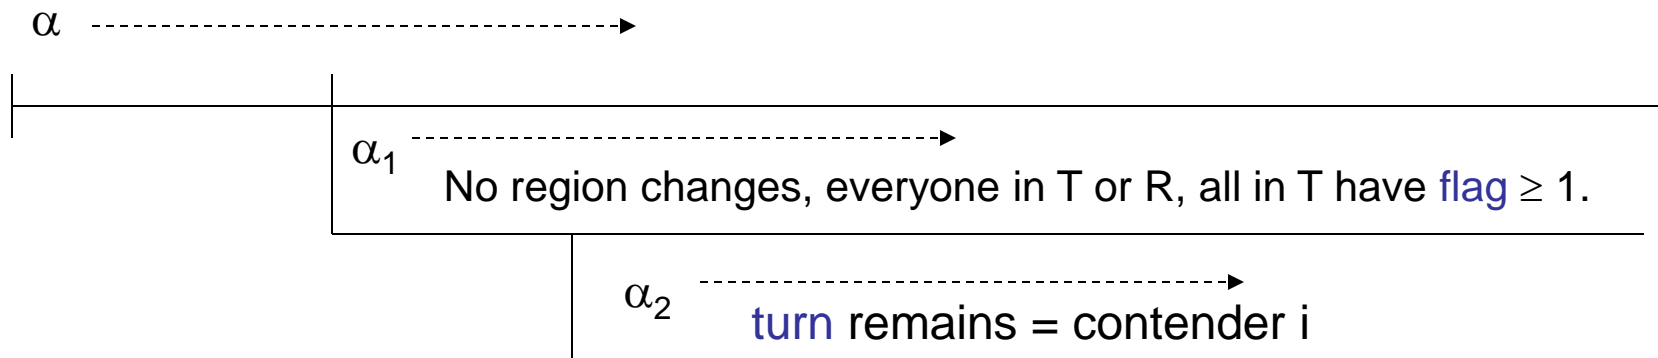
Progress, cont'd



- Then whenever turn is reset in α_1 , it must be set to a contender's index.
- **Claim:** In α_1 , turn eventually acquires a contender's index.
- **Proof:**
 - Suppose not---it stays non-contender forever.
 - Consider any contender i .
 - If it ever reaches test-turn, then it will set $\text{turn} := i$, since it sees a non-contending process, yielding a contradiction.
 - Why must process i reach test-turn?
 - Either that, or it succeeds in reaching C.
 - But we have assumed no one reaches C.

Progress, cont'd

- In α_1 , once **turn** = contender's index, it is thereafter always = some contender's index.
 - Because contenders are the only processes that can change **turn**.
- May change several times.
- Eventually, **turn** stops changing (because tests come out negative), stabilizes to some contender index, say i .



- Thereafter, all contenders $j \neq i$ wind up looping in Stage 1.
 - If j reaches Stage 2, it returns to Stage 1, since it doesn't go to C.
 - But then j 's tests always fail, so j stays in Stage 1.
- But then nothing stops process i from entering C.

Mutual exclusion, Proof 2

- Use invariants.
- Must show they hold after any number of steps.
- Main goal invariant: $|\{i : pc_i = \text{crit}\}| \leq 1$.
- To prove by induction, need more:
 1. If $pc_i = \text{crit}$ (or leave-try or reset) then $|S_i| = n$.
 2. There do not exist $i, j, i \neq j$, with i in S_j and j in S_i .
- 1 and 2 easily imply mutual exclusion.
- **Proof of 1:** Easy induction
- **Proof of 2:**
 - Needs some easy auxiliary invariants saying what S -values go with what $flag$ values and what pc values.
 - Key step: When j gets added to S_i , by $check(j)_i$ event.
 - Then must have $flag(j) \neq 2$.
 - But then $S_j = \emptyset$ (by auxiliary invariant), so $i \notin S_j$, can't break invariant.

Running Time

- Upper bound on time from when **some process** is in T until **some process** is in C.
- Assume upper bound of l on successive turns for each process task (here, all steps of each process are in one task).
- Time upper bound for [Dijkstra]: $O(l n)$.
- **Proof:** LTTR (see p. 275)

Mutual Exclusion Algorithms with Fairness Guarantees [Peterson]

Adding fairness guarantees

- Dijkstra's algorithm does not guarantee fairness in granting the resource to different users.
- This might not be important in practice, if contention is rare.
- Other theoretical algorithms add fairness guarantees.
- **[Peterson]**: a suite of algorithms guaranteeing lockout-freedom.
- **Lockout-freedom**: In any (low-level) fair execution:
 - If all users always return the resource then any user that enters T eventually enters C.
 - Any user that enters E eventually enters R.

Peterson 2-process algorithm

- Shared variables:
 - **turn**, in $\{0,1\}$, 2W2R read/write register, initially arbitrary.
 - for each process i in $\{0,1\}$:
 - **flag(i)**, in $\{0,1\}$, 1W1R register, initially 0
 - Written by i , read by $1-i$.

- Process i 's trying protocol:
 - Sets **flag(i)** := 1, sets **turn** := i .
 - Waits for either **flag(1-i)** = 0 or **turn** $\neq i$.

Other process not active.

Other process has the **turn** variable.

- Toggles between the two tests.
- Exit protocol:
 - Sets **flag(i)** := 0

Precondition/effect code

Shared variables:

$\text{turn} \in \{0,1\}$, initially arbitrary

for every $i \in \{0,1\}$:

$\text{flag}(i) \in \{0,1\}$, initially 0

Actions of process i :

Input: try_i , exit_i

Output: crit_i , rem_i

Internal: set-flag_i , set-turn_i , check-flag_i , check-turn_i , reset_i

Precondition/effect code, Peterson 2P, process i

try_i

Eff: $pc := \text{set-flag}$

set-flag_i

Pre: $pc = \text{set-flag}$

Eff: $\text{flag}(i) := 1$

$pc := \text{set-turn}$

set-turn_i

Pre: $pc = \text{set-turn}$

Eff: $\text{turn} := i$

$pc := \text{check-flag}$

check-flag_i

Pre: $pc = \text{check-flag}$

Eff: if $\text{flag}(1-i) = 0$ then $pc := \text{leave-try}$
else $pc := \text{check-turn}$

check-turn_i

Pre: $pc = \text{check-turn}$

Eff: if $\text{turn} \neq i$ then $pc := \text{leave-try}$
else $pc := \text{check-flag}$

More code, Peterson 2P, process i

crit_i

Pre: **pc** = leave-try

Eff: **pc** := crit

exit_i

Eff: **pc** := reset

reset_i

Pre: **pc** = reset

Eff: **flag(i)** := 0

pc := leave-exit

rem_i

Pre: **pc** = leave-exit

Eff: **pc** := rem

Correctness: Mutual exclusion

- **Key invariant:**
 - If $pc_i \in \{\text{leave-try}, \text{crit}, \text{reset}\}$ (essentially in C), and $pc_{1-i} \in \{\text{check-flag}, \text{check-turn}, \text{leave-try}, \text{crit}, \text{reset}\}$ (engaged in the competition or has won the competition), then $\text{turn} \neq i$.
- That is:
 - If i has won and 1-i is currently competing then turn is set favorably for i---which means it is set to 1-i.
- **Implies mutual exclusion:** If both are in C then turn must be set both ways, contradiction.
- **Proof of invariant:** All cases of inductive step are easy.
 - E.g.: a successful check-turn_i , causing i to advance to leave-try.
 - This explicitly checks that $\text{turn} \neq i$, as needed.

Correctness: Progress

- By contradiction:
 - Suppose someone is in T, and no one is ever thereafter in C.
 - Then the execution eventually stabilizes so no new region changes occur.
 - After stabilization:
 - If exactly one process is in T, then it sees the other's **flag** = 0 and enters C.
 - If both processes are in T, then **turn** is set favorably to one of them, and it enters C.

Correctness: Lockout-freedom

- Argue that neither process can enter C three times while the other stays in T, after setting its **flag** := 1.
- **Bounded bypass.**
- **Proof:** By contradiction.
 - Suppose process i is in T and has set **flag** := 1, and subsequently process $(1-i)$ enters C three times.
 - In each of the second and third times through T, process $(1-i)$ sets **turn** := $1-i$ but later sees **turn** = i .
 - That means process i must set **turn** := i at least twice during that time.
 - But process i sets **turn** := i only once during its one execution of T.
 - Contradiction.
- Bounded bypass + progress imply lockout-freedom.

Time complexity

- Time from when any particular process i enters T until it enters C : $c + O(l)$, where:
 - c is an upper bound on the time any user remains in the critical section, and
 - l is an upper bound on local process step time.
- Detailed proof: See book, p. 283.
- Rough idea:
 - Either process i either enters immediately, or has to wait for $(1-i)$.
 - But in that case, it only has to wait for one critical-section time, since if $(1-i)$ reenters, it will set turn favorably for i .

Peterson n-process algorithms

- Extend 2-process algorithm for lockout-free mutual exclusion to an n-process algorithm, in two ways:
 - Using linear sequence of competitions, or
 - Using binary tree of competitions.

Sequence of competitions

- Competitions $1, 2, \dots, n-1$.
- Competition k has one loser, up to $n-k$ winners.
- Thus, only one can win in competition $n-1$, implying mutual exclusion.
- Shared vars:
 - For each competition k in $\{1, 2, \dots, n-1\}$:
 - $\text{turn}(k)$ in $\{1, 2, \dots, n\}$, mWmR register, written and read by all, initially arbitrary.
 - For i in $\{1, 2, \dots, n\}$:
 - $\text{flag}(i)$ in $\{0, 1, 2, \dots, n-1\}$, 1WmR register, written by i and read by all, initially 0.
- Process i trying protocol:
 - For each level k :
 - Set $\text{flag}(i) := k$, indicating that i is competing at level k .
 - Set $\text{turn}(k) := i$.
 - Wait for either $\text{turn}(k) \neq i$, or everyone else's $\text{flag} < k$ (check flags one at a time).
- Exit protocol:
 - Set $\text{flag}(i) := 0$

Correctness: Mutual exclusion

- **Definition:** Process i is a winner at level k if either:
 - $\text{level}_i > k$, or
 - $\text{level}_i = k = n-1$ and $\text{pc}_i \in \{\text{leave-try}, \text{crit}, \text{reset}\}$.
- **Definition:** Process i is a competitor at level k if either:
 - Process i is a winner at level k , or
 - $\text{level}_i = k$ and $\text{pc}_i \in \{\text{check-flag}, \text{check-turn}\}$.
- **Invariant 1:** If process i is a winner at level k , and process $j \neq i$ is a competitor at level k , then $\text{turn}(k) \neq i$.
- **Proof:** By induction, similar to 2-process case.
 - Complication: More steps to consider.
 - Now have many flags, checked in many steps.
 - Need auxiliary invariants saying something about what is true in the middle of checking a set of flags.

Correctness: Mutual exclusion

- **Invariant 2:** For any k , $1 \leq k \leq n-1$, there are at most $n-k$ winners at level k .
- **Proof:** By induction, on level number, for a particular reachable state (not induction on number of steps).
 - **Basis:** $k = 1$:
 - Suppose false, for contradiction.
 - Then all n processes are winners at level 1.
 - Then Invariant 1 implies that $\text{turn}(1)$ is unequal to all indices, contradiction.
 - **Inductive step:** Assume for k , $1 \leq k \leq n-2$, show for $k+1$.
 - ...

Correctness: Mutual exclusion


- **Invariant 2:** For any k , $1 \leq k \leq n-1$, there are at most $n - k$ winners at level k .
- **Inductive step:** Assume for k , $1 \leq k \leq n-2$, show for $k+1$.
 - Suppose false, for contradiction.
 - Then more than $n - (k + 1)$ processes, that is, at least $n - k$ processes, are winners at level $k + 1$: $|Win_{k+1}| \geq n - k$.
 - Every level $k+1$ winner is also a level k winner: $Win_{k+1} \subseteq Win_k$.
 - By inductive hypothesis, $|Win_k| \leq n-k$.
 - So $Win_{k+1} = Win_k$, and $|Win_{k+1}| = |Win_k| = n - k$.
 - **Q:** What is the value of $turn(k+1)$?
 - Can't be the index of any process in Win_{k+1} , by Invariant 1.
 - Must be the index of some competitor at level $k+1$ (Invariant, LTTR).
 - But every competitor at level $k+1$ is a winner at level k , so is in Win_k .
 - Contradiction, since $Win_{k+1} = Win_k$.

Progress, Lockout-freedom

- Lockout-freedom proof idea:
 - Let k be the highest level at which some process, say i , gets stuck.
 - Then $\text{turn}(k)$ must remain $= i$.
 - That means no one else ever reenters the competition at level k .
 - Eventually, winners from level k will finish, since k is the highest level at which anyone gets stuck.
 - Then all other flags will be $< k$, so i advances.
- Alternatively, prove lockout-freedom by showing a time bound for each process, from $\rightarrow T$ until $\rightarrow C$. (See book)
 - Define $T(0)$ = maximum time from when a process $\rightarrow T$ until $\rightarrow C$.
 - Define $T(k)$, $1 \leq k \leq n-1$ = max time from when a process wins at level k until $\rightarrow C$.
 - $T(n-1) \leq l$.
 - $T(k) \leq 2 T(k+1) + c + (3n+2) l$, by detailed analysis.
 - Solve recurrences, get exponential bound, good enough for showing lockout-freedom.


Peterson Tournament Algorithm

- Assume $n = 2^h$.
- Processes = leaves of binary tree of height h .
- **Competitions** = internal nodes, labeled by binary strings.
- Each process engages in $\log n$ competitions, following path up to root.
- Each process i has:
 - A unique competition x at each level k .
 - A unique role in x (0 = left, 1 = right).
 - A set of potential opponents in x .



Peterson Tournament Algorithm

- Shared variables:
 - For each process i , $\text{flag}(i)$ in $\{0, \dots, h\}$, indicating level, initially 0
 - For each competition x , $\text{turn}(x)$, a Boolean, initially arbitrary.
- Process i 's trying protocol: For each level k :
 - Set $\text{flag}(i) := k$.
 - Set $\text{turn}(x) := b$, where:
 - x is i 's level k competition,
 - b is i 's "role", 0 or 1
 - Wait for either:
 - $\text{turn}(x) = \text{opposite role}$, or
 - all flags of potential opponents in x are $< k$.
- Exit protocol:
 - Set $\text{flag}(i) := 0$.



Correctness

- **Mutual exclusion:**
 - Similar to before.
 - Key invariant: At most one process from any particular subtree rooted at level k is currently a winner at level k .
- **Time bound** (from $\rightarrow T$ until $\rightarrow C$): $(n-1) c + O(n^2 l)$
 - Implies progress, lockout-freedom.
 - Define: $T(0) = \text{max time from } \rightarrow T \text{ until } \rightarrow C$.
 - $T(k), 1 \leq k \leq \log n = \text{max time from winning at level } k \text{ until } \rightarrow C$.
 - $T(\log n) \leq l$.
 - $T(k) \leq 2 T(k+1) + c + (2^{k+1} + 2^k + 7) l$ (see book).
 - Roughly: Might need to wait for a competitor to reach C , then finish C , then for yourself to reach C .
 - Solve recurrences.

Bounded Bypass?

- Peterson's Tournament algorithm has a low time bound from $\rightarrow T$ until $\rightarrow C$:
$$(n - 1) c + O(n^2 l)$$
- Implies lockout-freedom, progress.
- **Q:** Does it satisfy bounded bypass?
- **No!** There's no upper bound on the number of times one process could bypass another in the trying region. E.g.:
 - Process 0 enters, starts competing at level 1, then pauses.
 - Process 7 enters, quickly works its way to the top, enters C, leaves C.
 - Process 7 enters again...repeats any number of times.
 - All while process 0 is paused.
- No contradiction between small time bound and unbounded bypass.
 - Because of the way we're modeling timing of asynchronous executions, using upper bound assumptions.
 - When processes go at very different speeds, we say that the slow processes are going at normal speed, faster processes are going very fast.

Lamport's Bakery Algorithm



Lamport's Bakery Algorithm

- Like taking tickets for service in a bakery.
- Nice features:
 - Uses only single-writer, multi-reader registers.
 - Extends to even weaker (“safe”) registers, in which operations have durations, and a read that overlaps a write receives an arbitrary response.
 - Guarantees lockout-freedom, in fact, almost-FIFO behavior.
- But:
 - Registers are of unbounded size.
 - Algorithm can be simulated using bounded registers, but not easily (uses “Bounded Concurrent Timestamps”).
- Shared variables:
 - For each process i :
 - $\text{choosing}(i)$, a Boolean, written by i , read by all, initially 0
 - $\text{number}(i)$, a natural number, written by i , read by all, initially 0

Bakery Algorithm

- First part, up to $\text{choosing}(i) := 0$ (the “Doorway”, D):
 - Process i chooses a number $>$ all the numbers it reads for the other processes; writes this in $\text{number}(i)$.
 - While doing this, keeps $\text{choosing}(i) = 1$.
 - Two processes could choose the same number (unlike in a real bakery).
 - Break ties with process ids.
- Second part:
 - Wait to see that no others are choosing, and no one else has a smaller number.
 - That is, wait to see that your ticket is the smallest.
 - Never go back to the beginning of this part---just proceed step by step, waiting when necessary.

Code

Shared variables:

for every $i \in \{1, \dots, n\}$:

$\text{choosing}(i) \in \{0, 1\}$, initially 0, writable by i , readable by all $j \neq i$

$\text{number}(i)$, a natural number, initially 0, writable by i , readable by $j \neq i$.

try_i

$\text{choosing}(i) := 1$

$\text{number}(i) := 1 + \max_{j \neq i} \text{number}(j)$

$\text{choosing}(i) := 0$

for $j \neq i$ do

 waitfor $\text{choosing}(j) = 0$

 waitfor $\text{number}(j) = 0$ or $(\text{number}(i), i) < (\text{number}(j), j)$

crit_i

exit_i

$\text{number}(i) := 0$

rem_i

Correctness: Mutual exclusion

- **Key invariant:** If process i is in C , and process $j \neq i$ is in $(T - D) \cup C$,


Trying region after doorway, or critical region

then $(\text{number}(i), i) < (\text{number}(j), j)$.

- **Proof:**
 - Could prove by induction.
 - Instead, give argument based on events in executions.

Correctness: Mutual exclusion

- **Invariant:** If i is in C , and $j \neq i$ is in $(T - D) \cup C$, then $(\text{number}(i), i) < (\text{number}(j), j)$.
- **Proof:**
 - Consider a point where i is in C and $j \neq i$ is in $(T - D) \cup C$.
 - Then before i entered C , it must have read $\text{choosing}(j) = 0$, event π .



- **Case 1:** j sets $\text{choosing}(j) := 1$ (starts choosing) after π .
 - Then $\text{number}(i)$ is set before j starts choosing.
 - So j sees the “correct” $\text{number}(i)$ and chooses something bigger.
 - That suffices.
- **Case 2:** j sets $\text{choosing}(j) := 0$ (finishes choosing) before π .
 - Then when i reads $\text{number}(j)$ in its second waitfor loop, it gets the “correct” $\text{number}(j)$.
 - Since i decides to enter C , it must see $(\text{number}(i), i) < (\text{number}(j), j)$.

Correctness: Mutual exclusion

- **Invariant:** If i is in C , and $j \neq i$ is in $(T - D) \cup C$, then $(\text{number}(i), i) < (\text{number}(j), j)$.
- **Proof of mutual exclusion:**
 - Apply invariant both ways.
 - Contradictory requirements.

Liveness Conditions

- **Progress:**
 - By contradiction.
 - If not, eventually region changes stop, leaving everyone in T or R, and at least one process in T.
 - Everyone in T eventually finishes choosing.
 - Then nothing blocks the smallest (number, index) process from entering C.
- **Lockout-freedom:**
 - Consider any i that enters T.
 - Suppose for contradiction that i never reaches C.
 - Eventually it finishes the doorway.
 - Thereafter, any newly-entering process picks a bigger number.
 - Progress implies that some processes continue to enter C, as long as i is still in T.
 - In fact, this must happen infinitely many times!
 - But those with bigger numbers can't get past i , contradiction.

FIFO Condition

- Not really FIFO ($\rightarrow T$ vs. $\rightarrow C$), but almost:
 - **FIFO after the doorway**: if j leaves D before $i \rightarrow T$, then $j \rightarrow C$ before $i \rightarrow C$.
- But the “doorway” is an artifact of this algorithm, so this isn’t a meaningful way to evaluate the algorithm!
- Maybe say “there exists a doorway such that”...
- But then we could take D to be the entire trying region, making the property trivial.
- To make the property nontrivial:
 - Require D to be “wait-free”: a process is guaranteed to complete D if it keeps taking steps, regardless of what other processes do.
 - D in the Bakery Algorithm is wait-free.
- The algorithm is **FIFO after a wait-free doorway**.

Impact of Bakery Algorithm

- Originated some important ideas:
 - Wait-freedom
 - Fundamental notion for theory of fault-tolerant asynchronous distributed algorithms.
 - Weakly coherent memories
 - Beginning of formal study: definitions, and some algorithmic strategies for coping with them.



Next time...

- More Mutual Exclusion algorithms:
 - Burns' algorithm
- Number of registers needed for mutual exclusion.
- Reading:
 - Sections 10.6,10.8
- Generalized resource allocation and exclusion problems
- Reading:
 - Chapter 11