

6.852: Distributed Algorithms


Fall, 2015

Lecture 17

Today's plan


- Dining Philosophers (briefly)
- **Reading:** Chapter 11
- Asynchronous shared-memory systems with failures.
- Consensus in asynchronous shared-memory systems.
- Impossibility of consensus [Fischer, Lynch, Paterson]
- **Reading:** Chapter 12
- **Next:**
 - Atomic objects
 - Reading: Chapter 13

Dining Philosophers




Generalized resource allocation

- **Mutual exclusion:** Problem of allocating a single non-sharable resource.
- Can generalize to more resources, some sharing.
- **Exclusion specification \mathcal{E}** (for a given set of users):
 - Any collection of sets of users, closed under superset.
 - Expresses which users are incompatible, can't coexist in the critical section.
- **Example: k-exclusion** (any k users are OK, but not $k+1$)
 - $\mathcal{E} = \{ E : |E| > k \}$
- **Example: Reader-writer locks**
 - Relies on classification of users as readers vs. writers.
 - $\mathcal{E} = \{ E : |E| > 1 \text{ and } E \text{ contains a writer} \}$
- **Example: Dining Philosophers (Dijkstra)**
 - $\mathcal{E} = \{ E : E \text{ includes a pair of neighbors} \}$



Resource specifications

- Some exclusion specs can be described conveniently in terms of requirements for concrete resources.
- **Resource specification:** Each user needs a certain particular subset of the resources.
- Can't share: Users with intersecting sets exclude each other.




- **Example: Dining Philosophers (Dijkstra)**
 - $\mathcal{E} = \{ E : E \text{ includes a pair of neighbors} \}$
 - Forks (resources) between adjacent philosophers; each needs both adjacent forks in order to eat.
 - Only one can hold a particular fork at a time, so adjacent philosophers must exclude each other.
- Not every exclusion problem can be expressed in this way.
 - k-exclusion cannot.

Resource allocation problem, for a given exclusion specification \mathcal{E}

- Same shared-memory architecture as for mutual exclusion (processes and shared variables).
- **Well-formedness:** As before.
- **Exclusion:** There is no reachable state in which the set of users in C is a set in \mathcal{E} .
- **Progress:** As before.
- **Lockout-freedom:** As before.
- But these don't capture concurrency requirements.
 - Any lockout-free mutual exclusion algorithm also satisfies all these conditions (provided that \mathcal{E} doesn't contain any singleton sets).
- Can add concurrency conditions, e.g.:
 - **Independent progress:** If $i \in T$ and every j that could conflict with i remains in R , then eventually $i \rightarrow C$.
 - **Time bound:** Obtain better bounds from $i \rightarrow T$ to $i \rightarrow C$, even in the presence of conflicts, than we can for mutual exclusion.


Dining Philosophers Problem

- Like Mutual Exclusion, with a different exclusion condition:
 - No two neighbors are in C at once (exclusion specification), or
 - Forks on edges, each philosopher needs both adjacent forks to eat (explicit resource specification).
- Can use progress and fairness conditions as for Mutex.
- Can add new conditions to capture concurrent access to C.
- Dijkstra posed the problem, gave a solution for a strong shared-memory model.
 - Globally-shared variables, atomic access to all of shared memory.
- Distributed version: Assume the only shared variables are RMW variables corresponding to the forks, accessible only by processes at the endpoints.




Impossibility Result

- **Theorem 1:** If all processes are identical and refer to forks by local names L and R, and all shared variables have the same initial value, then we can't guarantee DP exclusion + progress.
- **Proof:** Can't break symmetry:
 - Consider only executions that work in synchronous rounds, prove by induction on rounds that symmetry is preserved.
 - Assume all processes $\rightarrow T$.
 - By progress, someone $\rightarrow C$.
 - By symmetry, all do, violating DP exclusion.
- **Example symmetric algorithm:** Wait for R fork first, then L fork.
 - Guarantees DP exclusion.
 - Progress fails---all processes might get R fork, then wait forever for L fork (deadlock).




DP Algorithms


- So we need a mechanism to break symmetry.
- **Solution 1:** Number forks in increasing order around the table; every process picks up its smaller numbered fork first.
 - Yields DP exclusion, progress, lockout freedom, independent progress.
 - But the time isn't good---we can have a long “waiting chain” of processes waiting for neighbors to release forks.



Creating a long waiting chain




DP Algorithms




- **Solution 2:** Right/Left algorithm (**Burns**):
 - Classify processes as R or L (at least one of each).
 - R processes pick up right fork first, L processes pick up left fork first.
 - Yields DP exclusion, progress, lockout freedom, independent progress.
 - In even-sized rings in which R and L alternate, the lengths of waiting chains are limited to 2.
 - Yields a good (constant) time bound, LTTR.
- Generalize to solve any resource problem:
 - Represent the problem as an undirected graph.
 - Nodes = resources.
 - Edge between two resources if some user wants both.
 - Color the nodes of the graph; order colors.
 - All processes acquire resources in order of colors.

Asynchronous shared-memory systems with failures



Asynchronous shared-memory systems with failures

- Process stopping failures.
- Architecture as for Mutual Exclusion.
 - Processes + shared variables, one system automaton.
 - Users
- Add **stop_i** inputs.
 - Effect is to disable all future locally controlled actions of process *i*.
- Fair executions:
 - Every process that doesn't fail gets infinitely many turns to perform locally-controlled steps.
 - Just ordinary fairness---**stop** means that nothing further is enabled.
 - Users also get turns.



The consensus problem in asynchronous shared-memory systems with failures

Consensus in Asynchronous Shared-Memory Systems



- Recall: Consensus in synchronous networks.
 - Algorithms for stopping failures:
 - FloodSet, FloodMin, Optimizations: $f+1$ rounds, any number of processes, low communication
 - Lower bounds: $f+1$ rounds
 - Algorithms for Byzantine failures
 - EIG: $f+1$ rounds, $n > 3f$, exponential communication
 - Lower bounds: $f+1$ rounds, $n > 3f$
- Asynchronous networks: Impossible
- Asynchronous shared memory:
 - Read/write variables: Impossible
 - Read-modify-write variables: Simple algorithms exist
- Impossibility results hold even if n is large and f is just 1!

Consequences of impossibility results

- Can't solve problems like transaction commit, agreement on choice of leader, fault diagnosis,...in the purely asynchronous model with failures.
- But these problems must be solved...what to do?
- Can strengthen the assumptions:
 - Rely on timing assumptions: Upper and lower bounds on message delivery time, on step time.
 - Probabilistic assumptions
- And/or weaken the guarantees:
 - Allow a small probability of violating safety properties, or of not terminating.
 - Conditional termination, based on stability for a “sufficiently long” interval of time.
- We'll see some of these strategies.
- But, first, the impossibility result!

Architecture

- V , set of consensus values
- Interaction between user U_i and process (agent) p_i :
 - User U_i submits initial value v with $\text{init}(v)_i$.
 - Process p_i returns decision v with $\text{decide}(v)_i$.
 - I/O handled slightly differently from synchronous setting, where inputs and outputs were in local variables.
 - Assume each user performs at most one $\text{init}(v)_i$ in an execution.
- Shared variable types:
 - Read/write registers (for now)





Requirements 1

- **Well-formedness:**
 - At most one $\text{decide}()_i$, appears, and only after an $\text{init}()_i$.
- **Agreement:**
 - All decision values are identical.
- **Validity:**
 - If all init actions that occur contain the same v , then that v is the only possible decision value.
 - Stronger version: Any decision value is someone's initial value.
- **Termination:**
 - Failure-free termination (most basic requirement):
 - In any fair failure-free (ff) execution in which init events occur on all "ports", decide events occur on all ports.
- **Basic problem requirements:** Well-formedness, agreement, validity, failure-free termination.

Requirements 2: Fault-tolerance

- **Wait-free termination** (strongest condition):
 - In any fair execution in which init events occur on all ports, a decide event occurs on every port i for which no stop_i occurs.
 - Similar to wait-free doorway in Lamport's Bakery algorithm: says i finishes regardless of whether the other processes stop or not.
- Also consider tolerating a limited number of failures.
- Should be easier to achieve, so impossibility results are stronger.
- **f -failure termination, $0 \leq f \leq n$:**
 - In any fair execution in which init events occur on all ports, **if there are stop events on at most f ports**, then a decide event occurs on every port i for which no stop_i occurs.
- Wait-free termination = n -failure termination = $(n-1)$ -failure termination.
- **1-failure termination:** The interesting special case we will consider in our main proof.

Impossibility results for consensus in asynchronous shared-memory systems with failures



Impossibility of agreement

- **Main Theorem** [Fischer, Lynch, Paterson], [Loui, Abu-Amara]:
 - For $n \geq 2$, there is no algorithm in the read/write shared memory model that solves the agreement problem and guarantees **1-failure termination**.
- **A Weaker Theorem** [Herlihy]:
 - For $n \geq 2$, there is no algorithm in the read/write shared memory model that solves the agreement problem and guarantees **wait-free termination**.
- We'll prove Herlihy's theorem first.

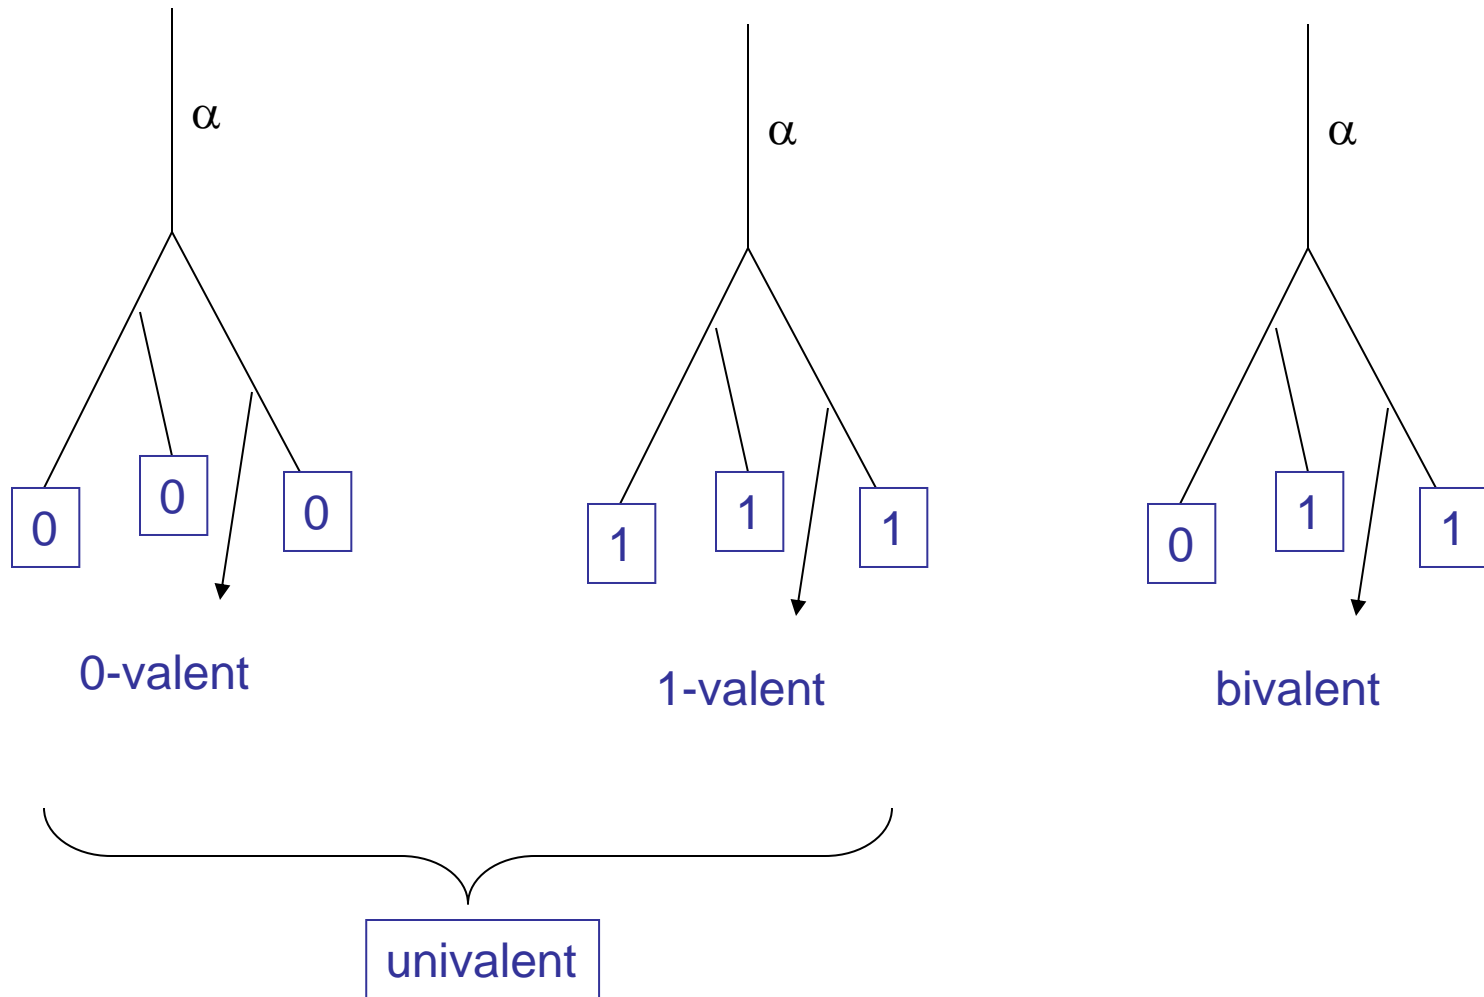
Restrictions (WLOG)

- $V = \{ 0, 1 \}$
- One task per process
- Processes are deterministic:
 - Unique start state.
 - From any state, any process has at most one locally-controlled action enabled.
 - From any state, for any enabled action, there is exactly one new state.
- Non-halting:
 - Every non-failed process always has some locally-controlled action enabled, even after it decides.

Terminology

- **Initialization:**
 - Sequence of n init steps, one per port, in index order: $\text{init}(v_1)_1, \text{init}(v_2)_2, \dots, \text{init}(v_n)_n$
- **Input-first execution:**
 - Begins with an initialization.
- A finite execution α is:
 - **0-valent**, if 0 is the only decision value appearing in α or any extension of α , and 0 actually does appear in α or some extension.
 - **1-valent**, if 1 is the only decision value appearing in α or any extension of α , and 1 actually does appear in α or some extension.
 - **Univalent**, if α is 0-valent or 1-valent.
 - **Bivalent**, if each of 0, 1 occurs in some extension of α .

Univalence and Bivalence



Exhaustive classification

- **Lemma 1:**
 - If A solves agreement with ff-termination, then each finite ff execution of A is either univalent or bivalent.
- **Proof:**
 - Can extend to a fair execution, in which everyone is required to decide.


Bivalent Initializations

Bivalent initializations


- From now on, fix A to be an algorithm solving agreement with (at least) 1-failure termination.
 - Could also satisfy stronger conditions, like f -failure termination, or wait-free termination.
- **Lemma 2:** A has a bivalent initialization.
- That is, the final decision value cannot always be determined from the inputs only.
- Contrast: In non-fault-tolerant case, final decision can be determined from the inputs only; e.g., take majority.
- **Proof:**
 - Same argument used (later) by [Aguilera, Toueg].
 - Suppose not. Then all initializations are univalent.
 - Define initializations $\alpha_0 = \text{all 0s}$, $\alpha_1 = \text{all 1s}$.
 - α_0 is 0-valent, α_1 is 1-valent, by validity.

Bivalent initializations

- A solves agreement with 1-failure termination.
- **Lemma 2:** A has a bivalent initialization.
- **Proof, cont'd:**
 - Construct chain of initializations, spanning from α_0 to α_1 , each differing in the initial value of just one process.
 - There must be two consecutive initializations, say α and α' , where α is 0-valent and α' is 1-valent.
 - Differ only in initial value of some process i .
 - Consider a fair execution extending α , in which i fails right after α .
 - All but i must eventually decide, by 1-failure termination; since α is 0-valent, all must decide 0.
 - Extend α' in the same way, all but i still decide 0, by indistinguishability.
 - Contradicts 1-valence of α' .



Impossibility for Wait-Free Termination




Impossibility for wait-free termination

- **Weaker Theorem [Herlihy]:**
 - For $n \geq 2$, there is no algorithm in the read/write shared memory model that solves the agreement problem and guarantees wait-free termination.
- **Proof:**
 - We already assumed A solves agreement with 1-failure termination (which yielded a bivalent initialization).
 - Now assume, for contradiction, that A (also) satisfies the stronger wait-free termination condition.
 - Proof is based on pinpointing exactly how a decision gets determined, that is, how the execution moves from bivalence to univalence.

Impossibility for wait-free termination

- **Definition:** A **decider execution** α is a finite, failure-free, input-first execution such that:
 - α is bivalent.
 - For every i , $\text{ext}(\alpha, i)$ is univalent.

Extension of α with one step of i




- **Lemma 3:** A (with wait-free termination) has a decider execution.

Impossibility for wait-free termination

- **Lemma 3:** A (with w-f termination) has a decider.



- **Proof:**

- Suppose not. Then any bivalent ff input-first execution has a 1-step bivalent ff extension.
- Start with a bivalent initialization (Lemma 2), and produce an infinite ff execution α all of whose prefixes are bivalent.
 - At each stage, start with a bivalent ff input-first execution and extend by one step to another bivalent ff execution.
 - Possible by assumption.
- α must contain infinitely many steps of some process, say i .
- Claim i must decide in α :
 - Add stop events for all processes that take only finitely many steps.
 - Result is a fair execution α' .
 - Wait-free termination says i must decide in α' .
 - α is indistinguishable from α' , by i , so i must decide in α also.
- Contradicts bivalence.





Impossibility for wait-free termination

- **Proof of theorem, cont'd:**
 - Fix a decider, α .
 - Since α is bivalent and all 1-step extensions are univalent, there must be two processes, say i and j , leading to 0-valent and 1-valent states, respectively.
 - Case analysis yields a contradiction:
 1. i 's step is a read
 2. j 's step is a read
 3. Both writes, to different variables.
 4. Both writes, to the same variable.





Case 1: i's step is a read

- Run all but i after $\text{ext}(\alpha, j)$.
- Looks like a fair execution in which i fails.
- So all others must decide; since $\text{ext}(\alpha, j)$ is 1-valent, they decide 1.
- Now run the same extension, starting with j's step, after $\text{ext}(\alpha, i)$.
- They cannot see i's read.
- So they behave the same, decide 1.
- Contradicts 0-valence of $\text{ext}(\alpha, i)$.




Case 2: j 's step is a read

- Symmetric.



Case 3: Writes to different shared variables


- Then the two steps are completely independent.
- Could be performed in either order, and the result should be the same.
- $\text{ext}(\alpha, ij)$ and $\text{ext}(\alpha, ji)$ are indistinguishable to **all** processes, and end up in the same system state.



- But $\text{ext}(\alpha, ij)$ is 0-valent, since it extends the 0-valent execution $\text{ext}(\alpha, i)$.
- And $\text{ext}(\alpha, ji)$ is 1-valent, since it extends the 1-valent execution $\text{ext}(\alpha, j)$.
- Contradictory requirements.

Case 4: Writes to the same shared variable x .


- Run all but i after $\text{ext}(\alpha, j)$; they must decide.
- Since $\text{ext}(\alpha, j)$, is 1-valent, they decide 1.
- Run the same extension, starting with j 's step, after $\text{ext}(\alpha, i)$.
- Cannot see i 's write to x , because j 's write overwrites it.
- They behave the same, decide 1.
- Contradicts 0-valence of $\text{ext}(\alpha, i)$.



Impossibility for wait-free termination

- So we have proved:
- **Weaker Theorem:** [Herlihy]
 - For $n \geq 2$, there is no algorithm in the read/write shared memory model that solves the agreement problem and guarantees wait-free termination.

Impossibility for 1-failure termination




Impossibility for 1-failure termination

- **Q:** Why doesn't the previous proof yield impossibility for 1-failure termination?
- Lemma 2 (bivalent initialization) works for $f = 1$.
- In the proof of Lemma 3 (existence of decider), wait-free termination is used to say that a process i must decide in **any fair execution in which i doesn't fail**.
- 1-failure termination makes a termination guarantee only when **at most one process fails**.
- **Main Theorem:**
 - For $n \geq 2$, there is no algorithm in the read/write shared memory model that solves the agreement problem and guarantees 1-failure termination.


Impossibility for 1-failure termination

- From now on, assume A satisfies 1-failure termination, not necessarily wait-free termination (weaker requirement).
- Initialization lemma still works:
 - **Lemma 2:** A has a bivalent initialization.
- New key lemma, replacing Lemma 3:
- **Lemma 4:** If α is any bivalent, ff, input-first execution of A , and i is any process, then there is some ff-extension α' of α such that $\text{ext}(\alpha', i)$ is bivalent.





Lemma 4 \Rightarrow Main Theorem

- **Lemma 4:** If α is any bivalent, ff, input-first execution of A , and i is any process, then there is some ff-extension α' of α such that $\text{ext}(\alpha', i)$ is bivalent.
- **Proof of Main Theorem using Lemma 4:**
 - Construct a fair, ff, input-first execution in which no process ever decides, contradicting the basic ff-termination requirement.
 - Start with a bivalent initialization.
 - Then cycle through the processes round-robin: 1, 2, ..., n , 1, 2, ...
 - At each step, say for i , use Lemma 4 to extend the execution, including at least one step of i , while maintaining bivalence and avoiding failures.
 - Contradiction.




Proof of Lemma 4

- **Lemma 4:** If α is any bivalent, ff, input-first execution of A , and i is any process, then there is some ff-extension α' of α such that $\text{ext}(\alpha', i)$ is bivalent.
- **Proof:**
 - By contradiction. Suppose there is some bivalent, ff, input-first execution α of A and some process i , such that for every ff extension α' of α , $\text{ext}(\alpha', i)$ is univalent.
 - In particular, $\text{ext}(\alpha, i)$ is univalent, WLOG 0-valent.
 - Since α is bivalent, there is some extension of α in which someone decides 1, WLOG failure-free.




Proof of Lemma 4

- There is some ff-extension of α in which someone decides 1.
- Consider letting i take one step at each point along the “spine”.
- By assumption, results are all univalent.
- 0-valent at the beginning, 1-valent at the end.
- So there are two consecutive results, one 0-valent and the other 1-valent:
- A new kind of “decider”.






New “Decider”

- **Claim:** $j \neq i$.
- **Proof:**
 - If $j = i$ then:
 - 1 step of i yields 0-valence
 - 2 steps of i yield 1-valence
 - But process i is deterministic, so this can't happen.
 - “Child” of a 0-valent state can't be 1-valent.
- The rest of the proof is a case analysis, similar to before...





Case 1: i's step is a read

- Run j after i.
- Executions ending with ji and ij are indistinguishable to everyone but i (because this is a read step of i).
- Run all processes except i in the same order after ji and ij.
- In each case, they must decide, by 1-failure termination.
- After ji, they must decide 1.
- After ij, they must decide 0.
- But indistinguishable, contradiction!





Case 2: j 's step is a read

- Executions ending with j and i are indistinguishable to everyone but j (because this is a read step of j).
- Run all processes except j in the same order after j and i .
- In each case, they must decide, by 1-failure termination.
- After j , they must decide 1.
- After i , they must decide 0.
- But indistinguishable, contradiction!





Case 3: Writes to different shared variables

- As for the wait-free case.
- The steps of i and j are independent, could be performed in either order, indistinguishable to everyone.
- But the execution ending with ji is 1-valent, whereas the execution ending with ij is 0-valent.
- Contradiction.



Case 4: Writes to the same shared variable x.

- As for Case 2.
- Executions ending with j_i and i are indistinguishable to everyone but j (because i overwrites the write step of j).
- Run all processes except j in the same order after j_i and i .
- After j_i , they must decide 1.
- After i , they must decide 0.
- Indistinguishable, contradiction!



Impossibility for 1-failure termination

- So we have proved:
- **Main Theorem:** [Fischer, Lynch, Paterson]
[Loui, Abu-Amara]
 - For $n \geq 2$, there is no algorithm in the read/write shared memory model that solves the agreement problem and guarantees 1-failure termination.

Variations, extensions,
significance,...

Extension to networks

- Result also holds in asynchronous networks---see this soon.
- [Fischer, Lynch, Paterson 82, 85] proved first for networks; 2001 Dijkstra Prize.
- [Loui, Abu-Amara 87] extended result and proof to shared memory.

Significance of FLP impossibility result

- For distributed computing practice:
 - Reaching agreement is important in practice:
 - Agreeing on aircraft altimeter readings.
 - Database transaction commit.
 - Agreeing on updates to data replicas.
 - Agreeing on BitCoin transaction logs.
 - FLP shows limitations on the kind of algorithm one can look for.
- For distributed computing theory:
 - Variations:
 - [Loui, Abu-Amara 87] Read/write shared memory.
 - [Herlihy 91] Stronger fault-tolerance requirement (wait-free termination); simpler proof.
 - Circumventing the impossibility result:
 - Strengthening the assumptions.
 - Weakening the guarantees.

Strengthening the assumptions

- Using limited timing information [Dolev, Dwork, Stockmeyer 87].
 - Bounds on message delays, processor step time.
 - Makes the model more like the synchronous model.
- Using randomness [Ben-Or 83][Rabin 83] [Attiya, Censor].
 - Allow random choices in local transitions.
 - Also weakens guarantees:
 - Small probability of a wrong decision, or
 - Small probability of not terminating, in any bounded time (Probability of not terminating approaches 0 as time approaches infinity.)

Weakening the guarantees

- Agreement, validity must always hold.
- Termination guaranteed if system behavior “stabilizes”:
 - No new failures.
 - Timing (of process steps, messages) within “normal” bounds.
- Good solutions have been developed, both theoretical and practical.
- [Dwork, Lynch, Stockmeyer 88] Dijkstra Prize, 2007
 - Keeps trying to choose a leader, who tries to coordinate agreement.
 - Coordination attempts can fail.
 - Once system stabilizes, a unique leader is chosen, coordinates agreement.
 - Tricky part: Ensuring failed attempts don’t lead to inconsistent decisions.
- [Lamport 89] Paxos algorithm.
 - Improves on [DLS] by allowing more concurrency.
 - Refined, engineered for practical use.

Weakening the guarantees

- Agreement, validity must always hold.
- Termination required if system behavior “stabilizes”:
 - No new failures.
 - Timing (of process steps, messages) within “normal” bounds.
- Good solutions, both theoretically and in practice.
- [Dwork, Lynch, Stockmeyer 88]: Dijkstra Prize, 2007
- [Lamport 89] Paxos algorithm.
- [Chandra, Hadzilacos, Toueg 96] Failure detectors (FDs)
 - Services that encapsulate use of time for detecting failures.
 - Develop similar algorithms to [DLS 88] and [Lamport 89] using FDs.
 - Studied properties of FDs, identified weakest FD to solve consensus.

Extension to k-consensus

- At most k different decisions may occur overall.
- Solvable for $k-1$ process failures but not for k failures.
 - Algorithm for $k-1$ failures: [Chaudhuri 93].
 - Impossibility result:
 - [Herlihy, Shavit 93], [Borowsky, Gafni 93], [Saks, Zaharoglu 93]
 - Godel Prize, 2004.
 - Techniques from algebraic topology: Sperner's Lemma.
 - Similar to those used for lower bound on rounds for k -agreement, in synchronous model.
- Question (recent results):
 - What is the weakest failure detector to solve k -consensus with k failures?

Importance of the read/write data type

- Consensus impossibility result doesn't hold for more powerful data types.
- **Example:** Read-modify-write shared memory
 - Very strong primitive.
 - In one step, can read variable, do local computation, and write back a value.
 - Easy algorithm:
 - One shared variable x , value in $V \cup \{\perp\}$, initially \perp .
 - Each process i accesses x once.
 - If it sees:
 - \perp , then it changes the value in x to its own initial value and decides on that value.
 - Some v in V , then it decides on that value.
- Read/write registers are similar to asynchronous FIFO reliable channels---we'll see the precise connection later.

Next time...

- Atomic objects
- **Reading:** Chapter 13