



# 6.852: Distributed Algorithms

## Fall, 2015

Class 18

# Impossibility results for consensus in asynchronous shared-memory systems with failures



# Impossibility of agreement

- **Main Theorem** [Fischer, Lynch, Paterson], [Loui, Abu-Amara]:
  - For  $n \geq 2$ , there is no algorithm in the read/write shared memory model that solves the agreement problem and guarantees **1-failure termination**.
- **A Weaker Theorem** [Herlihy]:
  - For  $n \geq 2$ , there is no algorithm in the read/write shared memory model that solves the agreement problem and guarantees **wait-free termination**.

# Importance of the read/write data type


- Consensus impossibility result doesn't hold for more powerful data types.
- **Example:** Read-modify-write shared memory
  - Very strong primitive.
  - In one step, can read variable, do local computation, and write back a value.
  - Easy algorithm:
    - One shared variable  $x$ , value in  $V \cup \{\perp\}$ , initially  $\perp$ .
    - Each process  $i$  accesses  $x$  once.
    - If it sees:
      - $\perp$ , then it changes the value in  $x$  to its own initial value and decides on that value.
      - Some  $v$  in  $V$ , then it decides on that value.
- Read/write registers are similar to asynchronous FIFO reliable channels---we'll see the precise connection later.

# Today's plan

- Atomic objects:
  - Basic definitions
  - Canonical atomic objects
  - Atomic objects vs. shared variables
- Reading: Sections 13.1-13.2
- Next time:
  - Algorithms to implement atomic objects:
    - Atomic snapshots
    - Atomic read/write registers
  - Reading: Sections 13.3-13.4


# Atomic Objects

- Atomic objects are fundamental building blocks for fault-tolerant distributed systems and multiprocessor systems.
- An atomic object is a version of a shared variable, with **invocation and response events separated**, rather than being combined into one indivisible event.
- Also consider (stopping) failures.
- Separating invocations and responses allows us to consider lower-level implementations of these objects.
  - These can use shared memory, or distributed network algorithms.
  - For shared memory, we can “layer” the development through several levels.



# Shared memory model


- A single I/O automaton with processes and variables “inside”.
  - Separation is expressed by locality restrictions on the actions and transitions.
  - Processes and variables aren’t separate automata.
  - Doesn’t exploit I/O automaton composition.
  - We can’t talk about “implementing” shared variables with lower-level distributed algorithms.



- **Q:** Can we model each process and variable as a separate I/O automaton?
  - Split operations on variables into separate invocation and response actions.
  - But we still want an (invocation, response) pair to “look like” an instantaneous access.

# Atomic Objects


- **Q:** Can we model each process and variable as a separate I/O automaton?
- Define **atomic objects**.
- Atomic object of a given type is similar to an ordinary shared variable of that type, but it also allows concurrent accesses.
- Interface has invocation inputs and response outputs.
- Invocation/response behavior “looks like” that of an instantaneous-access shared variable.
- Looks “as if” operations occur one at a time, sequentially, in some order consistent with order of invocations and responses.
- AKA **linearizable objects** [Herlihy, Wing]







# Atomic Objects

- Interface has invocation inputs and response outputs.
- Looks as if operations occur one at a time, sequentially, in some order consistent with order of invocations and responses.
- Also, consider fault-tolerance conditions (stopping failures only), as for consensus:
  - Wait-free termination
  - $f$ -failure termination
  - Etc.
- Separating invocations and responses allows us to consider lower-level implementations of atomic objects.
  - Shared-memory algorithms, or distributed network algorithms.
  - For shared memory algorithms, we can develop algorithms hierarchically, using several levels.
- Atomic objects are important building blocks for multiprocessor systems and distributed systems.





# Replacing variables with atomic objects

- Now processes and objects are all I/O automata, combined using ordinary automata composition.
- Interactions:
  - Processes access atomic objects via invocations, get responses.
  - Invocations are outputs of processes, inputs of objects.
  - Responses are outputs of objects, inputs of processes.
  - May be a time delay between invocation and response.





# Replacing variables with atomic objects

- Locality constraints are now expressed by the I/O automata decomposition.
- More complicated than shared variables:
  - More actions (invocations/responses instead of complete accesses).
  - Algorithms have more steps, more bookkeeping.
  - More stuff to reason about.
- More realistic system model.



# Replacing variables with atomic objects

- Q: But how do we know that this replacement doesn't introduce new behaviors?
- We need some restrictions to get “equivalence”.
- Q: What can we say about failure behavior, when we make this replacement?




# Atomic Objects: Basic Definitions

# Variable Types

- **Variable type:**  $(V, v_0, \text{invs}, \text{resps}, f)$ 
  - $V$ : Set of values
  - $v_0$ : Initial value
  - $\text{invs}$ : Set of invocations
  - $\text{resps}$ : Set of responses
  - $f: \text{invs} \times V \rightarrow \text{resps} \times V$ 
    - Describes responses to an invocation and changes to the variable.
- AKA **State machine** [Lamport]
- AKA **Sequential specification** [Herlihy]
- **Execution:**  $v_0, a_1, b_1, v_1, a_2, b_2, v_2, a_3, b_3, v_3, a_4, b_4, v_4, \dots$ 
  - $v_i$  is value;  $a_i$  is invocation;  $b_i$  is response
  - Ends with value (if finite).
  - $(b_i, v_i) = f(a_i, v_{i-1})$  for  $i > 0$ .
- **Trace:**  $a_1, b_1, a_2, b_2, a_3, b_3, a_4, b_4, \dots$  (i.e., just invocations and responses, hide the variable values)

# Atomic objects

- Atomic object  $A$  of a given variable type is an I/O automaton with a particular kind of interface, satisfying some conditions:
  - Well-formedness
  - Atomicity
  - Liveness (termination)
- External interface:
  - Assume “ports”  $1, 2, \dots, n$  (one for each process).
  - May restrict so that some invocations are allowed on only some of the ports.
  - Also allow **stop** inputs on all ports, as before.
- Compose with users  $U_i$ , assumed to preserve well-formedness (alternating invocations and responses at each port, starting with an invocation).



# Conditions satisfied by A

- **Preserves well-formedness** (alternating invocations and responses at each port, starting with an invocation).
- **Atomicity:**
  - First define when a well-formed **sequence**  $\beta$  of invocations and responses (at all ports) is **atomic**.
  - Then **A satisfies atomicity** iff all well-formed executions of  $A \times U$ , where  $U = \Pi U_i$  (for any users), have atomic traces.
- First suppose that all invocations have matching responses (that is, the sequence  $\beta$  is **complete**).
- Then we say  $\beta$  is **atomic** provided that it's possible to insert a **serialization point** (dummy event) somewhere between each invocation and matching response, such that, if all the invs and resps are moved to their serialization points, the result is a trace of the (serial) variable type.





# Atomicity for complete sequences

- Suppose  $\beta$  is a **complete** well-formed sequence of invocations and responses.

Then  $\beta$  is **atomic** provided that one can insert a **serialization point** between each invocation and matching response, such that, if all the invs and resps are moved to their serialization points, the result is a trace of the (serial) variable type.

- Examples:** Initial value 0.



- read, 0, write(8), ack** is correct for serial specification.
- write(8), ack, read, 8** is also correct.

# Alternative definition [Herlihy]

- Suppose  $\beta$  is a **complete** well-formed sequence of invocations and responses. Then  $\beta$  is **atomic** provided that it can be reordered to a trace of the variable type, while preserving:
  - The order of events at each process, and
  - The order of any response and following invocation (anywhere).
- Equivalent.

# Complication: Incomplete operations

- **Q:** What about sequences  $\beta$  containing some incomplete operations? Which ops should get serialization points?
- We can't require that we **include** serialization points for **all** such operations (operation might fail right after invocation).
- We can't require that we **exclude all** such operations (operation might fail just before returning).
- So, we leave it optional...
- **Require that it's possible to:**
  - Insert serialization points for all complete operations (between invocations and responses).
  - Select **some arbitrary subset**  $\Phi$  of incomplete operations.
  - For each operation in  $\Phi$ , insert a serialization point somewhere after the invocation, and make up a response.
  - In such a way that moving all matched invocations and responses to their serialization points (and removing other invocations) yields a trace of the variable type.

# Atomic sequences, in general



- Suppose  $\beta$  is any well-formed sequence of invocations and responses.

Then  $\beta$  is **atomic** provided that one can

- Insert serialization points for all complete operations.
- Select a subset  $\Phi$  of incomplete operations.
- For each operation in  $\Phi$ , insert a serialization point somewhere after the invocation, and make up a response.
- In such a way that moving all matched invs and their resps to the serialization points (and removing other invs) yields a trace of the variable type.

# More atomicity examples


- Initial value 0.



- read, 0 is correct for serial specification.
- write(8), ack, read, 8 is correct.

# Another atomicity example


- Initial value 0.




- read, 0, read, 0,...(forever) is correct.
- The write does not (cannot) get a serialization point.

# Some non-atomic sequences

- Write not seen:



- Out-of-order reads



# Note on the atomicity property

- [Well-formedness + atomicity] is a **safety property**.
- More precisely, let  $P$  be the trace property, for sequences of invocations and responses, expressing:
  - Well-formedness for every port, plus
  - Atomicity.
- Then  $P$  is a safety property.
- In other words, if this combination doesn't hold, the violation occurs at some particular point in the sequence.
- Plausible, but not completely obvious---proved in book, p. 405.
  - Uses Konig's Lemma to show limit-closure.
  - That is, if we can assign serialization points correctly to successively-extended finite sequences, then there is a way to assign them to their infinite limiting sequence.



# Back to the conditions satisfied by an atomic object A...

- Preserves well-formedness.
- Atomicity:
  - We just defined when a well-formed **sequence**  $\beta$  of invocations and responses (at all ports) is **atomic**.
  - Then **A satisfies atomicity** iff all well-formed executions of  $A \times U$ , where  $U = \prod U_i$  (for any users) have atomic traces.
- Liveness (termination):

# Liveness

- **Failure-free termination** (basic requirement for atomic objects):
  - In any fair failure-free execution of  $A \times U$ , every invocation has a matching response.
  - “Fair” here refers to fairness in the underlying I/O automata model---A keeps taking steps.
- **Definition:** Automaton A (with the right interface) is an **atomic object** if it satisfies well-formedness, atomicity, and failure-free termination (for all U).

# Other liveness conditions

- As for consensus, we sometimes consider other liveness conditions, expressing fault-tolerance properties.
- **Wait-free termination:** In any fair execution of  $A \times U$ , every invocation on a non-failing port gets a response.
- **f-failure termination,  $0 \leq f \leq n$ :** In any fair execution of  $A \times U$  in which failures occur on  $\leq f$  ports, every invocation on a non-failing port gets a response.

# Example: A wait-free atomic object

- Variable type:
  - Natural numbers, initial value 0.
  - **read** and **increment** operations.
- Atomic object supports **read** and **increment** operations on all ports.
- Implement with an n-process shared-memory system.
- Shared read/write registers
  - $x(i)$ ,  $1 \leq i \leq n$ , natural number, initially 0.
  - $x(i)$  writable by  $i$ , readable by all.
- To implement **increment<sub>i</sub>**:
  - Process  $i$  increments its own variable  $x(i)$ .
  - Can do this using a **write** operation, by remembering the previous value written.
- To implement **read<sub>i</sub>**:
  - Process  $i$  reads all the shared variables, one at a time, in any order, and returns the sum.
- **Q:** Why does this work?

# Read/Increment algorithm

- **increment<sub>i</sub>**: Increment  $x(i)$ .
- **read<sub>i</sub>**: Read all the shared variables, one at a time, in any order, and return the sum.
- **Proof**:
  - **Well-formed, wait-free**: Immediate.
  - **Atomic**: Say where to put the serialization points.
    - For an **increment**: At the actual write step.
    - For a complete **read**:
      - Must be somewhere between invocation and response.
      - Read returns a value  $v$  such that  $\text{sum of the } x\text{'s at the beginning} \leq v \leq \text{sum at the end}$ .
      - Since the sum increases by one each time, there is some point where  $\text{sum of the } x\text{'s} = v$ .
      - Put the serialization point there.
    - For an incomplete **read**: Don't bother.
- Correctness depends on the particular kinds of operations.

# Canonical Atomic Object Automata

# Canonical atomic object automaton

- Express the set of traces acceptable for a wait-free atomic object as the fair traces of a particular **canonical object automaton**; see Section 13.1.2.
- Could generalize to f-failure termination (we'll see this later).
- Canonical object automaton keeps internal copy of the variable, plus delay buffers for invocations and responses.
- Behavior: 3 kinds of steps:
  - **Invoke**: Invocation arrives, gets put into **in-buffer**.
  - **Perform**: Invoked operation gets performed on the internal copy of the variable, response gets put into **resp-buffer**.
  - **Respond**: Response returned to user.
- Internal perform step is convenient, even though we're interested only in specifying external behavior.
- Perform step corresponds to serialization point.

# Canonical atomic object automaton

- **Liveness:**
  - One task for each port  $i$ .
  - Use the usual I/O automata convention that tasks keep getting turns to take steps.
  - To model the effects of failures, we include a specially **dummy<sub>i</sub>** action in each task  $i$ , which gets enabled when **stop<sub>i</sub>** occurs.




# Canonical atomic object automaton

- Equivalent to the original serialization-point-based specification for a wait-free atomic object, in a precise sense.
- Can be used to prove correctness of algorithms that implement atomic objects, e.g., using simulation relations to prove safety.
- **Theorem 1:** Every fair trace of the canonical automaton (with well-formed  $U$ ) satisfies the properties that define a wait-free atomic object.
- **Theorem 2:** Every trace allowed by a wait-free atomic object (with well-formed  $U$ ) is a fair trace of the canonical automaton.

# Atomic Objects vs. Shared Variables


# Atomic objects vs. shared vars

- Atomic objects aren't the same as shared variables.
- But an important basic result says we can substitute atomic objects for shared variables in a shared-memory system, and the resulting system “behaves the same”.
- Enables hierarchical construction of shared-memory systems.
- This is not completely obvious, although the research literature just assumes that it works.



# Atomic objects vs. shared vars


- Atomic objects aren't the same as shared variables.
- But an important basic result says we can substitute atomic objects for shared variables in a shared-memory system, and the resulting system “behaves the same”.
- Enables hierarchical construction of shared-memory systems.



- **The substitution:**
  - Given  $A$ , a shared-memory system, and
  - For each shared variable  $x$  of  $A$ , given an atomic object  $B_x$  (same type, interface corresponding to the allowed connections).
  - $\text{Trans}(A)$  is a composition of I/O automata, one for each process and one for each shared variable.


# Atomic objects vs. shared vars

- Given shared-memory system  $A$ , and for each shared variable  $x$  of  $A$ , given atomic object  $B_x$ .
- $\text{Trans}(A)$  is a composition of I/O automata, one for each process and one for each shared variable:
  - For **variable  $x$** , use **atomic object  $B_x$** .
  - For **process  $i$** , use **automaton  $P_i$** , where:
    - Inputs of  $P_i$  are:
      - Inputs of  $A$  on port  $i$ ,
      - Responses of all the  $B_x$ s on port  $i$ ,
      - **stop <sub>$i$</sub>** .
    - Outputs of  $P_i$  are:
      - Outputs of  $A$  on port  $i$ , and
      - Invocations to all the  $B_x$ s on port  $i$ .
    - Steps of  $P_i$  simulate those of process  $i$  of  $A$  directly, except:
      - When process  $i$  of  $A$  accesses  $x$ , then  $P_i$  invokes the operation on  $B_x$  and pauses, waiting for a response. When a response arrives,  $P_i$  resumes simulating process  $i$ .
      - When **stop <sub>$i$</sub>**  occurs, all tasks of  $P_i$  (which correspond to tasks of process  $i$  in  $A$ ) are disabled.



# Atomic objects vs. shared vars

- A note on failure actions:
  - $\text{stop}_i$  is an input both to  $P_i$ , and to all objects  $B_x$  that  $P_i$  is connected to.



# What is preserved by this transformation?

- **Theorem:** For any execution  $\alpha$  of  $\text{Trans}(A) \times U$ , there is an execution  $\alpha'$  of  $A \times U$  (that is, of the original shared-memory system) such that:
  - $\alpha \upharpoonright U = \alpha' \upharpoonright U$  (looks the same to the users), and
  - $\text{stop}_i$  events occur for the same  $i$  in  $\alpha$  and  $\alpha'$  (the same processes fail).
- **Technicality:** Need a little assumption about  $A$ : At any point, for each  $i$ , we don't have both process  $i$  and the user at  $i$  enabled to perform locally-controlled actions.
- **Proof:** Given  $\alpha$ , construct  $\alpha'$ :
  - Introduce serialization points and responses for operations of  $B_x$  in  $\alpha$ , as guaranteed by the atomicity definition.
  - Then commute the invocation and responses events with other events until they appear next to their serialization points.

# What is preserved?

- **Theorem:** For any execution  $\alpha$  of  $\text{Trans}(A) \times U$ , there is an execution  $\alpha'$  of  $A \times U$  such that:
  - $\alpha \mid U = \alpha' \mid U$  and
  - $\text{stop}_i$  events occur for the same  $i$  in  $\alpha$  and  $\alpha'$ .
- **Proof:** Given  $\alpha$ , construct  $\alpha'$ :
  - Add serialization points, responses for operations of  $B_x$ .
  - Commute invocation and responses events with other events until they appear next to their serialization points.
  - OK as far as the  $B_x$ s are concerned.
  - What about the  $P_i$ s? We aren't allowed to reorder events of the same  $P_i$ .
  - But no such reordering happens, because:
    - $P_i$  pauses when it performs invocations, and
    - No inputs arrive at  $P_i$  from  $U$  while  $P_i$  is waiting for a response to an inv (follows from the technical assumption---it's the system's turn)
  - Result is still an execution of  $\text{Trans}(A) \times U$  (using composition results), but now **all invocations and responses occurring in consecutive pairs**.
  - Now replace each pair with a single access step.



# Liveness


- Construction also preserves some liveness properties:
- We want to show that **if  $\alpha$  is fair then  $\alpha'$  is fair**, i.e., any fair execution of  $\text{Trans}(A) \times U$  emulates a fair execution of  $A \times U$ .
- **A difficulty:** Objects sometimes don't respond to invocations, whereas shared variable accesses always return. So the objects could possibly introduce new blocking.
- We need an assumption that implies that the objects don't introduce new blocking.
- E.g., we could assume that the  $B_x$  objects are wait-free.
- E.g., we could assume that at most  $f$  failures occur in  $\alpha$  and each  $B_x$  guarantees  $f$ -failure termination (**Theorem 13.7**).
  - “The failures that happen are tolerated by the objects.”
  - That's good enough to ensure that the objects always respond to non-failed processes.

# Application 1 of Trans results

- Implementing atomic objects using other atomic objects:
  - Suppose  $A$  is an atomic object implementation, using shared memory.
  - Say  $A$  and all the  $B_x$ s guarantee  $f$ -failure termination.
  - Then  $\text{Trans}(A)$  also implements an atomic object (of the same type), and guarantees  $f$ -failure termination.
  - See Corollary 13.9, p. 417, for details.


# Application 2 of Trans results

- Building shared-memory systems hierarchically.
  - Suppose the  $B_x$ s are themselves shared-memory systems.
  - Then  $\text{Trans}(A)$  yields a 2-level system:
  - If we compose each  $P_i$  at the top level with all the i-port agent processes within the  $B_x$  implementations, we get an actual shared-memory system (consisting of individual processes and variables).



# Combining the two applications


- Building shared-memory implementations of atomic objects hierarchically.
  - Like Application 2, but where:
    - Low-level systems implement atomic objects  $B_x$ , and
    - High-level system implements an atomic object, as in Application 1.
  - Shows how to combine shared-memory implementations of atomic objects at two levels to get a single shared-memory implementation of the high-level atomic object.
  - Used implicitly throughout the research literature.



# Algorithms to implement RMW atomic objects

# Read-Modify-Write Atomic Object

- Can we implement a general RMW atomic object using just read/write shared variables?
- **Non-fault-tolerant implementation:**
  - Implement the RMW variable using a single shared read/write register.
  - Access the register only within a critical region, using two operations, a read followed by a write.
  - To implement the critical region, use a lockout-free mutual exclusion algorithm, e.g., one of Peterson's.
- **Q: Fault-tolerant implementation?**



# Read-Modify-Write Atomic Object

- Fault-tolerant implementation?
- Say, 1-failure termination.
- **Theorem:** There is no shared memory system using only read/write shared variables that implements a general RMW atomic object and guarantees 1-failure termination.
- **Proof:** By contradiction.
  - Suppose there is, system B.
  - Let A be any agreement algorithm that uses shared RMW variables and guarantees 1-failure termination.
    - Earlier, we saw how to guarantee even wait-free termination.
  - Substitute B for each of the RMW shared variables in A.
  - The resulting system solves agreement in the read/write shared-memory model, with 1-failure termination.
  - Contradicts impossibility result for consensus in the read/write shared-memory model.

# Next time:

- More algorithms to implement atomic objects:
  - Atomic snapshots
  - Atomic read/write registers
- **Reading:** Sections 13.3-13.4