

6.852: Distributed Algorithms

Fall, 2015

Class 20

Today's plan

- Wait-free synchronization.
- The wait-free consensus hierarchy
- Universality of consensus
- **Reading:**
 - [Herlihy, Wait-free synchronization] (another Dijkstra Prize paper)
 - [Attiya, Welch, Chapter 15]
- **Next time:**
 - More on wait-free computability
 - Wait-free vs. f -fault-tolerant computability
 - Reading:
 - [Borowsky, Gafni, Lynch, Rajsbaum]
 - [Attiya, Welch, Section 5.3.2]
 - [Attie, Guerraoui, Kouznetsov, Lynch]

Overview

- General goal:
 - **Classify atomic object types**: Which types can be used to implement which others, for which numbers of processes and failures?
 - A theory of relative computability, for objects in distributed systems.
- Herlihy considers **wait-free termination** only ($n-1$ failures).
- Considers **particular object types**:
 - Primitives used in multiprocessor memories: test-and-set, fetch-and-add, compare-and-swap.
 - Standard programming data types: counters, queues, stacks.
 - Consensus, k -consensus.
- Defines a **hierarchy of types**, with:
 - Read/write registers at the bottom, level 1.
 - Consensus (viewed as an atomic object) at the top, level ∞ .
 - Others in between (but mostly at levels 1 and 2).
- **Universality result**: Consensus for n processes can be used to implement any object for n processes.

Herlihy's Hierarchy

- Defines hierarchy in terms of:
 - How many processes can solve consensus using only objects of the given type, plus registers (thrown in for free).
- Shows that no object type at one “level” of the hierarchy can implement any object at a higher level.
- Shows:
 - Read/write registers are at level 1.
 - Stacks, queues, fetch-and-add, test-and-set are at level 2.
 - Consensus, compare-and-swap are at “level ∞ ”.
- Hierarchy has limitations:
 - All of the interesting types are at level 1, 2 or ∞ .
 - Gives no information about relative computability of objects at the same level.
 - Lacks some basic, desirable “robustness” properties.
- Yields some interesting classification results.
- But doesn't give a complete story---more work is needed.

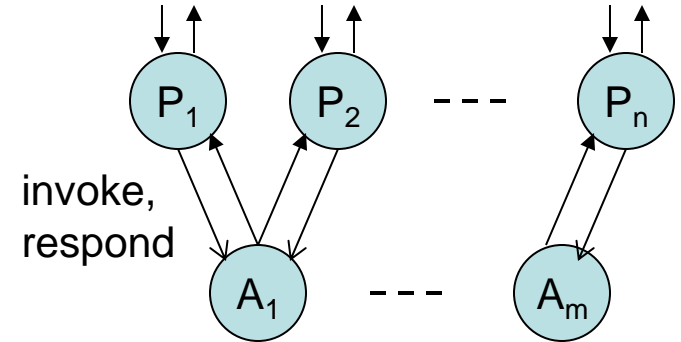
Outline

- Basic definitions
 - Consensus as an atomic object
 - Consensus numbers and the consensus hierarchy
- Queue types
 - Consensus number = 2
- Compare-and-swap types
 - Consensus number = ∞
- Universality of consensus

Basic definitions

The Model

- **Concurrent system model:**
 - Processes + atomic objects
- Herlihy models everything using I/O automata.
 - Uses a slight variant of tasks to define fair executions.
 - We'll keep using tasks.
- **Sequential specification** = variable type
- Use a concurrent system to implement an atomic object of a specified type.
- **Warning:** Herlihy's definition of implementation says only one object is used, but his results sometimes use many objects (of the same type).



Consensus as an atomic object

- **Consensus variable type** $(X, x_0, \text{invs}, \text{resps}, \delta)$:
 - V = consensus domain, $X = V \cup \{\perp\}$.
 - $x_0 = \perp$
 - $\text{invs} = \{\text{init}(v) \mid v \in V\}$
 - $\text{resps} = \{\text{decide}(v) \mid v \in V\}$
 - $\delta(\text{init}(v), \perp) = (\text{decide}(v), v)$, for any v in V
 - $\delta(\text{init}(w), v) = (\text{decide}(v), v)$, for any v, w in V
- That is, the first value anyone provides in an $\text{init}()$ operation is everyone's decision.
- Herlihy's **consensus object** is simply a wait-free atomic object for the consensus variable type.
- Enables him to consider atomic objects everywhere:
 - For high-level objects being implemented, and
 - For low-level objects used in the implementations.
- But, he generally treats low-level objects as shared variables.

Herlihy's consensus object vs. our consensus definition

- Herlihy's consensus atomic object is “almost the same” as our notion of consensus:
 - Satisfies well-formedness, agreement, strong validity (every decision is someone's initial value).
 - Wait-free termination.
 - Every `init()` on a non-failing port eventually receives a `decide()` response.
 - Doesn't add any new constraints.
- Some (unimportant) differences:
 - Allows repeated operations on the same port; but all get the same response value v .
 - Inputs needn't arrive everywhere; equivalent requirement (**Exercise 12.1**).

Binary vs. arbitrary consensus

- Herlihy's paper talks about “implementing consensus”, without specifying the domain.
- It doesn't matter:
- **Theorem:** Let \mathbf{T} be the consensus type with domain $\{0,1\}$, and \mathbf{T}' a consensus type with some other finite value domain V .

Then there is a wait-free implementation of an n -process atomic object of type \mathbf{T}' from n -process shared variables of type \mathbf{T} and read/write registers.

General consensus from binary

- **Shared variables:**
 - Binary consensus objects, $\text{Cons}(1), \dots, \text{Cons}(k)$, where k is the length of a bit string representation for elements of V .
 - Read/write registers $\text{Init}(1), \dots, \text{Init}(n)$ over $V \cup \{\perp\}$, where V is the consensus domain, initially all \perp .
- **Process i :**
 - Post (write) initial value in $\text{Init}(i)$, as a bit string.
 - Maintain current preferred value internally, initialized to initial value.
 - For $l = 1$ to k do:
 - Engage in binary consensus using $\text{Cons}(l)$, with the l -order bit of your current preference as input.
 - If your bit loses, then:
 - Read all $\text{Init}(j)$ registers to find some value whose first $l-1$ bits agree with your current preference, and whose l 'th bit is the winning bit from $\text{Cons}(l)$.
 - Reset your preference to this value.
 - Return your final preference.

What about an infinite set V ?

- **Theorem:** Let \mathbf{T} be the consensus type with domain $\{0,1\}$, \mathbf{T}' a consensus type with **any value domain V** .

Then there is a wait-free implementation of an n -process atomic object of type \mathbf{T}' from n -process shared variables of type \mathbf{T} and read/write registers.

- **Proof:**
 - Similar algorithm.
 - But now reach consensus on the **index j** for some active process, rather than the actual value (**active** means that it writes $\text{Init}(j)$).
 - Then return that j 's initial value, read from $\text{Init}(j)$.
- **Moral:** When we talk about solvability of consensus, we needn't specify V .

Consensus Numbers

- **Definition:** The **consensus number** of a variable type **T** is the largest number n such that shared variables of type **T** and read/write registers can be used to implement an n -process wait-free atomic consensus object.
- That is, **T + registers solve n -process consensus**.
- Note that read/write registers are thrown in for free.
 - Helpful in writing algorithms.
 - Reasonable because they are at the bottom of the hierarchy, consensus number 1. (Why?)
 - Follows from [Loui, Abu-Amara]: can't be used to solve even 2-process consensus.
- **Definition:** If **T + registers solve n -process consensus** for every n , then we say that **T** has **consensus number ∞** .

Consensus Numbers

- Consensus numbers yield a way of showing that one variable type \mathbf{T} cannot be used (by itself, plus registers) to implement another type \mathbf{T}' , for certain numbers of processes.
- **Theorem 1:** Suppose $\text{cons-number}(\mathbf{T}) = m$, and $\text{cons-number}(\mathbf{T}') > m$. Then there is no (wait-free) implementation of an atomic object of type \mathbf{T}' for $n > m$ processes, from shared variables of type \mathbf{T} and registers.
- **Proof:**

Consensus Numbers

- **Theorem 1:** Suppose $\text{cons-number}(\mathbf{T}) = m$, and $\text{cons-number}(\mathbf{T}') > m$. Then there is no (wait-free) implementation of an atomic object of type \mathbf{T}' for $n > m$ processes, from shared variables of type \mathbf{T} and registers.
- **Proof:**
 - Enough to show for $n = m+1$.
 - By contradiction. Suppose there is an $(m+1)$ -process implementation of an atomic object of type \mathbf{T}' from objects of type \mathbf{T} + registers.
 - Since $\text{cons-number}(\mathbf{T}') > m$, there is an $(m+1)$ -process consensus algorithm C using objects of type \mathbf{T}' + registers.
 - Replace the \mathbf{T}' shared variables in C with the assumed implementation of \mathbf{T}' objects from \mathbf{T} + registers.
 - By our composition theorem for shared-memory algorithms, this yields an $(m+1)$ -process consensus algorithm using \mathbf{T} + registers.
 - Contradicts assumption that $\text{cons-number}(\mathbf{T}) = m$.

Example: Read/write register types

- **Theorem 2:** Any read/write register type, for any value domain V and any initial value v_0 , has consensus number 1.
- **Proof:**
 - Clearly, can be used to solve 1-process consensus (trivial).
 - Cannot solve 2-process consensus [book, Theorem 12.6].
- **Corollary 3:** Suppose $\text{cons-number}(\mathbf{T}') > 1$. Then there is no (wait-free) implementation of an atomic object of type \mathbf{T}' for $n > 1$ processes, from registers only.
- **Proof:**
 - By Theorems 1 and 2.

Example: Snapshot types

- **Corollary 3:** Suppose $\text{cons-number}(\mathbf{T}') > 1$. Then there is no (wait-free) implementation of an atomic object of type \mathbf{T}' for $n > 1$ processes, from registers only.
- **Theorem 4:** Any snapshot type, for any underlying domain (W, w_0) , has consensus number 1.
- **Proof:**
 - By contradiction.
 - Suppose there is a snapshot type \mathbf{T}' with $\text{cons-number}(\mathbf{T}') > 1$.
 - Thus, it can be used to solve 2-process consensus.
 - Then by Corollary 3, there is no wait-free implementation of an atomic object of type \mathbf{T}' for > 1 processes, from registers only.
 - Contradicts known implementation of snapshots from registers.

Queue Types

Queue types

- FIFO queue type $\text{queue}(V, q_0)$:
 - V is some value domain.
 - q_0 is a finite sequence giving the initial queue contents.
 - Operations:
 - $\text{enqueue}(v)$, v in V : Add v to end of queue, return ack.
 - $\text{dequeue}()$: Return head of queue if nonempty, else \perp .
- Most commonly: $q_0 = \lambda$, empty sequence.
- **Theorem 5:** There is a queue type \mathbf{T} with $\text{cons-number}(\mathbf{T}) \geq 2$.
- **Proof:**

Queue types

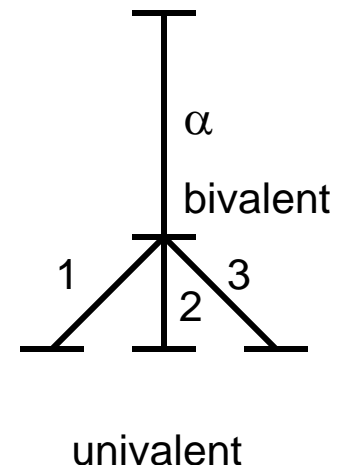
- **Theorem 5:** There is a queue type \mathbf{T} with $\text{cons-number}(\mathbf{T}) \geq 2$.
- **Proof:**
 - Construct a 2-process consensus algorithm for an arbitrary domain V , using queue shared variables.
 - Shared variables:
 - One queue of integers, initially = sequence consisting of one element, 0.
 - Registers $\text{Init}(1)$ and $\text{Init}(2)$ over $X = V \cup \{\perp\}$, initially \perp .
 - **Process i :**
 - Post initial value in $\text{Init}(i)$.
 - Perform **enqueue()**.
 - If you get 0, then return your initial value.
 - Else (you get \perp), read and return $\text{Init}(j)$, for the other process j .
 - First dequeuer wins.

Queue types

- **Theorem 5:** There is a queue type T with $\text{cons-number}(T) \geq 2$.
- **Corollary 6:** There is no wait-free implementation of an n -process atomic object of the above queue type using registers only, for any $n \geq 2$.
- **Proof:**
 - By Corollary 3.
 - Essentially: Suppose there is. Plug it into the above 2-process consensus algorithm and get a 2-process consensus algorithm using registers only, contradiction.
- **Q:** What about queues with other initial values q_0 ?
- **E.g.,** initially-empty queues?
 - Claim there's an algorithm, but it's more complicated. HW?
- **Q:** What about other, known initial values?

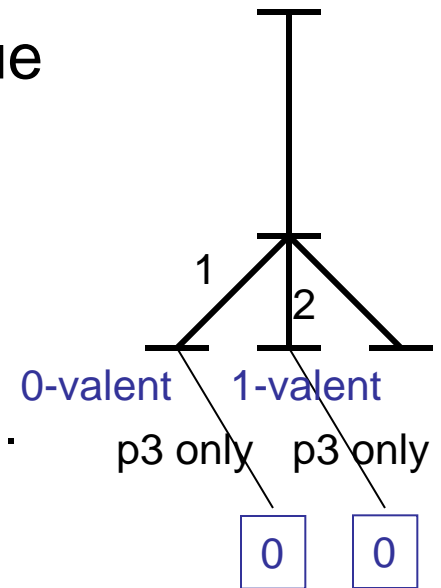
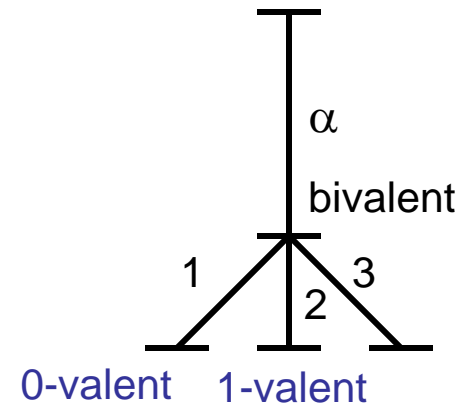
Queue lower bound

- **Theorem 7:** Every queue type \mathbf{T} has $\text{cons-number}(\mathbf{T}) \leq 2$.
- **More strongly:** No combination of queue variables, with any queue types, initialized in any way, plus registers, can implement 3-process consensus.
- **Proof:**
 - Suppose such an algorithm, A , exists.
 - As for the register-only case, we can show that A has a bivalent initialization.
 - Furthermore, we can maneuver as before to a decider configuration:



Queue impossibility

- Suppose WLOG that process 1 yields 0-valence, process 2 yields 1-valence.
- Consider what p1 and p2 might do in their steps.
- If they access different variables, or both access the same register, we get contradictions as in the pure read/write case.
- So assume they both access the same queue q; consider cases based on types of operations they perform.
- **Case 1: p1 and p2 both dequeue:**
 - Then resulting states look the same to p3.
 - Running p3 alone after both yields a contradiction.



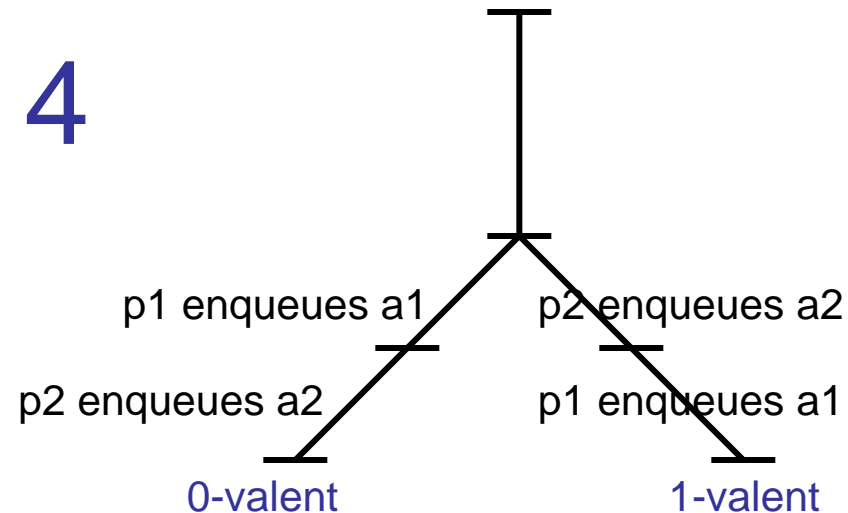
Case 2

- Case 2: p1 enqueues and p2 dequeues:
 - If the queue is nonempty after α , the two steps commute---same system state after p1 p2 or p2 p1, yielding a contradiction.
 - If the queue is empty after α , then the states after p1 and p2 p1 look the same to all but p2 (and the queue is the same in both cases).
 - Running p3 (or p1) alone after both yields a contradiction.
- Case 3: p1 dequeues and p2 enqueues:
 - Symmetric.

Case 4

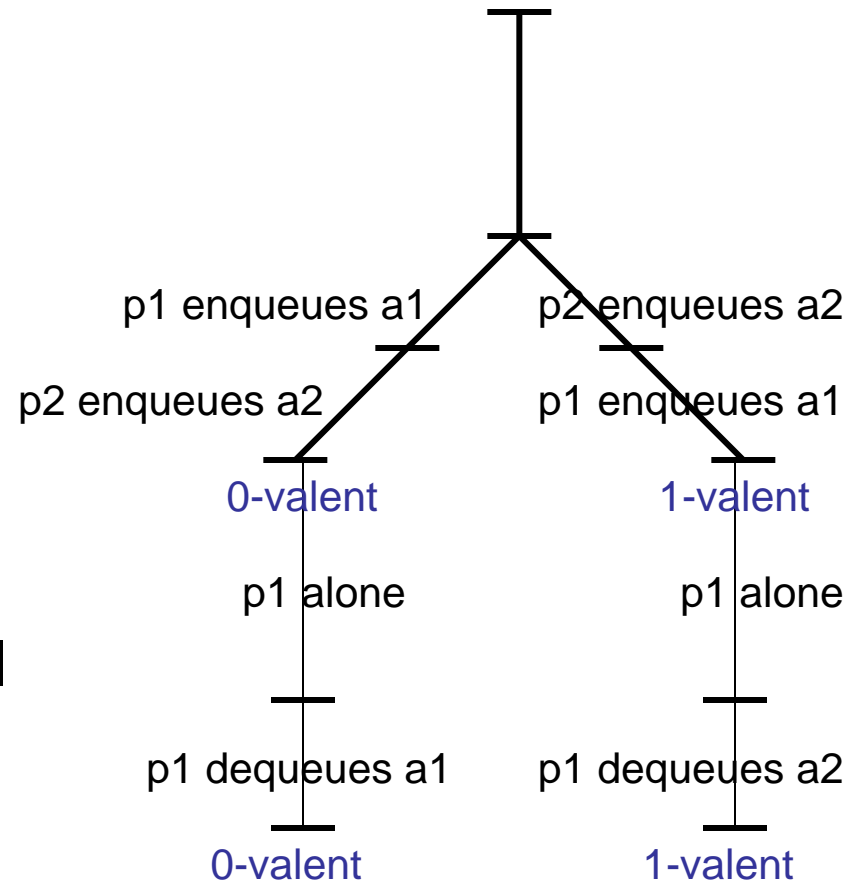
- Case 4: p1 and p2 both enqueue:

- Consider two possible orders:
- We will construct two executions:
 - After p1 p2, p1 runs alone until it dequeues a1, then p2 runs alone until it dequeues a2.
 - After p2 p1, p1 runs alone until it dequeues a2, then p2 runs alone until it dequeues a1.
- These two executions are indistinguishable by p3, leading to the usual sort of contradiction.
- But how do we construct these two executions?
 - Q: What is different after p1 p2 and p2 p1?
 - Only the queue q, which ends with a1 a2 in first case, a2 a1 in second.
 - States of all processes, values of other objects, are the same in both.



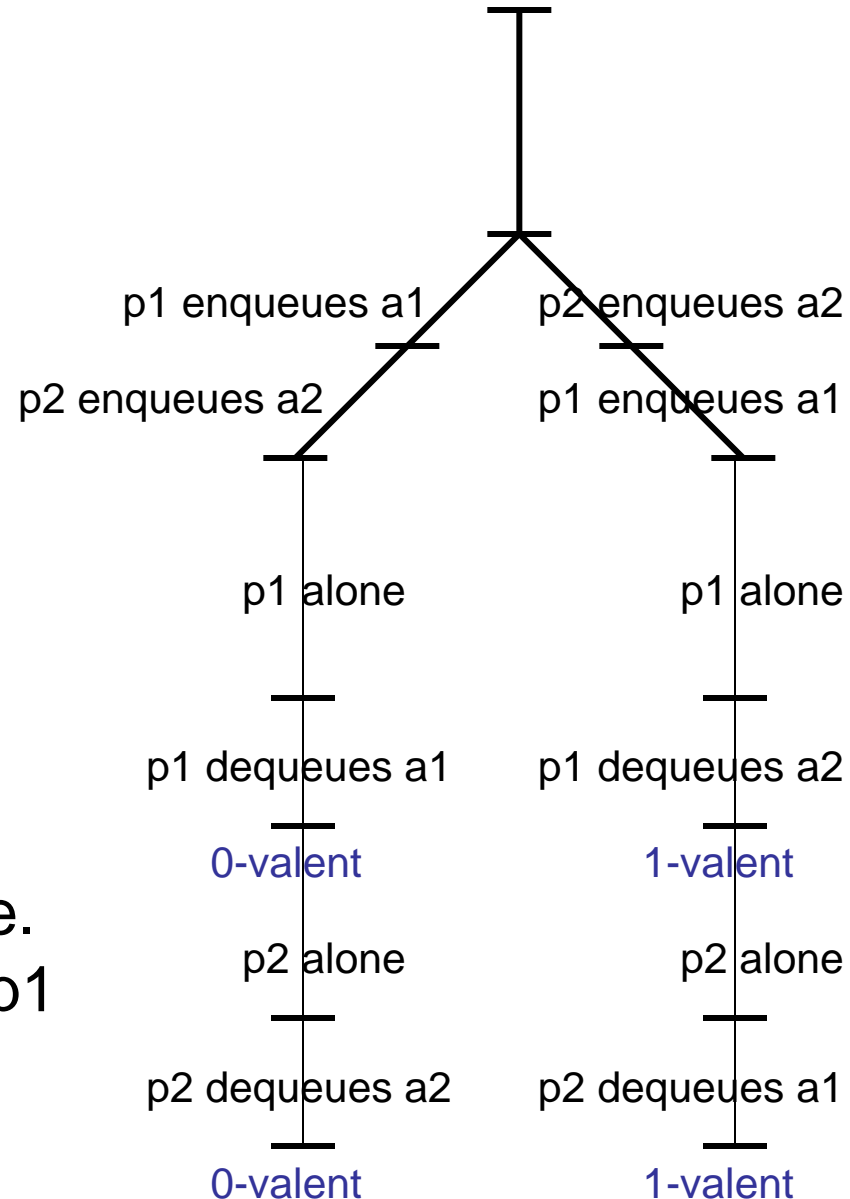
Constructing the executions

- Run p1 alone after p1 p2 and after p2 p1.
- Must eventually decide, differently in these two situations.
- But p1 can't distinguish until it dequeues a1 or a2 from q, so it must eventually do so.
- So we can run p1 alone just until it dequeues a1 or a2 from q.
- **Q:** Now what is different?
- q starts with a2 on left branch, a1 on right branch, later elements are the same.
- States of all other objects are the same.
- States of p2 and p3 are the same, but p1 may be different.



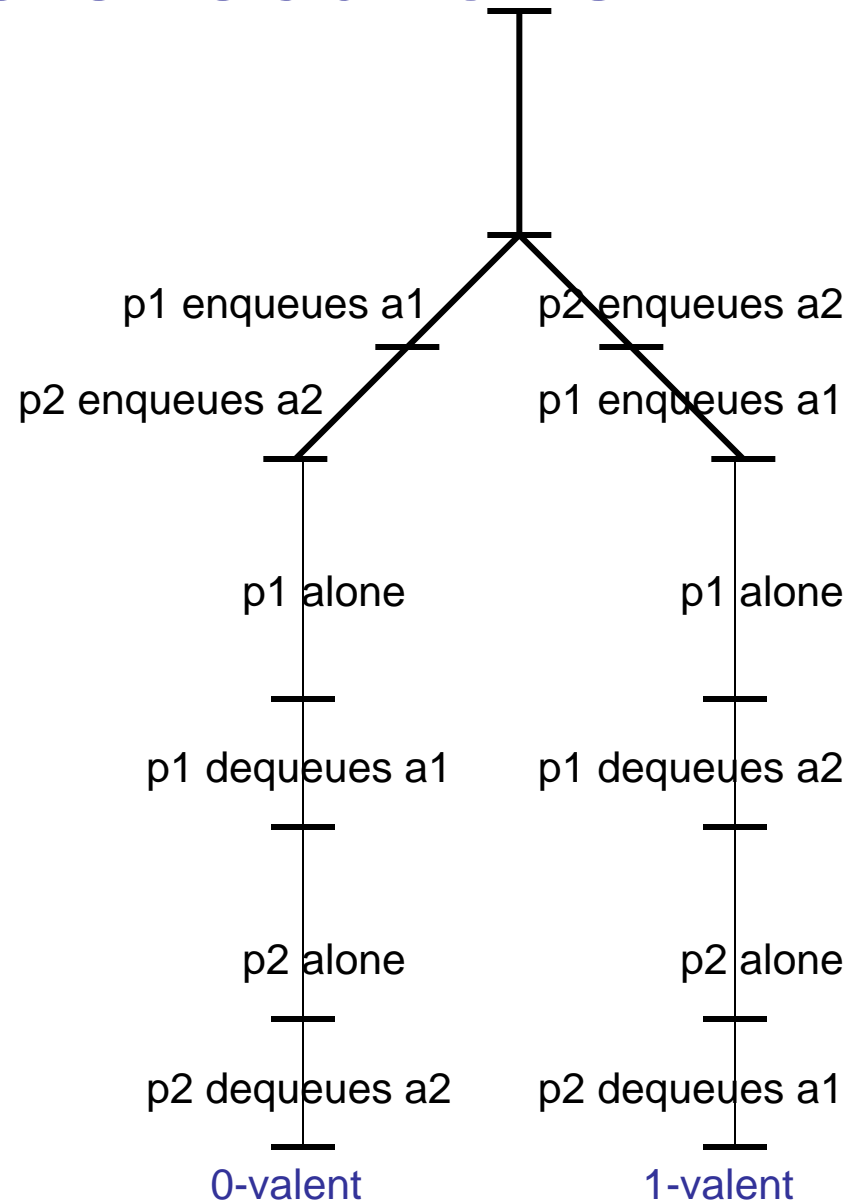
Constructing the executions

- Now run p2 alone after both branches.
- Must decide differently in the two executions.
- But p2 can't distinguish until it dequeues from q, so it must eventually do so.
- So run p2 alone just until it dequeues from q.
- **Q:** Now what is different?
- All objects, including q, are same.
- State of p3 is the same, though p1 and p2 may be different.



Constructing the executions

- This gives the needed executions:
 - After p1 p2, p1 runs alone until it dequeues a1, then p2 runs alone until it dequeues a2.
 - After p2 p1, p1 runs alone until it dequeues a2, then p2 runs alone until it dequeues a1.
- As described earlier, just run p3 alone after both to get the contradiction.



Queue types: Recap

- We just showed:
 - **Theorem 7:** Every queue type \mathbf{T} has $\text{cons-number}(\mathbf{T}) \leq 2$.
 - In fact (stronger statement), all queue types together can't solve 3-process consensus.
 - So $\text{cons-number}(\mathbf{T})$ definition doesn't tell the entire story.
- Also:
 - **Theorem 5:** There is a queue type \mathbf{T} with $\text{cons-number}(\mathbf{T}) \geq 2$.
- Gives quite a bit of information about the power of queue types.

Compare-and-Swap (CAS) Types

Compare-and-swap types

- Compare-and-swap type:
 - V , the value domain.
 - v_0 , initial value.
 - $\text{invs} = \{ \text{compare-and-swap}(u,v) \mid u, v \text{ in } V \}$
 - $\text{resps} = V$
 - $\delta(\text{compare-and-swap}(u,v), w) =$
 - (w, v) if $u = w$,
 - (w, w) if not.
- That is, if the variable value is equal to the first argument, change it to the second argument; otherwise leave the variable alone.
- In either case, return the former value of the variable.

Compare-and-swap types

- **Theorem 8:** Let \mathbf{T} be the consensus type with value domain V . Then there is a compare-and-swap type \mathbf{T}' that can be used to implement an n -process consensus object with type \mathbf{T} , for any n .
- That is, \mathbf{T}' can be used to solve n -process consensus for any n ; so $\text{cons-number}(\mathbf{T}') = \infty$.
- **Proof:**
 - Use just a single C&S shared variable, value domain = $V \cup \{\perp\}$, initial value = \perp .
 - Process i :
 - If initial value = v , then access the C&S shared variable with **compare-and-swap**(\perp , v), obtain the previous value w .
 - If $w = \perp$ then decide v . (You are first).
 - Otherwise, decide w . (Someone else was first and proposed w .)

Compare-and-swap types

- **Corollary 9:** It is impossible to implement an atomic object of this C&S type \mathbf{T}' (from Theorem 8) for $n \geq 3$ processes using just queues and read/write registers.
- **Proof:** Like proof of Theorem 1.
 - Enough to show for $n = 3$.
 - By contradiction. Suppose there is a 3-process implementation of an atomic object of type \mathbf{T}' using queues + registers.
 - By Theorem 8, there is a 3-process consensus algorithm C using just \mathbf{T}' + registers.
 - Replace the \mathbf{T}' shared variables in C with the assumed implementation of \mathbf{T}' from queues + registers.
 - Yields a 3-process consensus algorithm using just queues + registers.
 - Contradicts (the stronger version of) Theorem 7.
- **[Herlihy]** classifies other data types similarly, LTTR.

Universality of Consensus

Universality of consensus

- Consensus variables and registers can implement a wait-free n -process atomic object of any variable type, for any number n .
- Algorithm in [Herlihy] combines:
 - A basic unfair, non-wait-free algorithm.
 - A fairness mechanism, to ensure that every operation is completed.
 - Optimizations, to reuse memory, save time.
- [Attiya, Welch, Chapter 15] separate these three aspects.
- Here, we'll simplify by ignoring the optimizations.
- Assume arbitrary data type $T = (V, v_0, \text{invs}, \text{resps}, \delta)$.
- Fix n .

1. Non-wait-free algorithm

- **Shared variables:**
 - An infinite sequence of n-process consensus variables, $\text{Cons}(1)$, $\text{Cons}(2)$, ...
 - Each consensus variable's domain is $\{ (j, k, a) \text{ where: } \begin{array}{l} \bullet j \text{ is a process id, } 1 \leq j \leq n, \\ \bullet k \text{ is a positive integer, a local sequence number,} \\ \bullet a \in \text{invs, the set of invocations for the } T \text{ object } \end{array} \}$
- $\text{Cons}(r)$ is used to decide which proposed invocation on the implemented object is the r^{th} one to be performed.
- The consensus objects explicitly decide on the sequence of invocations, and it's consistently observed everywhere.
- **Process i:**
 - Participates in consensus executions in order 1,2,3,...
 - Keeps track locally of the decision values for all consensus variables; these are triples of the form (j,k,a) .
 - Knowing the sequence of consensus decisions allows process i to “run” the invocations in the sequence and compute the new states and responses for the implemented object.

Non-wait-free algorithm, process i

- When new invocation **a** arrives:
 - Record it in local variable **current-inv**, as a triple **(i, k, a)**, where k is the first unused local sequence number.
 - For each **Cons(r)**, starting from the first one that i hasn't yet participated in:
 - Invoke **init(current-inv)** on **Cons(r)**.
 - Record decision in local variable **decision(r)**.
 - If **decision(r) = current-inv** then
 - Run the sequence of invocations in **decision(1), ..., decision(r)** to compute the response.
 - Return response to the user (and become idle).
 - Else continue on to r+1.

Algorithm properties

- **Well-formed:** Yes
- **Atomic:** Yes
 - Everyone sees a consistent sequence of operations.
 - Serialization point for an operation can be the point where it wins at some consensus shared variable **Cons(r)**.
- **Wait-free:** No
 - Process *i* could submit the same operation to infinitely many **Cons** variables, and it could always lose.

2. Wait-free algorithm

- Add a simple **priority mechanism** to ensure that each operation completes.
- For **Cons(r)**, $r \equiv i \bmod n$, any current invocation of process i gets priority.
- Priority is managed outside the consensus variables:
 - A process i sometimes “helps” another process j , by invoking consensus objects with j ’s invocation instead of i ’s own.
- **Additional shared variables:**
 - **announce(i)**, for each process i , a single-writer multi-reader register, written by i , read by everyone
 - Value domain: $\{ (i, k, a) \text{ as above} \} \cup \{ \perp \}$.
 - Initial value: \perp

Wait-free algorithm, process i

- When new invocation **a** arrives:
 - Record it in local variable **current-inv** as before, as triple (i, k, a).
 - Write value of **current-inv** into **announce(i)**.
 - Then proceed as in the non-wait-free algorithm, except:
 - Before participating in **Cons(r)**, read **announce(j')**, where $j' \equiv r \pmod n$.
 - If **announce(j')** contains a triple **inv** (not \perp), and **inv** has not already won any of **Cons(1)**, **Cons(2)**, ..., **Cons(r-1)**, then invoke **init(inv)** on **Cons(r)**.
 - Otherwise, invoke **init(current-inv)** on **Cons(r)**.
 - Handle decisions as before.
 - Just before returning a response to the user, reset **announce(i)** $:= \perp$.

Algorithm properties

- **Well-formed, Atomic:** Yes, as before.
- **Wait-free:** Yes:
 - Claim every operation eventually completes.
 - If not, then consider some (i,k,a) that gets stuck.
 - Then after `announce(i)` is set to (i,k,a) , it keeps this value forever.
 - Process i participates in infinitely many consensus executions on behalf of this (i,k,a) , losing all of them.
 - Choose any r such that:
 - $r \equiv i \pmod n$, and
 - r is sufficiently large so that no process accesses `Cons(r)`, or even reads `announce(i)` in preparation for accessing `Cons(r)`, before `announce(i)` is set to (i,k,a) .
 - Then for this j , everyone who participates will choose to help i by submitting (i,k,a) as input.
 - At least one process participates (i itself).
 - So the decision must be (i,k,a) .

Complexity

- **Shared-memory size:**
 - Infinitely many shared variables, each of unbounded size.
- **Time:**
 - Unbounded, because:
 - A process i may start with a $\text{Cons}(r)$ that is far out of date, and have to access $\text{Cons}(r)$, $\text{Cons}(r+1)$, ... to catch up.
- **Herlihy:**
 - Formulates the algorithm somewhat differently, in terms of a linked list of operations, so it's hard to compare.
 - Time:
 - Claims a nice $O(n)$ bound.
 - Avoids the catch-up time by allowing processes to survey others to get recent information.
 - Shared memory:
 - Still uses unbounded sequence numbers.
 - Still uses infinitely many consensus objects---seems unavoidable since each is good for only one decision.
 - "Garbage-collects" to reclaim space taken by old objects.

Robustness

- [Jayanti] defined a robustness property for the hierarchy:
 - **Robustness:** If **T** is a type at level n , and **S** is a set of types, all at levels $< n$, then **T** has no implementation from **S** for n processes.
- But he did not determine whether the hierarchy is robust.
- Herlihy's results don't imply this; what they say is:
 - If **T** is a type at level n , and **S** is a single type at a level $< n$, then **T** has no implementation from **S** and registers.
- But it's still conceivable that combining low-consensus-number types could allow implementation of a higher-consensus-number type.
- Later papers give both positive and negative results.
 - Based on technical issues.

Summary

- Basic definitions for wait-free consensus and the consensus hierarchy.
- Read/write register types, snapshot types, are at level 1.
- Queue data types are at level 2.
- Compare-and-Swap data types are at level ∞ .
- Universality of consensus.
- Work is still needed to achieve our original goals:
 - Determine which types of objects can be used to implement which other types, for which numbers of processes and failures.
 - A comprehensive theory of relative computability, for objects in distributed systems.

Next time...

- More on wait-free computability
- Wait-free vs. f -fault-tolerant computability
- **Reading:**
 - [Borowsky, Gafni, Lynch, Rajsbaum]
 - [Attiya, Welch, Section 5.3.2]
 - [Attie, Guerraoui, Kouznetsov, Lynch]