

6.852: Distributed Algorithms

Fall, 2015

Lecture 6

Today's plan

- Fault-tolerant synchronous distributed algorithms
- Fault-tolerant consensus
- Link failures:
 - The Two Generals problem
- Process(or) failures:
 - Stopping and Byzantine failure models
 - Algorithms for agreement with stopping and Byzantine failures
 - Exponential Information Gathering
- **Reading:** Sections 5.1, 6.1-6.3
- **Next:**
 - Lower bounds for Byzantine agreement:
 - Number of processors
 - Number of rounds

Distributed consensus

- Abstract problem of reaching agreement among processes in a distributed system, all of which start with their own “opinions”.
- Complications: Failures (process, link); timing uncertainties.
- **Motivation:**
 - Database transactions: Commit or abort
 - Aircraft control:
 - Agree on value of altimeter reading (SIFT)
 - Agree on which plane should go up/down, in resolving encounters (TCAS)
 - Resource allocation: Agree on who gets priority for obtaining a resource, doing the next database update, etc.
 - Replicated State Machines: To emulate a virtual machine consistently, agree on next step.


Distributed consensus

- Abstract problem of reaching agreement among processes in a distributed system, all of which start with their own “opinions”.
- Complications: Failures (process, link); timing uncertainties.
- Fundamental problem
- We'll revisit it several times:
 - In synchronous, asynchronous, and partially synchronous settings.
 - In insect colony algorithms.
 - With link failures, processor failures.
 - Algorithms, impossibility results.

Consensus with Link Failures

Informal Scenario

- Several generals plan a coordinated attack.
- All should agree to attack:
 - Absolutely must agree.
 - Should attack if possible.
- Each has an initial opinion about his/her army's readiness.
- Nearby generals can communicate using foot messengers:
 - Unreliable, can get lost or captured
 - Connected, undirected communication graph, known to all generals
 - Known bound on time for successful messenger to deliver message.
- Motivation: Transaction commit
- Can show no algorithm exists!




Formal problem statement

- $G = (V, E)$, undirected graph (bidirectional edges)
- Synchronous model, n processes
- Each process has input 1 (attack) or 0 (don't attack).
- Any subset of the messages can be lost.
- All should eventually set **decision** output variables to 0 or 1.
 - In practice, would need this to happen by some deadline.
- Correctness conditions:
 - **Agreement:**
 - No two processes decide differently.
 - **(Weak) Validity:**
 - If all start with 0, then 0 is the only allowed decision.
 - If all start with 1 and all messages are successfully delivered, then 1 is the only allowed decision.

Alternatively:

- Stronger validity condition:
 - If anyone starts with 0 then 0 is the only allowed decision.
 - If all start with 1 and all messages are successfully delivered, then 1 is the only allowed decision.
 - Typical for transaction commit (1 = commit, 0 = abort).
- Guidelines:
 - For designing algorithms, try to use stronger correctness conditions (better algorithm).
 - For impossibility results, use weaker conditions (better impossibility result).


Impossibility for 2 Generals [Gray]



- Other cases similar, LTTR.
- **Proof:** By contradiction.
 - Suppose we have a solution---a process (states, transitions) for each index 1, 2.
 - Assume that both processes send messages at every round.
 - WLOG, could add dummy messages.
 - Proof is based on limitations of local knowledge.
 - Start with α , the execution where both start with input 1 and all messages are received.
 - By the termination condition, both eventually decide.
 - Say, by the end of r rounds.
 - By the validity condition, both decide on 1.

2-Generals Impossibility

- α_1 : Same as α , but lose all messages after round r .
 - Doesn't matter, since both processes have already decided by round r .
 - So, both decide 1 in α_1 .
- α_2 : Same as α_1 , but lose the last message from process 1 to process 2.
 - Claim α_1 is indistinguishable from α_2 by process 1, $\alpha_1 \sim^1 \alpha_2$.
 - Formally, 1 sees the same sequence of states, incoming and outgoing messages.
 - So process 1 also decides 1 in α_2 .
 - By termination, process 2 decides in α_2 .
 - By agreement, process 2 decides 1 in α_2 .




A fine point:

- In α_2 , process 2 must decide 1 at some point, not necessarily by round r .
- That's good enough...

Continuing...

- α_3 : Same as α_2 , but lose the last message from process 2 to process 1.
 - Then $\alpha_2 \sim^2 \alpha_3$.
 - So process 2 decides 1 in α_3 .
 - By termination, process 1 decides 1 in α_3 .
 - By agreement, process 1 decides 1 in α_3 .
- α_4 : Same as α_3 , but lose the last message from process 1 to process 2.
 - Then $\alpha_3 \sim^1 \alpha_4$.
 - So process 1 decides 1 in α_4 .
 - So process 2 decides 1 in α_4 .
- Keep removing edges, get to:



The contradiction

- α_{2r+1} : Both start with 1, no messages received.
 - Still both must eventually decide 1.
- α_{2r+2} : process 1 starts with 1, process 2 starts with 0, no messages received.
 - Then $\alpha_{2r+1} \sim^1 \alpha_{2r+2}$.
 - So process 1 decides 1 in α_{2r+2} .
 - So process 2 decides 1 in α_{2r+2} .
- α_{2r+3} : Both start with 0, no messages received.
 - Then $\alpha_{2r+2} \sim^2 \alpha_{2r+3}$.
 - So process 2 decides 1 in α_{2r+3} .
 - So process 1 decides 1 in α_{2r+3} .
- But α_{2r+3} contradicts (weak) validity!

Impossibility Result

- We have proved:
- **Theorem:** There is no algorithm to solve the coordinated attack problem for a 2-node graph.
- So what can we do?
 - Use randomized algorithms, get probabilistic guarantees.
 - E.g., see Section 5.2 [Varghese].

Consensus with Process(or) Failures

Consensus with process failures

- Stopping failures (crashes) and Byzantine failures (arbitrary processor malfunction, possibly malicious)
- Agreement problem:
 - n -node connected, undirected graph, known to all processes.
 - Input v from a set V , in a special state variable.
 - Output v from V , by setting $\text{decision} := v$.
 - Bounded number $\leq f$ of processors may fail.
- Why a bounded number of failures?
 - A typical way of describing limited amounts of failure.
 - Alternatives: Bounded rate of failure; probabilistic bounds on failure.

Stopping agreement

- Assume process may stop at any point:
 - Between rounds.
 - While sending messages at a round; any subset of intended messages may be delivered.
 - After sending, but before changing state.
- Correctness conditions:
 - **Agreement:** No two processes (failing or not) decide on different values.
 - “Uniform agreement”
 - **Validity:** If all processes start with the same v , then v is the only allowable decision.
 - **Termination:** All nonfaulty processes eventually decide.
- Alternatively:
 - **Stronger validity condition:** Every decision value must be some process' initial value.
 - Use this later, for k -agreement.

Byzantine agreement

- “Byzantine Generals Problem” [Lamport, Pease, Shostak]
 - Originally “Albanian Generals”
- Faulty processes may exhibit “arbitrary behavior”:
 - Can start in arbitrary states, send arbitrary messages, perform arbitrary transitions.
 - But can’t affect anyone else’s state or outgoing messages.
 - Often called “malicious” (but they need not be).
- Correctness conditions:
 - **Agreement:** No two **nonfaulty** processes decide on different values.
 - **Validity:** If all **nonfaulty** processes start with the same v , then v is the only allowable decision **for nonfaulty processes**.
 - **Termination:** All nonfaulty processes eventually decide.

Technicality about stopping vs. Byzantine agreement

- A Byzantine agreement algorithm doesn't necessarily solve stopping agreement:
- For stopping, all processes that decide, even ones that later fail, must agree (uniformity condition).
- Too strong for Byzantine setting.
- Implication holds in some special cases, e.g., when decisions always happen at the end.


Complexity measures

- **Time:** Number of rounds until all nonfaulty processes decide.
- **Communication:** Number of messages, or number of bits.
 - For Byzantine case, just count those sent by nonfaulty processes.

Consensus with Process(or) Stopping Failures

Simple algorithm for stopping agreement

- Assume complete n-node graph.
- Idea:
 - Processes keep sending all V values they've ever seen.
 - Use simple decision rule at the end.
- Specifically:
 - Process i maintains $W \subseteq V$, initially containing just i 's initial value.
 - Repeatedly: Broadcast W , add received elements to W .
 - After k rounds:
 - If $|W| = 1$ then decide on the unique value.
 - Else decide on a default value $v_0 \in V$.
- Q: What should k be?



How many rounds?

- Depends on number f of failures to be tolerated.
- $f = 0$:
 - $k = 1$ works.
 - All get the same W .
- $f = 1$:
 - $k = 1$ doesn't work:
 - Say process 1 has initial value u , others have initial value v .
 - Process 1 fails during round 1, sends to some and not others.
 - So some have $W = \{v\}$, others $\{u,v\}$, may decide differently.
 - $k = 2$ works:
 - If someone fails in round 1, then no one fails in round 2.
- General f :
 - $k = f + 1$

Correctness proof (for $k = f+1$)

- **Claim 1:** Suppose $1 \leq r \leq f+1$ and no process fails during round r . Let i and j be two processes that haven't failed by the end of round r . Then $W_i = W_j$ right after round r .
- **Proof:** Each gets exactly the union of all the W 's of the processes that have not failed by the beginning of round r .
- "Clean round"---allows everyone to resolve their differences.
- **Claim 2:** Suppose all the W sets are identical just after round r , for all processes that are still non-failed. Then the same is true for any $r' > r$.
- **Proof:** Obvious.

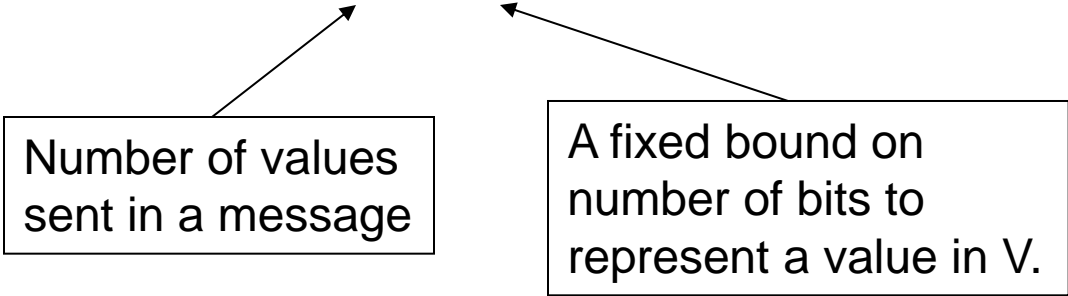
Check correctness conditions

- **Agreement:**
 - There must be some round r , $1 \leq r \leq f+1$, at which no process fails (since $\leq f$ failures)---a clean round.
 - Claim 1 says all that haven't yet failed have same W after round r .
 - Claim 2 implies that all have same W after round $f + 1$.
 - So nonfaulty processes pick the same value.
- **Validity:**
 - If everyone starts with v , then v is the only value that anyone ever gets, so $|W| = 1$ and v is chosen.
- **Termination:**
 - Obvious from decision rule.

Complexity bounds

- **Time:** $f+1$ rounds
- **Communication:**
 - Messages: $\leq (f + 1) n^2$
 - Message bits: Multiply by $n b$

Number of values
sent in a message



The diagram consists of two rectangular boxes. The left box contains the text 'Number of values sent in a message'. The right box contains the text 'A fixed bound on number of bits to represent a value in V.'. From the top-right corner of the left box, an arrow points diagonally upwards and to the right, ending at the 'n' in the 'Message bits' term of the communication list. From the top-left corner of the right box, an arrow points diagonally upwards and to the left, ending at the 'b' in the 'Message bits' term.

A fixed bound on
number of bits to
represent a value in V .

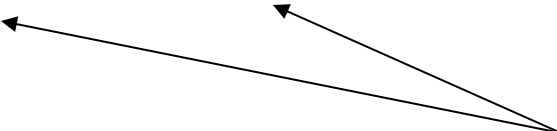
- **Can improve communication:**
 - Messages: $\leq 2 n^2$
 - Message bits: Multiply by b

Improved algorithm (Opt)

- Each process broadcasts its own value in round 1.
- May broadcast at **one other round**, just after it first learns about some value different from its own.
- In that case, it chooses just **one such value** to rebroadcast.
- After $f + 1$ rounds, use same rule as before:
 - If $|W| = 1$ then decide on the unique value.
 - Else decide on default value v_0 .

Correctness

- Relate behavior of Opt to that of the original algorithm.
- Specifically, relate executions of both algorithms with the same inputs and same failure pattern.
- Let OW denote the W set in the optimized algorithm.
- Relation between states of the two algorithms:
 - For every i :
 - $OW_i \subseteq W_i$.
 - If $|W_i| = 1$ then $OW_i = W_i$.
 - If $|W_i| > 1$ then $|OW_i| > 1$.



Not necessarily the same set,
but both > 1 .


- Relation after $f+1$ rounds implies same decisions.

Proof of correspondence

- Induction on number of rounds (p. 107)
- Key ideas:
 - $OW_i \subseteq W_i$
 - Obvious, since Opt just suppresses sending of some messages from Unopt.
 - If $|W_i| = 1$ then $OW_i = W_i$.
 - Nothing suppressed in this case.
 - Actually, follows from the first property and the fact that OW_i is always nonempty.
 - If $|W_i| > 1$ then $|OW_i| > 1$.
 - Inductive step, for some round r :
 - If in Unopt, process i receives messages only from processes with $|W| = 1$, then in Opt, it receives the same sets. So after round r , $OW_i = W_i$. So in this case, if $|W_i| > 1$ then $|OW_i| > 1$.
 - Otherwise, in Unopt, process i receives a message from some process j with $|W_j| > 1$, and so (by induction), $|OW_j| > 1$. Then after round r , $|W_i| > 1$ and $|OW_i| > 1$.

Exponential Information Gathering (EIG)

- A strategy for consensus algorithms, which works for Byzantine agreement as well as stopping agreement.
- Based on EIG tree data structure.
- EIG tree $T_{n,f}$, for n processes, f failures:
 - $f+2$ levels
 - Paths from root to leaf correspond to strings of $f+1$ distinct process names.
- Example: $T_{4,2}$




EIG Stopping agreement algorithm


- Each process i uses the same EIG tree, $T_{n,f}$.
- Decorates nodes of the tree with values in V , level by level.
- **Initially:** Decorate root with i 's input value.
- **Round $r \geq 1$:**
 - Send all level $r-1$ decorations for nodes whose labels don't include i , to everyone.
 - Including yourself---simulate locally.
 - Use received messages to decorate level r nodes---to determine label, append sender's id at the end.
 - If no message is received, use \perp .
- **The decoration for node $(i_1, i_2, i_3, \dots, i_k)$ in i 's tree is the value v such that $(i_k$ told $i)$ that $(i_{k-1}$ told $i_k)$ that ...that $(i_1$ told $i_2)$ that i_1 's initial value was v .**
- **Decision rule for stopping case (trivial):**
 - Let W = set of all values decorating the local EIG tree.
 - If $|W| = 1$ decide that value, else default v_0 .

Example


- 3 processes, 1 failure
- Use $T_{3,1}$:




Initial values:



Process 1





Process 2




Process 3

Example


- Process 2 is faulty, fails after sending to process 1 at round 1.
- After round 1:



Process 1





Process 2




Process 3

Example


- After round 2:



Process 1



Process 2



Process 3

p3 discovers that p2's value is 0 after round 2, by hearing it from p1.

Correctness and complexity

- Correctness similar to previous algorithms.
- Time: $f+1$ rounds, as before.
- Messages: $\leq (f + 1) n^2$
- Bits: Exponential in number of failures, $O(n^{f+1} b)$
- Can improve as before by relaying only the first two messages with distinct values.
- **Extension:**
 - The simple EIG stopping algorithm, and its optimized variant, can be used to tolerate worse types of failures.
 - Not full Byzantine model---that will require more work...
 - Rather, a restricted version of the Byzantine model, in which processes can **authenticate messages**.
 - Removes ability of process to make false claims about what other processes said.

Consensus with Byzantine Failures

Byzantine agreement algorithm


- Recall correctness conditions:
 - **Agreement:** No two **nonfaulty** processes decide on different values.
 - **Validity:** If all **nonfaulty** processes start with the same v , then v is the only allowable decision **for nonfaulty processes**.
 - **Termination:** All **nonfaulty** processes eventually decide.
- Present EIG algorithm for Byzantine agreement, using:
 - Exponential communication (in f)
 - $f+1$ rounds
 - $n > 3f$
- Expensive!
 - Time bound: Inherent. (Lower bound, next time)
 - Number-of-processors bound: Inherent. (Lower bound, next time)
 - Communication: Can be improved to polynomial.

Bad example: $n = 3, f = 1$

- Consider three executions of an EIG algorithm, with any decision rule.
- α_1 : p1 and p2 nonfaulty, initial value 1, p3 faulty, initial value 0
 - Round 1: All truthful
 - Round 2: p3 lies, telling p1 that “p2 said 0”; all other communications are truthful.
 - Validity requires that p1 and p2 decide 1.
- α_2 : p2 and p3 nonfaulty, initial value 0, p1 faulty, initial value 1
 - Round 1: All truthful
 - Round 2: p1 lies, telling p3 that “p2 said 1”; all other communications are truthful.
 - Validity requires that p2 and p3 decide 0.
- α_3 : p1 nonfaulty, initial value 1, p3 nonfaulty, initial value 0, p2 faulty, initial value doesn't matter.
 - Round 1: p2 tells p1 its initial value is 1, tells p3 its initial value is 0 (inconsistent).
 - Round 2: All truthful.
- $\alpha_3 \sim^1 \alpha_1$, so p1 behaves the same in both, decides 1 in α_3 .
- $\alpha_3 \sim^3 \alpha_2$, so p3 behaves the same in both, decides 0 in α_3 .
- Contradicts agreement!


Bad example

- α_1 : p1 and p2 nonfaulty, initial value 1, p3 faulty, initial value 0
 - Round 1: All truthful
 - Round 2: p3 lies, telling p1 that “p2 said 0”; all other communications are truthful.
 - Validity requires that p1 and p2 decide 1.




Bad example

- α_2 : p2 and p3 nonfaulty, initial value 0, p1 faulty, initial value 1
 - Round 1: All truthful
 - Round 2: p1 lies, telling p3 that “p2 said 1”; all other communications are truthful.
 - Validity requires that p2 and p3 decide 0.



Bad example

- α_3 : p1 nonfaulty, initial value 1, p3 nonfaulty, initial value 0, p2 faulty, initial value doesn't matter.
 - Round 1: p2 tells p1 its initial value is 1, tells p3 its initial value is 0 (inconsistent).
 - Round 2: All truthful.



Notes on the example


- The correct processes can tell something is wrong, but that doesn't help:
 - E.g., in α_1 , p1 sees that p2 sends 1, but p3 says that p2 said 0.
 - So p1 knows that either p2 or p3 is faulty, but doesn't know which.
 - By termination, p1 has to decide something, but neither value works right in all cases.
- Impossibility of solving Byzantine agreement with 3 processes, 1 failure:
 - This is not a proof--- maybe there's a non-EIG algorithm, or one that takes more rounds,...
 - Come back to this next time...

EIG algorithm for Byzantine agreement

- Assume $n > 3f$.
- Same EIG tree as before.
- Relay messages for $f+1$ rounds, as before.
- Decorate the tree with values from V , replacing any garbage messages with default value v_0 .
- Call the decorations $\text{val}(x)$, where x is any node label.
- **New decision rule:**
 - **Redecorate** the tree bottom-up, defining $\text{newval}(x)$.
 - Leaf: $\text{newval}(x) = \text{val}(x)$
 - Non-leaf: $\text{newval}(x) =$
 - newval of strict majority of children in the tree, if majority exists,
 - v_0 otherwise.
 - Final decision: $\text{newval}(\lambda)$ (newval at root)

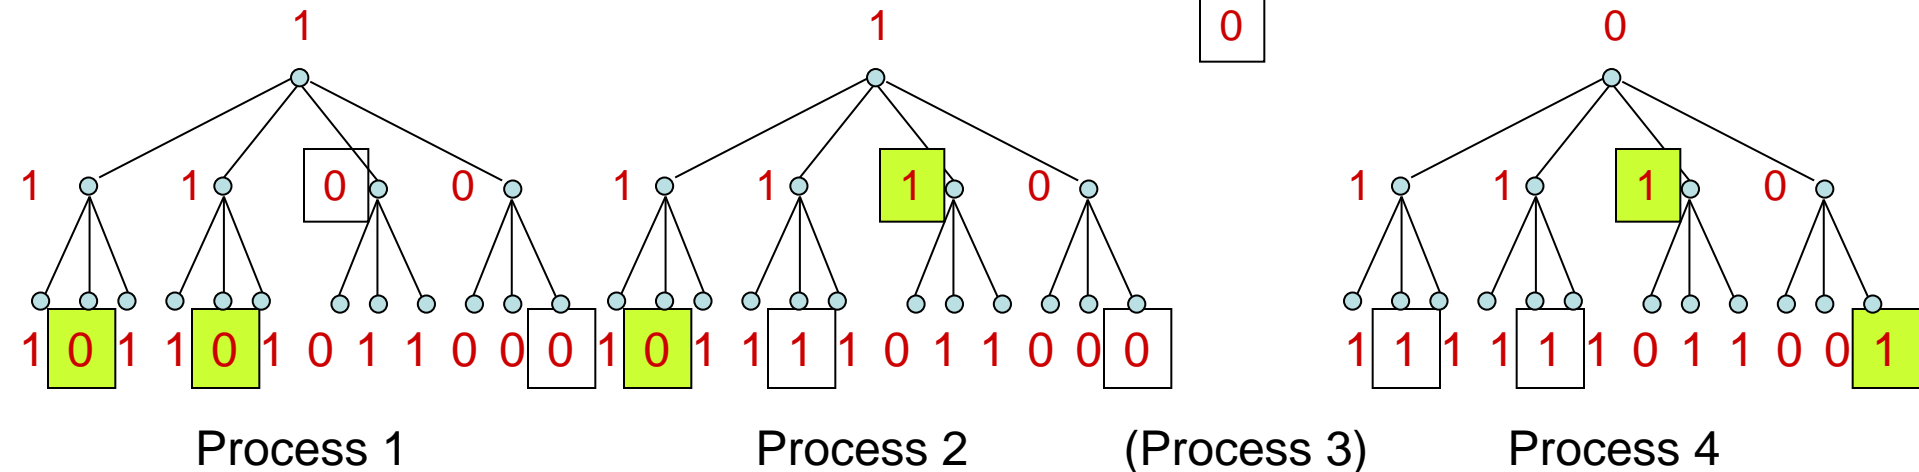
Example: $n = 4, f = 1$

- $T_{4,1}$:
- Consider a possible execution in which p3 is faulty.
- Initial values 1 1 0 0
- Round 1
- Round 2



Lies



0





Example: $n = 4, f = 1$

- Now calculate newvals, bottom-up, choosing majority values, $v_0 = 0$ if no majority.

Corrected by taking majority



0



Process 1


Process 2

(Process 3)

Process 4

Correctness proof

- **Lemma 1:** If i, j, k are nonfaulty, then $\text{val}(x)_i = \text{val}(x)_j$ for every node label x ending with k .
- In example, such nodes are (in red):



- **Proof:** k sends same message to i and j and they decorate accordingly.

Proof, cont'd


- **Lemma 2:** If x ends with nonfaulty process index then $\exists v \in V$ such that $\text{val}(x)_i = \text{newval}(x)_i = v$ for every nonfaulty i .
- **Proof:** Induction on lengths of labels, bottom up.
 - **Basis:** Leaf.
 - Lemma 1 implies that all nonfaulty processes have same $\text{val}(x)$.
 - $\text{newval} = \text{val}$ for each leaf.
 - **Inductive step:** $|x| = r \leq f$ ($|x| = f+1$ at leaves)
 - Lemma 1 implies that all nonfaulty processes have same $\text{val}(x)$, say v .
 - We need $\text{newval}(x) = v$ everywhere also.
 - Every nonfaulty process j broadcasts same v for x at round $r+1$, so $\text{val}(xj)_i = v$ for every nonfaulty j and i .
 - By inductive hypothesis, also $\text{newval}(xj)_i = v$ for every nonfaulty j and i .
 - A majority of labels of x 's children end with nonfaulty process indices:
 - Number of children of node x is $\geq n - f > 3f - f = 2f$.
 - At most f are faulty.
 - So, majority rule applied by i leads to $\text{newval}(x)_i = v$, for all nonfaulty i .

Main correctness conditions

- **Validity:**
 - If all nonfaulty processes begin with v , then all nonfaulty processes broadcast v at round 1, so $\text{val}(j)_i = v$ for all nonfaulty i, j .
 - By Lemma 2, also $\text{newval}(j)_i = v$ for all nonfaulty i, j .
 - Majority rule implies $\text{newval}(\lambda)_i = v$ for all nonfaulty i .
 - So all nonfaulty i decide v .
- **Termination:**
 - Obvious.
- **Agreement:**
 - Requires a bit more work:

Agreement


- **Path covering:** Subset of nodes containing at least one node on each path from root to leaf:



- **Common node:** One for which all nonfaulty processes have the same newval.
 - If a node's label ends in a nonfaulty process index, Lemma 2 implies it's common.
 - Others might be common too.

Agreement


- **Lemma 3:** There exists a path covering all of whose nodes are common.
- **Proof:**
 - Let C = nodes with labels of the form x_i , i nonfaulty.
 - By Lemma 2, all of these are common.
 - Claim these form a path covering:
 - There are at most f faulty processes.
 - Each path contains $f+1$ labels ending with $f+1$ distinct indices.
 - So at least one of these labels ends with a nonfaulty process index.



Agreement

- **Lemma 4:** If there's a common path covering of the subtree rooted at any node x , then x is common.
- **Proof:**
 - By induction, from the leaves up.
 - “Common-ness” propagates upward.


- **Example:**



Agreement

- **Lemma 4:** If there's a common path covering of the subtree rooted at any node x , then x is common
- **Proof:**
 - By induction, from the leaves up.
 - “Common-ness” propagates upward.


- **Example:**



Agreement

- **Lemma 4:** If there's a common path covering of the subtree rooted at any node x , then x is common
- **Proof:**
 - By induction, from the leaves up.
 - “Common-ness” propagates upward.

- **Example:**



Agreement

- **Lemma 3:** There exists a path covering all of whose nodes are common.
- **Lemma 4:** If there's a common path covering of the subtree rooted at any node x , then x is common
- **Lemma 5:** The root is common.
- **Proof:** By Lemmas 3 and 4.
- Thus, all nonfaulty processes get the same $\text{newval}(\lambda)$.
- Yields Agreement.

Complexity bounds

- As for EIG for stopping agreement:
 - Time: $f+1$
 - Communication: $O(n^{f+1})$
- Number of processes: $n > 3f$
- Q: Is $n > 3f$ necessary?

Next time...

- Lower bounds for Byzantine agreement:
 - Number of processors
 - Bounds for connectivity, weak Byzantine agreement.
 - Number of rounds
- Reading:
 - Sections 6.4-6.7
 - [Aguilera, Toueg]
 - (Optional) [Keidar-Rajsbaum]