

# 6.852: Distributed Algorithms

## Fall, 2015

### Lecture 7

# Today's plan

- Exponential Information Gathering (EIG) algorithm for Byzantine agreement.
- Lower bounds:
  - Number-of-processors lower bound for Byzantine agreement.
  - Connectivity bounds.
  - Weak Byzantine agreement.
  - Time lower bounds for stopping and Byzantine agreement.
- Reading:
  - Sections 6.3-6.7
  - [Aguilera, Toueg]
  - [Keidar, Rajsbaum]
- Next: Some other distributed agreement problems

# Byzantine agreement


- Recall correctness conditions:
  - **Agreement:** No two nonfaulty processes decide on different values.
  - **Validity:** If all nonfaulty processes start with the same  $v$ , then  $v$  is the only allowable decision for nonfaulty processes.
  - **Termination:** All nonfaulty processes eventually decide.
- EIG algorithm for Byzantine agreement, using:
  - Exponential communication (in  $f$ )
  - $f+1$  rounds
  - $n > 3f$

# EIG algorithm for Byzantine agreement

- Assume  $n > 3f$ .
- Same EIG tree as before.
- Relay messages for  $f+1$  rounds, as before.
- Decorate the tree with values from  $V$ , replacing any garbage messages with default value  $v_0$ .
- Call the decorations  $\text{val}(x)$ , where  $x$  is any node label.
- **New decision rule:**
  - **Redecorate** the tree bottom-up, defining  $\text{newval}(x)$ .
    - Leaf:  $\text{newval}(x) = \text{val}(x)$
    - Non-leaf:  $\text{newval}(x) =$ 
      - $\text{newval}$  of strict majority of children in the tree, if majority exists,
      - $v_0$  otherwise.
  - Final decision:  $\text{newval}(\lambda)$  ( $\text{newval}$  at root)

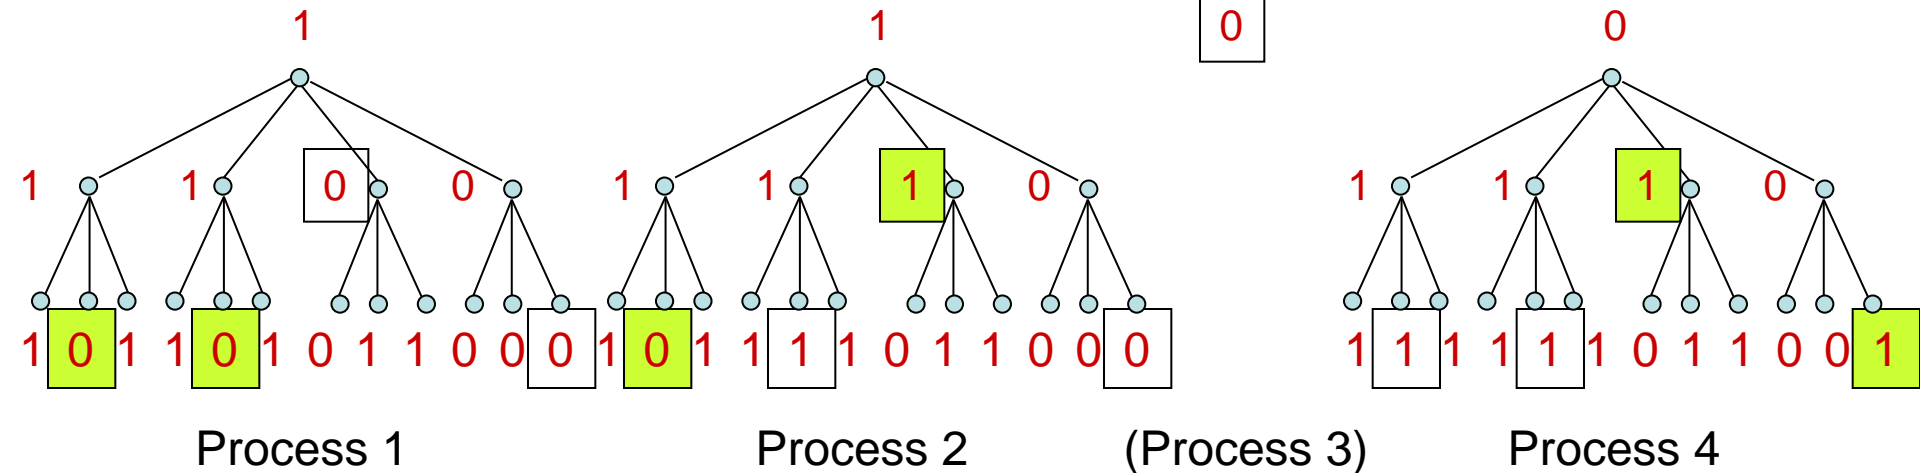
# Example: $n = 4, f = 1$

- $T_{4,1}$ :
- Consider a possible execution in which p3 is faulty.
- Initial values 1 1 0 0
- Round 1
- Round 2



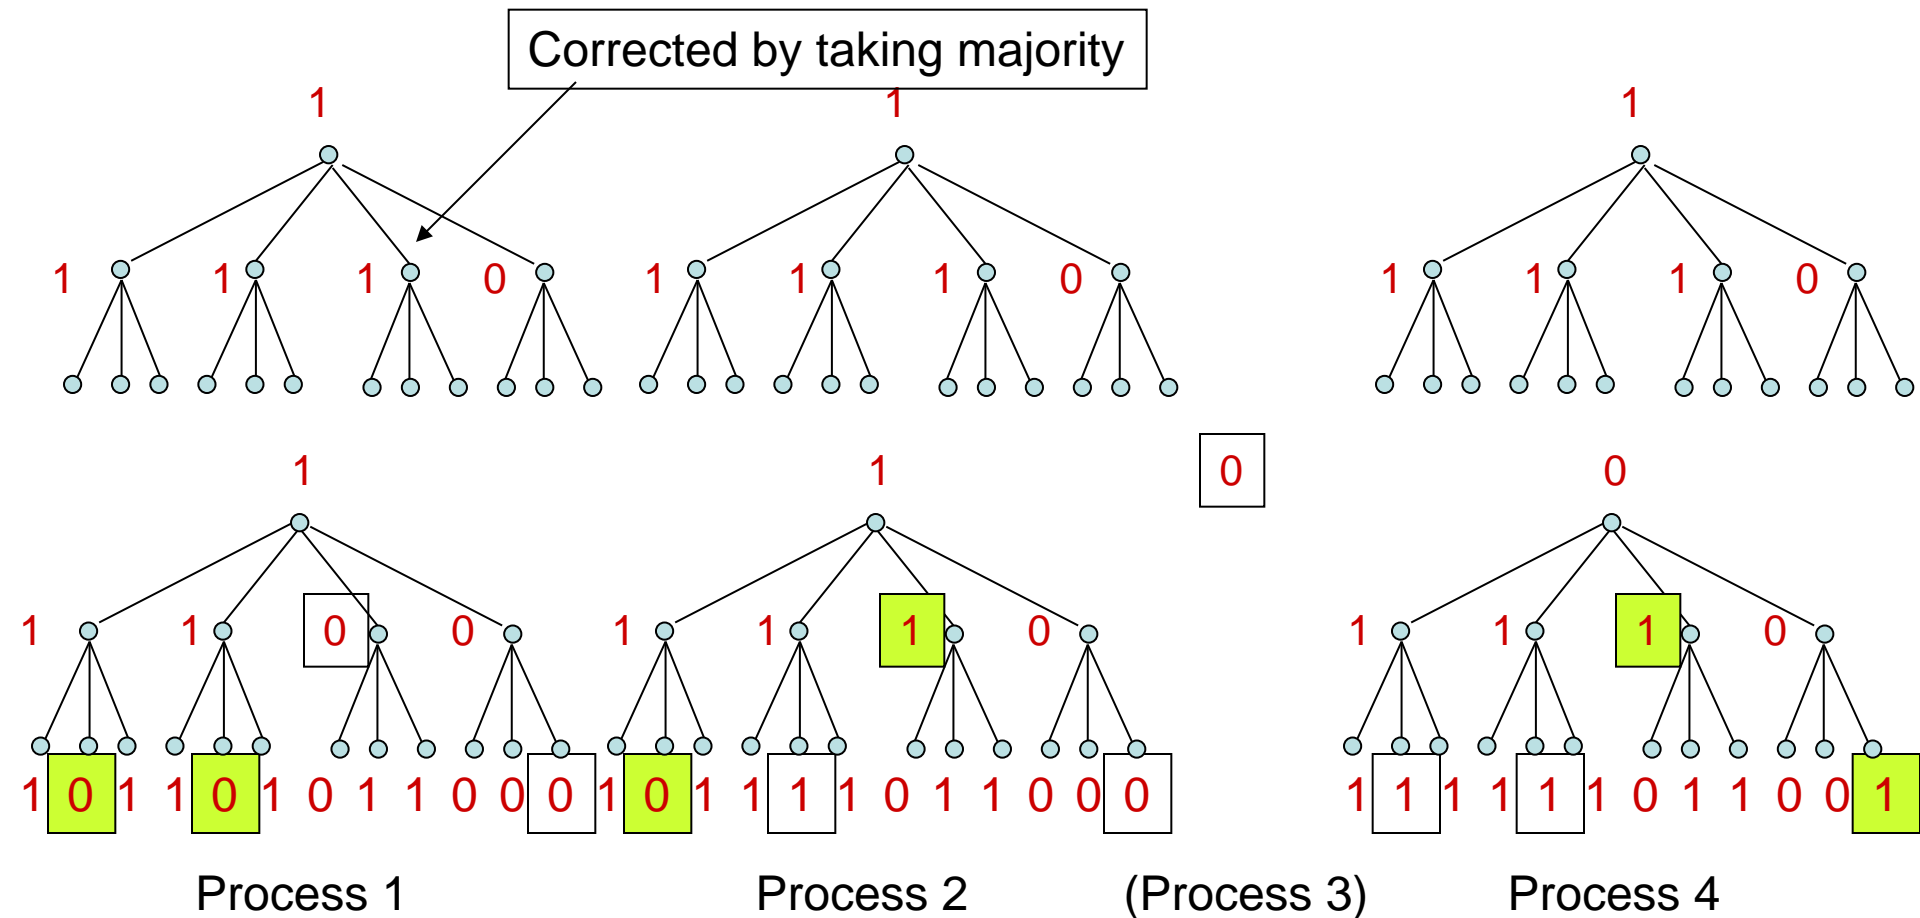
Lies

0




# Example: $n = 4, f = 1$

- Now calculate newvals, bottom-up, choosing majority values,  $v_0 = 0$  if no majority.



# Correctness proof

- **Lemma 1:** If  $i, j, k$  are nonfaulty, then  $\text{val}(x)_i = \text{val}(x)_j$  for every node label  $x$  ending with  $k$ .
- In example, such nodes are (in red):



- **Proof:**  $k$  sends same message to  $i$  and  $j$  and they decorate accordingly.

# Proof, cont'd

- **Lemma 2:** If  $x$  ends with a nonfaulty process index then  $\exists v \in V$  such that  $\text{val}(x)_i = \text{newval}(x)_i = v$  for every nonfaulty  $i$ .
- **Proof:** Induction on lengths of labels, bottom up.
  - **Basis:** Leaf.
    - Lemma 1 implies that all nonfaulty processes have same  $\text{val}(x)$ .
    - $\text{newval} = \text{val}$  for each leaf.
  - **Inductive step:**  $|x| = r \leq f$  ( $|x| = f+1$  at leaves)
    - Lemma 1 implies that all nonfaulty processes have same  $\text{val}(x)$ , say  $v$ .
    - We need  $\text{newval}(x) = v$  everywhere also.
    - Every nonfaulty process  $j$  broadcasts same  $v$  for  $x$  at round  $r+1$ , so  $\text{val}(xj)_i = v$  for every nonfaulty  $j$  and  $i$ .
    - By inductive hypothesis, also  $\text{newval}(xj)_i = v$  for every nonfaulty  $j$  and  $i$ .
    - A majority of labels of  $x$ 's children end with nonfaulty process indices:
      - Number of children of node  $x$  is  $\geq n - f > 3f - f = 2f$ .
      - At most  $f$  are faulty.
    - So, majority rule applied by  $i$  leads to  $\text{newval}(x)_i = v$ , for all nonfaulty  $i$ .




# Main correctness conditions

- **Validity:**
  - If all nonfaulty processes begin with  $v$ , then all nonfaulty processes broadcast  $v$  at round 1, so  $\text{val}(j)_i = v$  for all nonfaulty  $i, j$ .
  - By Lemma 2, also  $\text{newval}(j)_i = v$  for all nonfaulty  $i, j$ .
  - Majority rule implies  $\text{newval}(\lambda)_i = v$  for all nonfaulty  $i$ .
  - So all nonfaulty  $i$  decide  $v$ .
- **Termination:**
  - Obvious.
- **Agreement:**
  - Requires a bit more work:

# Agreement


- **Path covering:** Subset of nodes containing at least one node on each path from root to leaf:



- **Common node:** One for which all nonfaulty processes have the same newval.
- If a node's label ends in a nonfaulty process index, Lemma 2 implies it's common.
- Others might be common too.

# Agreement


- **Lemma 3:** There exists a path covering all of whose nodes are common.
- **Proof:**
  - Let  $C$  = nodes whose labels end with a nonfaulty process index.
  - By Lemma 2, every node in  $C$  is common.
  - Claim  $C$  is a path covering:
    - There are at most  $f$  faulty processes.
    - Each path contains  $f+1$  labels ending with  $f+1$  distinct indices.
    - So at least one of these labels ends with a nonfaulty process index.



# Agreement

- **Lemma 4:** If there's a common path covering of the subtree rooted at any node  $x$ , then  $x$  is common.
- **Proof:**
  - By induction, from the leaves up.
  - “Common-ness” propagates upward.


- **Example:**



# Agreement

- **Lemma 4:** If there's a common path covering of the subtree rooted at any node  $x$ , then  $x$  is common
- **Proof:**
  - By induction, from the leaves up.
  - “Common-ness” propagates upward.


- **Example:**



# Agreement

- **Lemma 4:** If there's a common path covering of the subtree rooted at any node  $x$ , then  $x$  is common
- **Proof:**
  - By induction, from the leaves up.
  - “Common-ness” propagates upward.

- **Example:**



# Agreement

- **Lemma 3:** There exists a path covering all of whose nodes are common.
- **Lemma 4:** If there's a common path covering of the subtree rooted at any node  $x$ , then  $x$  is common.
- **Lemma 5:** The root is common.
- **Proof:** By Lemmas 3 and 4.
- Thus, all nonfaulty processes get the same  $\text{newval}(\lambda)$ .
- Yields Agreement.

# Complexity bounds

- As for EIG for stopping agreement:
  - Time:  $f+1$
  - Communication:  $O(n^{f+1})$
- Number of processes:  $n > 3f$
- Q: Is  $n > 3f$  necessary?




# Lower bound on the number of processes for Byzantine Agreement

# Number of processes for Byzantine agreement

- $n > 3f$  is necessary!
  - Holds for any  $n$ -node (undirected) graph.
  - For graphs with low connectivity, may need even more processes.
  - Number of failures that can be tolerated for Byzantine agreement in an undirected graph  $G$  has been completely characterized, in terms of number of nodes and graph connectivity.
- **Theorem 1:** 3 processes cannot solve Byzantine Agreement with 1 possible failure.

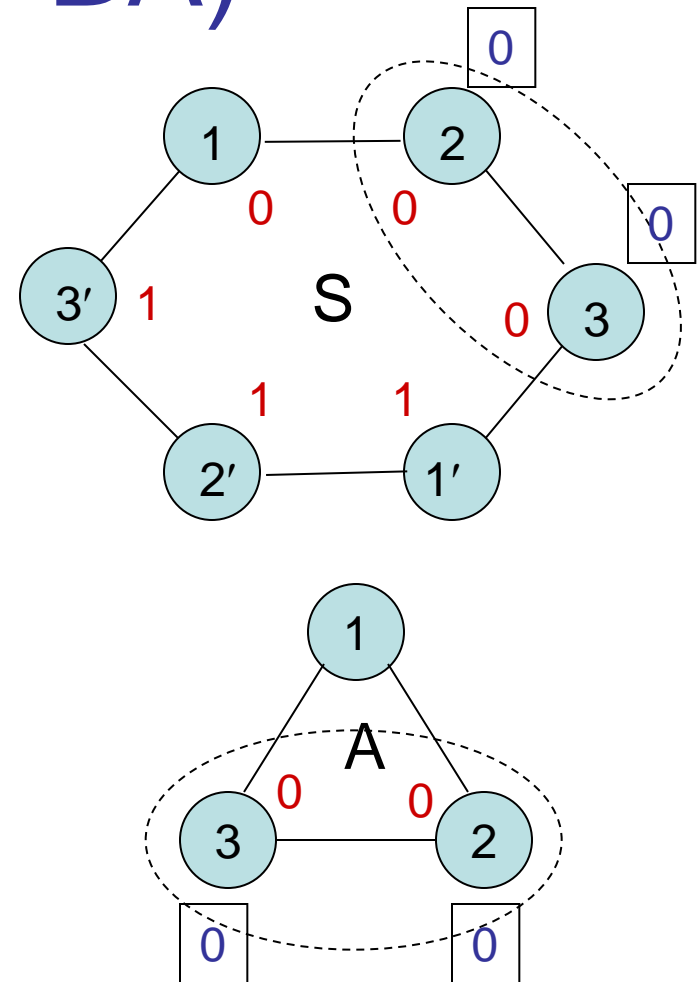
# Proof (3 vs. 1 BA)

- By contradiction. Suppose algorithm A, consisting of processes 1, 2, 3, solves BA with 1 possible failure.
- Construct new system S from 2 copies of A, with initial values as follows:
- **What is S?**
  - A synchronous system of some kind.
  - Not required to satisfy any particular correctness conditions.
  - Not necessarily a correct BA algorithm for the 6-node ring.
  - Just some synchronous system, which runs and does something.
  - We'll use it to get our contradiction.




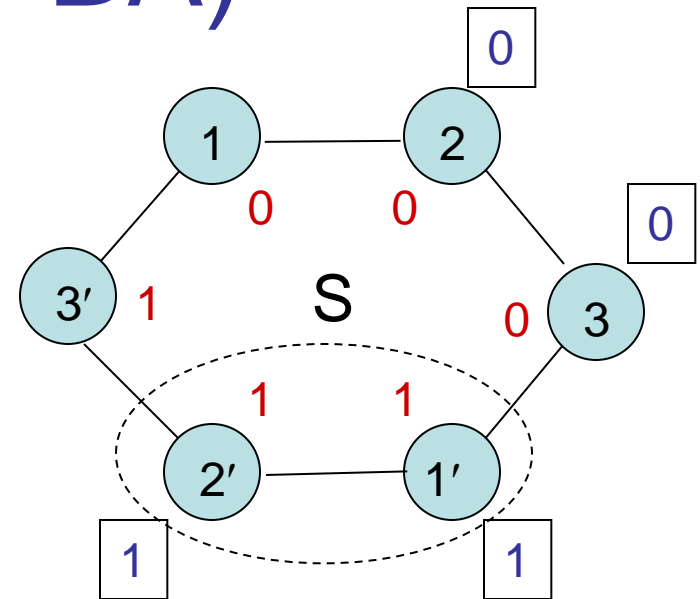
# Proof (3 vs 1 BA)

- Consider 2 and 3 in S:
- Looks to them like:
  - They're in A, with a faulty process 1.
  - 1 emulates 1'-2'-3'-1 from S.
- In A, 2 and 3 must decide 0
- So by indistinguishability, they decide 0 in S also.





# Proof (3 vs 1 BA)

- Now consider  $1'$  and  $2'$  in  $S$ .
- Looks to them like:
  - They're in  $A$  with a faulty process 3.
  - 3 emulates  $3'-1-2-3$  from  $S$ .
- They must decide 1 in  $A$ , so they decide 1 in  $S$  also.



# Proof (3 vs 1 BA)

- Finally, consider 3 and 1' in S:
  - Looks to them like:
    - They're in A, with a faulty process 2.
    - 2 emulates 2'-3'-1-2 from S.
  - In A, 3 and 1 must agree.
  - So by indistinguishability, 3 and 1' agree in S also.
- 
- But we already know that process 1' decides 1 and process 3 decides 0, in S.
  - Contradiction!



# Discussion

- We get this contradiction even if the original algorithm  $A$  is assumed to “know  $n$ ”.
- That simply means that:
  - The processes in  $A$  have the number 3 hard-wired into their state.
  - Their correctness properties are required to hold only when they are actually configured into a triangle.
- We are allowed to use these processes in a different configuration  $S$ ---as long as we don't claim any particular correctness properties for  $S$ .

# Impossibility for $n = 3f$


- **Theorem 2:**  $n$  processes can't solve BA, if  $n \leq 3f$ .
- **Proof:**
  - Similar construction, with  $f$  processes treated as a group.
  - Or, can use a **reduction**:
    - Show how to transform a solution for  $n \leq 3f$  to a solution for 3 vs. 1.
    - Since 3 vs. 1 is impossible, this yields a contradiction.
- Treat  $n = 2$  as a special case:
  - $n = 2, f = 1$
  - Each could be faulty, requiring the other to decide on its own value.
  - Or both nonfaulty, which requires agreement, contradiction.





# Transforming A to B

- Algorithm:
  - Partition A-processes into groups  $I_1, I_2, I_3$ , where  $1 \leq |I_1|, |I_2|, |I_3| \leq f$ .
  - Each  $B_i$  process simulates the entire  $I_i$  group.
  - $B_i$  initializes all processes in  $I_i$  with  $B_i$ 's initial value.
  - At each round,  $B_i$  simulates sending messages:
    - Local: Just simulate locally.
    - Remote: Package and send.
  - If any simulated process decides,  $B_i$  decides the same (use any).
- Show B satisfies correctness conditions:
  - Consider any execution of B with at most 1 faulty process.
  - Simulates an execution of A with at most  $f$  faulty processes.
  - Correctness conditions must hold in the simulated execution of A.
  - Show these all carry over to B's execution.





# B's correctness


- **Termination:**
  - If  $B_i$  is nonfaulty in B, then it simulates only nonfaulty processes of A (at least one).
  - Those terminate, so  $B_i$  does also.
- **Agreement:**
  - If  $B_i, B_j$  are nonfaulty processes of B, they simulate only nonfaulty processes of A.
  - Agreement in A implies all these agree.
  - So  $B_i, B_j$  agree.
- **Validity:**
  - If all nonfaulty processes of B start with  $v$ , then so do all nonfaulty processes of A.
  - Then validity of A implies that all nonfaulty A processes decide  $v$ , so the same holds for B.

# General graphs and connectivity bounds


- $n > 3f$  isn't the whole story:
  - 4 processes, can't tolerate 1 fault:
- **Theorem 3:** BA is solvable in an  $n$ -node graph  $G$ , tolerating  $f$  faults, if and only if both of the following hold:
  - $n > 3f$ , and
  - $\text{conn}(G) > 2f$ .
- **$\text{conn}(G)$**  = minimum number of nodes of  $G$  whose removal results in either a disconnected graph or a 1-node graph.



$\text{conn} = 1$



$\text{conn} = 3$




$\text{conn} = 3$

# Proof: “If” direction

- **Theorem 3:** BA is solvable in an  $n$ -node graph  $G$ , tolerating  $f$  faults, if and only if  $n > 3f$  and  $\text{conn}(G) > 2f$ .
- **Proof (“if”):**
  - Suppose both hold.
  - Then we can simulate a total-connectivity algorithm.
  - Key is to emulate reliable communication from any node  $i$  to any other node  $j$ .
  - Rely on **Menger’s Theorem**, which says that a graph is  $c$ -connected (that is, has  $\text{conn} \geq c$ ) if and only if each pair of nodes is connected by  $\geq c$  node-disjoint paths.
  - Since  $\text{conn}(G) \geq 2f + 1$ , we have  $\geq 2f + 1$  node-disjoint paths between  $i$  and  $j$ .
  - To send a message, send it on all these paths (assumes graph is known).
  - Majority must be correct, so take majority message.


# Proof: “Only if” direction

- **Theorem 3:** BA is solvable in an  $n$ -node graph  $G$ , tolerating  $f$  faults, if and only if  $n > 3f$  and  $\text{conn}(G) > 2f$ .
- **Proof (“only if”):**
  - We already showed  $n > 3f$ ; remains to show  $\text{conn}(G) > 2f$ .
  - Show key idea with simple case,  $\text{conn} = 2$ ,  $f = 1$ .
  - Canonical example:
    - Disconnect 1 and 3 by removing 2 and 4:
  - Proof by contradiction.
  - Assume some algorithm  $A$  that solves BA in this canonical graph, tolerating 1 failure.




# Proof (conn = 2, 1 failure)

- Now construct S from two copies of A.
- Consider 1, 2, and 3 in S:
  - Looks to them like they're in A, with a faulty process 4.
  - In A, 1, 2, and 3 must decide 0
  - So they decide 0 in S also.
- Similarly, 1', 2', and 3' decide 1 in S.



# Proof (conn = 2, 1 failure)

- Finally, consider 3', 4', and 1 in S:
  - Looks to them like they're in A, with a faulty process 2.
  - In A, they must agree, so they also agree in S.
  - But 3' decides 0 and 1 decides 1 in S, contradiction.
- Therefore, we can't solve BA in this canonical graph, with 1 failure.
- As before, we can generalize to  $\text{conn}(G) \leq 2f$ , or use a reduction.



# Other Byzantine processor bounds

- The bounds  $n > 3f$  and  $\text{conn} > 2f$  are fundamental for consensus-style problems with Byzantine failures.
- Same bounds hold, in synchronous settings with  $f$  Byzantine faulty processes, for:
  - Byzantine Firing Squad synchronization problem,
  - Weak Byzantine Agreement, and
  - Approximate agreement.
- Also, in timed (partially synchronous settings), for maintaining clock synchronization.
- Proofs all use similar “pasting” methods.



# Weak Byzantine Agreement

## [Lamport]



- Correctness conditions for BA:
  - **Agreement:** No two nonfaulty processes decide on different values.
  - **Validity:** If all nonfaulty processes start with the same  $v$ , then  $v$  is the only allowable decision for nonfaulty processes.
  - **Termination:** All nonfaulty processes eventually decide.
- Correctness conditions for Weak BA:
  - **Agreement:** Same as for BA.
  - **Validity:** If **all processes are nonfaulty** and start with the same  $v$ , then  $v$  is the only allowed decision value.
  - **Termination:** Same as for BA.
- Limits the situations where the decision is forced to go a certain way.
- Similar style to one direction of the validity condition for the 2-Generals problem.

# WBA Processor Bounds

- **Theorem 4:** Weak BA is solvable in an  $n$ -node graph  $G$ , tolerating  $f$  faults, if and only if  $n > 3f$  and  $\text{conn}(G) > 2f$ .
- Same bounds as for BA.
- **Proof:**
  - “If”: Follows from results for ordinary BA.
  - “Only if”:
    - By constructions like those for ordinary BA, but slightly more complicated.
    - Show 3 vs. 1 here, rest LTTR.


# Proof (3 vs. 1 Weak BA)

- By contradiction. Suppose algorithm A, consisting of procs 1, 2, 3, solves WBA with 1 fault.
- Let  $\alpha_0$  = execution in which everyone starts with 0 and there are no failures; results in decision 0.
- Let  $\alpha_1$  = execution in which everyone starts with 1 and there are no failures; results in decision 1.
- Let  $b$  = an upper bound on number of rounds for all processes to decide, in both  $\alpha_0$  and  $\alpha_1$ .
- Construct new system S from  $2b$  copies of A:




# Proof (3 vs. 1 Weak BA)

- **Claim:** Any two adjacent processes in  $S$  must decide the same thing..
  - Because it looks to them like they are in  $A$ , and they must agree in  $A$ .
- So everyone decides the same in  $S$ .
- WLOG, all decide 1.



# Proof (3 vs. 1 Weak BA)

- Now consider a block of  $2b + 1$  consecutive processes that begin with 0:



- Claims:
  - To all but the endpoints, the execution of  $S$  is indistinguishable from  $\alpha_0$ , the failure-free execution of  $A$  in which everyone starts with 0, for one round.
  - To all but two at each end, indistinguishable from  $\alpha_0$  for two rounds.
  - To all but three at each end, indist. from  $\alpha_0$  for three rounds.
  - ...
  - To midpoint, indistinguishable for  $b$  rounds.
- But  $b$  rounds are enough for the midpoint to decide 0, contradicting the fact that everyone decides 1 in  $S$ .

Lower bound on the number of rounds for Byzantine agreement

# Lower bound on number of rounds

- Notice that  **$f+1$  rounds** are used in all the agreement algorithms we've seen so far---both stopping and Byzantine.
- **That's inherent:**  $f+1$  rounds are needed in the worst-case, even for simple **stopping failures**.
- Assume an  $f$ -round stopping agreement algorithm  $A$  tolerating  $f$  faults, and get a contradiction.
- Restrictions on  $A$  (WLOG):
  - $n$ -node complete graph.
  - Decisions at end of round  $f$ .
  - $V = \{0,1\}$
  - All-to-all communication at every round  $\leq f$ .




# Special case: $f = 1$

- **Theorem 5:** Suppose  $n \geq 3$ . There is no  $n$ -process 1-fault stopping agreement algorithm in which nonfaulty processes always decide at the end of round 1.
- **Proof:** Suppose  $A$  exists.
  - Construct a chain of executions, each with at most one failure, such that:
    - First has (unique) decision value 0.
    - Last has decision value 1.
    - Any two consecutive executions in the chain are indistinguishable to some process  $i$  that is nonfaulty in both. So  $i$  must decide the same in both executions, and the two must have the same decision values.
  - So decision values in first and last executions must be the same.
  - Contradiction.







# Round lower bound, $f = 1$

- $\alpha_0$ : All processes have input 0, no failures.
- ...
- $\alpha_k$  (last one): All inputs 1, no failures.
- Start the chain from  $\alpha_0$ .
- Next execution,  $\alpha_1$ , removes message  $1 \rightarrow 2$ .
  - $\alpha_0$  and  $\alpha_1$  indistinguishable to everyone except 1 and 2; since  $n \geq 3$ , there is some other process.
  - These processes are nonfaulty in both executions.
- Next execution,  $\alpha_2$ , removes message  $1 \rightarrow 3$ .
  - $\alpha_1$  and  $\alpha_2$  indistinguishable to everyone except 1 and 3, hence to some nonfaulty process.
- Next, remove message  $1 \rightarrow 4$ .
  - Indistinguishable to some nonfaulty process.




# Continuing...

- Having removed all of process 1's messages, change 1's input from 0 to 1.
  - Looks the same to everyone else.
- We can't just keep removing messages, since we are allowed at most one failure in each execution.
- So, we continue by replacing missing messages, one at a time.
- Repeat with process 2, 3, and 4, eventually reach the last execution: all inputs 1, no failures.




# Special case: $f = 2$

- **Theorem 6:** Suppose  $n \geq 4$ . There is no  $n$ -process 2-fault stopping agreement algorithm in which nonfaulty processes always decide at the end of round 2.
- **Proof:** Suppose  $A$  exists.
  - Construct another chain of executions, each with at most 2 failures.
    - This chain is longer and more complicated.
  - Start with  $\alpha_0$ : All processes have input 0, no failures, 2 rounds:
  - Work toward  $\alpha_n$ , all 1's, no failures.
  - Each consecutive pair is indistinguishable to some nonfaulty process.
  - Use intermediate execs  $\alpha_i$ , in which:
    - Processes  $1, \dots, i$  have initial value 1.
    - Processes  $i+1, \dots, n$  have initial value 0.
    - No failures.




# Special case: $f = 2$

- Show how to connect  $\alpha_0$  and  $\alpha_1$ .
  - That is, change process 1's initial value from 0 to 1.
  - Other intermediate steps are essentially the same.
- Start with  $\alpha_0$ , work toward killing p1 at the beginning, to change its initial value, by removing messages.
- Then replace the messages, working back up to  $\alpha_1$ .
- Start by removing p1's round 2 messages, one by one.
- Q: Continue by removing p1's round 1 messages?
- No, because consecutive executions would not look the same to anyone:
  - E.g., removing  $1 \rightarrow 2$  at round 1 allows p2 to tell everyone about the failure, at round 2.




# Special case: $f = 2$

- Removing  $1 \rightarrow 2$  at round 1 allows p2 to tell all other processes about the failure:



vs.



- Distinguishable to everyone.
- So we must do something more elaborate.
- Recall that now we can allow 2 processes to fail in some executions.
- Use many steps to remove a single round 1 message  $1 \rightarrow i$ ; in these steps, both 1 and  $i$  will be faulty.

# Removing p1's round 1 messages

- Start with execution where p1 sends to everyone at round 1, and only p1 is faulty.
- Remove round 1 message  $1 \rightarrow 2$ :
  - p2 starts out nonfaulty, so sends all its round 2 messages.
  - Now make p2 faulty.
  - Remove p2's round 2 messages, one by one, until we reach an execution where  $1 \rightarrow 2$  at round 1, but p2 sends no round 2 messages.
  - Now remove the round 1 message  $1 \rightarrow 2$ .
    - Executions look the same to everyone but 1 and 2 (and they're all nonfaulty).
  - Replace the round 2 messages from p2, one by one, until p2 is no longer faulty.
- Repeat to remove p1's round 1 messages to p3, p4,...
- After removing all of p1's round 1 messages, change p1's initial value from 0 to 1, as needed.

# General case: Any $f$

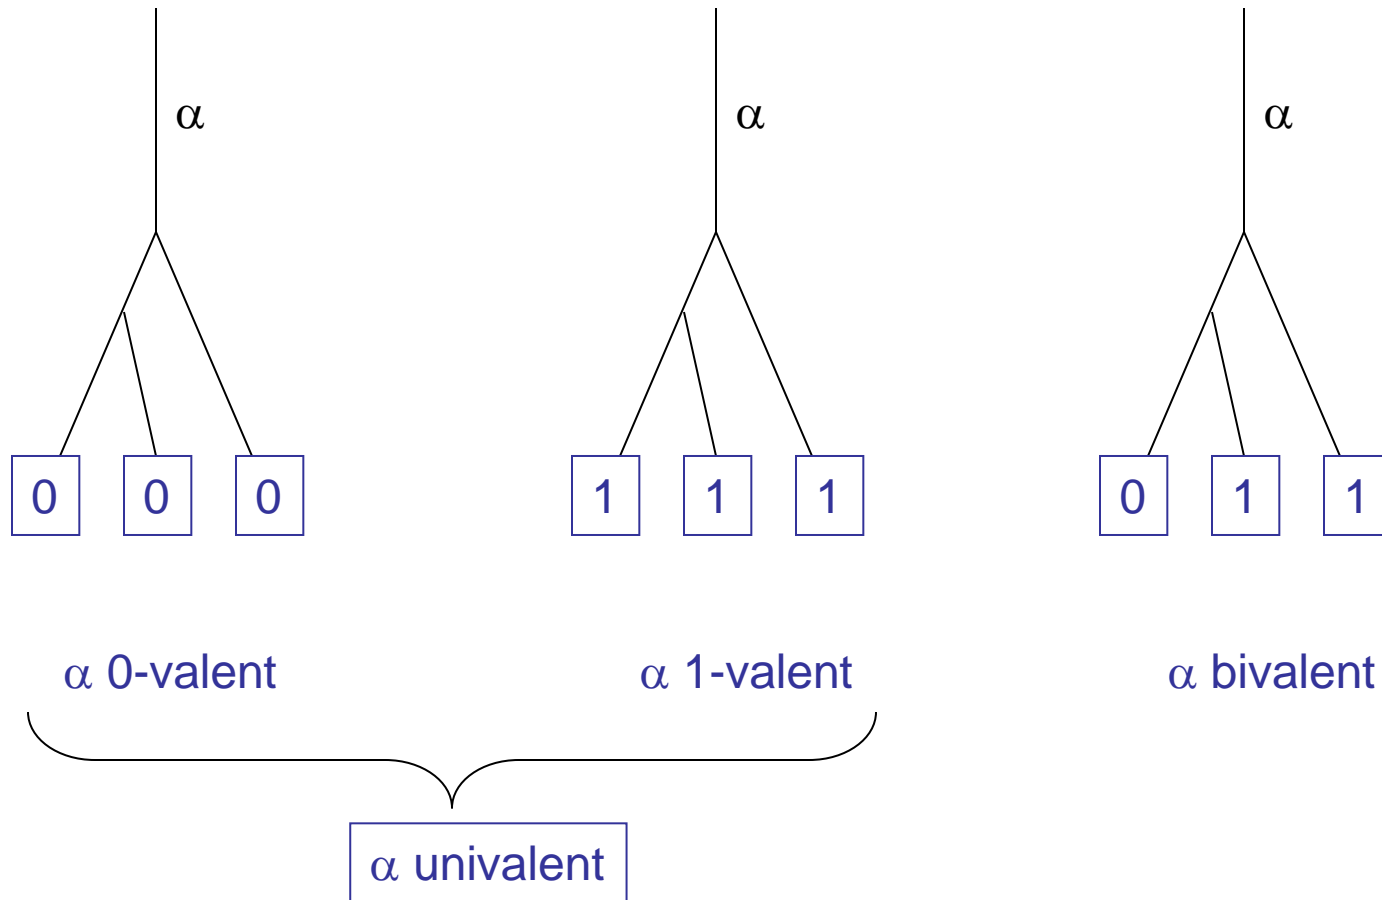
- **Theorem 7:** Suppose  $n \geq f + 2$ . There is no  $n$ -process  $f$ -fault stopping agreement algorithm in which nonfaulty processes always decide at the end of round  $f$ .
- **Proof:** Suppose  $A$  exists.
  - Same ideas, longer chain.
  - Must fail  $f$  processes in some executions in the chain, in order to remove all the required messages, at all rounds.
  - Construction in book, LTTR.
- **Newer proof [Aguilera, Toueg]:**
  - Uses ideas from [Fischer, Lynch, Paterson] impossibility of consensus (which you will see later).
  - They assume strong validity, but their proof works for our weaker validity condition also.

# [Aguilera, Toueg] proof

- By contradiction. Assume A solves stopping agreement for  $f$  failures and everyone decides after exactly  $f$  rounds.
- Consider only executions in which at most one process fails during each round.
- Recall: Failure at a round allows process to send any subset of the messages, or to send all but halt before changing state.
- Regard vector of initial values as a 0-round execution.
- **Definitions** (adapted from [FLP]):  $\alpha$ , an execution that completes some finite number (possibly 0) of rounds, is:
  - **0-valent**, if 0 is the only decision that can occur in any execution (of the kind we consider) that extends  $\alpha$ .
  - **1-valent**, if 1 is the only decision that can occur in...
  - **Univalent**, if  $\alpha$  is either 0-valent or 1-valent (essentially decided).
  - **Bivalent**, if both decisions occur in some extensions (undecided).



# Univalence and Bivalence




# Initial bivalence

- **Lemma 1:** There is some 0-round execution (vector of initial values) that is bivalent.
- **Proof** (derived from **[FLP]**):
  - Assume for contradiction that all 0-round executions are univalent.
  - 000...0 is 0-valent.
  - 111...1 is 1-valent.
  - So there must be two 0-round executions that differ in the value of just one process,  $i$ , such that one is 0-valent and the other is 1-valent.
  - But this is impossible, because if  $i$  fails at the start, no one else can distinguish the two 0-round executions.


# Bivalence through $f-1$ rounds

- **Lemma 2:** For every  $k$ ,  $0 \leq k \leq f-1$ , there is a bivalent  $k$ -round execution.
- **Proof:** By induction on  $k$ .
  - **Base ( $k=0$ ):** Lemma 1.
  - **Inductive step:** Assume for  $k$ , show for  $k+1$ , where  $k < f-1$ .
    - Assume a bivalent  $k$ -round execution  $\alpha$ .
    - Assume for contradiction that every 1-round extension of  $\alpha$  (with at most one new failure) is univalent.
    - Let  $\alpha^*$  be the 1-round extension of  $\alpha$  in which no new failures occur in round  $k+1$ .
    - By assumption, this is univalent, say WLOG that it's 1-valent.
    - Since  $\alpha$  is bivalent, there must be another 1-round extension of  $\alpha$ ,  $\alpha^0$ , that is 0-valent.



# Bivalence through $f-1$ rounds

- In  $\alpha^0$ , some single process, say  $i$ , fails in round  $k+1$ , by not sending to some set of processes, say  $J = \{j_1, j_2, \dots, j_m\}$ .
- Define a chain of  $(k+1)$ -round executions,  $\alpha^0, \alpha^1, \alpha^2, \dots, \alpha^m$ .
- Each  $\alpha^l$  in this sequence is the same as  $\alpha^0$  except that  $i$  also sends messages to  $j_1, j_2, \dots, j_l$ .
  - Adding in messages from  $i$ , one at a time.
- Each  $\alpha^l$  is univalent, by assumption.
- Since  $\alpha^0$  is 0-valent, either:
  - At least one of these is 1-valent, or
  - All are 0-valent.



## Case 1: At least one $\alpha^l$ is 1-valent

- Then there must be some  $l$  such that  $\alpha^{l-1}$  is 0-valent and  $\alpha^l$  is 1-valent.
- But  $\alpha^{l-1}$  and  $\alpha^l$  differ after round  $k+1$  only in the state of one process,  $j_l$ .
- We can extend both  $\alpha^{l-1}$  and  $\alpha^l$  by simply failing  $j_l$  at beginning of round  $k+2$ .
  - There is actually a round  $k+2$  because we've assumed  $k < f-1$ , so  $k+2 \leq f$ .
- And no one left alive can tell the difference!
- Contradiction for Case 1.

## Case 2: Every $\alpha^l$ is 0-valent


- Then compare:
  - $\alpha^m$ , in which  $i$  sends all its round  $k+1$  messages and then fails, with
  - $\alpha^*$ , in which  $i$  sends all its round  $k+1$  messages and does not fail.
- No other differences, since only  $i$  fails at round  $k+1$  in  $\alpha^m$ .
- $\alpha^m$  is 0-valent and  $\alpha^*$  is 1-valent.
- Extend to full  $f$ -round executions:
  - $\alpha^m$ , by allowing no further failures,
  - $\alpha^*$ , by failing  $i$  right after round  $k+1$  and then allowing no further failures.
- No one can tell the difference.
- Contradiction for Case 2.

# Bivalence through $f-1$ rounds

- So we've proved, so far:
- **Lemma 2:** For every  $k$ ,  $0 \leq k \leq f-1$ , there is a bivalent  $k$ -round execution.

# Disagreement after $f$ rounds


- **Lemma 3:** There is an  $f$ -round execution in which two nonfaulty processes decide differently.
- **Proof:**
  - Use Lemma 2 to get a bivalent  $(f-1)$ -round execution  $\alpha$  with  $\leq f-1$  failures.
  - In every 1-round extension of  $\alpha$ , everyone who hasn't failed must decide (and agree).
  - Let  $\alpha^*$  be the 1-round extension of  $\alpha$  in which no new failures occur in round  $f$ .
  - Everyone who is still alive decides after  $\alpha^*$ , and they must decide the same thing. WLOG, say they decide 1.
  - Since  $\alpha$  is bivalent, there must be another 1-round extension of  $\alpha$ , say  $\alpha^0$ , in which some nonfaulty process (and so, all nonfaulty processes) decide 0.





# Disagreement after f rounds

- In  $\alpha^0$ , some single process  $i$  fails in round  $f$ .
- Let  $j, k$  be two nonfaulty processes.
- Define a chain of three  $f$ -round executions,  $\alpha^0, \alpha^1, \alpha^*$ , where  $\alpha^1$  is identical to  $\alpha^0$  except that  $i$  sends to  $j$  in  $\alpha^1$  (it might not in  $\alpha^0$ ).
- Then  $\alpha^1 \sim^k \alpha^0$ .
- Since  $k$  decides 0 in  $\alpha^0$ ,  $k$  also decides 0 in  $\alpha^1$ .
- Also,  $\alpha^1 \sim^j \alpha^*$ .
- Since  $j$  decides 1 in  $\alpha^*$ ,  $j$  also decides 1 in  $\alpha^1$ .
- Yields disagreement in  $\alpha^1$ , contradiction!
- So we've proved:
- **Lemma 3:** There is an  $f$ -round execution in which two nonfaulty processes decide differently.
- Which immediately yields the lower bound result.



# Early-stopping agreement algorithms

- Tolerate  $f$  failures, but in executions with  $f' < f$  failures, terminate correspondingly faster.
- [Dolev, Reischuk, Strong 90] Gave a stopping agreement algorithm in which all nonfaulty processes terminate in at most  $\min(f' + 2, f+1)$  rounds.
  - If  $f' + 2 \leq f$ , decide “early”, within  $f' + 2$  rounds; in any case decide within  $f+1$  rounds.
- [Keidar, Rajsbaum 02] Lower bound of  $f' + 2$  for early-stopping agreement.
  - Not just  $f' + 1$ . Early stopping requires an extra round.

# Early-stopping agreement algorithms

- Tolerate  $f$  failures, but in executions with  $f' < f$  failures, terminate correspondingly faster.
- [Keidar, Rajsbaum 02] Lower bound of  $f' + 2$  for early-stopping agreement.
  - Not just  $f' + 1$ . Early stopping requires an additional round.
- **Theorem 1:** Assume  $0 \leq f' \leq f - 2$  and  $f < n$ . Every early-stopping agreement algorithm tolerating  $f$  failures has an execution with  $f'$  failures in which some nonfaulty process doesn't decide by the end of round  $f' + 1$ .

# Special case: $f' = 0$

- **Special Case Theorem 2:** Assume  $2 \leq f < n$ . Every early-stopping agreement algorithm tolerating  $f$  failures has a **failure-free execution** in which some nonfaulty process does not decide by the end of round 1.
- **Definition:** Let  $\alpha$  be an execution that completes some finite number (possibly 0) of rounds. Then  $\text{val}(\alpha)$  is the unique decision value in the extension of  $\alpha$  with no new failures.
  - Different from bivalence defs---now consider value in just one particular extension.

# Special case: $f' = 0$

- **Theorem 2:** Assume  $2 \leq f < n$ . Every early-stopping agreement algorithm tolerating  $f$  failures has a failure-free execution in which some nonfaulty process does not decide by the end of round 1.
- **Definition:**  $\text{val}(\alpha)$  is the decision value in the extension of  $\alpha$  with no new failures.
- **Proof of Theorem 2:**
  - Assume executions in which at most one process fails per round.
  - Identify 0-round executions with vectors of initial values.
  - Assume, for contradiction, that everyone decides by the end of round 1, in all failure-free executions.
  - $\text{val}(000\dots 0) = 0$ ,  $\text{val}(111\dots 1) = 1$ .
  - So there must be two 0-round executions  $\alpha^0$  and  $\alpha^1$ , that differ in the value of just one process  $i$ , such that  $\text{val}(\alpha^0) = 0$  and  $\text{val}(\alpha^1) = 1$ .

# Special case: $f' = 0$

- 0-round executions  $\alpha^0$  and  $\alpha^1$ , differing only in the initial value of process  $i$ , such that  $\text{val}(\alpha^0) = 0$  and  $\text{val}(\alpha^1) = 1$ .
- In the ff extensions of  $\alpha^0$  and  $\alpha^1$ , all nonfaulty processes decide by the end of round 1.
- **Define:**
  - $\beta^0$ , 1-round extension of  $\alpha^0$ , in which process  $i$  fails, sends only to  $j$ .
  - $\beta^1$ , 1-round extension of  $\alpha^1$ , in which process  $i$  fails, sends only to  $j$ .
- **Then:**
  - $\beta^0$  looks to  $j$  like ff extension of  $\alpha^0$ , so  $j$  decides 0 in  $\beta^0$  by round 1.
  - $\beta^1$  looks to  $j$  like ff extension of  $\alpha^1$ , so  $j$  decides 1 in  $\beta^1$  by round 1.
- $\beta^0$  and  $\beta^1$  are indistinguishable to all processes except  $i, j$ .
- **Define:**
  - $\gamma^0$ , infinite extension of  $\beta^0$ , in which process  $j$  fails right after round 1.
  - $\gamma^1$ , infinite extension of  $\beta^1$ , in which process  $j$  fails right after round 1.
- By agreement, all nonfaulty processes must decide 0 in  $\gamma^0$ , 1 in  $\gamma^1$ .
- But  $\gamma^0$  and  $\gamma^1$  are indistinguishable to all nonfaulty processes, so they can't decide differently, contradiction.

# Next time...

- Other kinds of consensus problems:
  - k-agreement
  - Approximate agreement (skim)
  - Distributed commit
- Reading:
  - Chapter 7 (just skim 7.2)