

6.852: Distributed Algorithms

Fall, 2015

Lecture 8

Today's plan

- Lower bound on number of rounds for agreement, cont'd.
- Early-stopping agreement.
- Other consensus-type problems:
 - k -agreement
 - Distributed commit
- Reading:
 - [Aguilera, Toueg]
 - [Keidar, Rajsbaum]
 - Chapter 7 (skim 7.2)
- Next:
 - Modeling asynchronous systems
 - I/O automata
- Reading:
 - Chapter 8

Lower bound on number of rounds
for agreement

Lower bound on number of rounds


- $f+1$ rounds are needed in the worst-case, for either Byzantine agreement or just stopping agreement.
- Assume an f -round stopping agreement algorithm A tolerating f faults, get a contradiction.
- Assume:
 - n -node complete graph.
 - Decisions at end of round f .
 - $V = \{0,1\}$
 - All-to-all communication at every round.

Special case: $f = 1$

- **Theorem 5:** Suppose $n \geq 3$. There is no n -process 1-fault stopping agreement algorithm in which nonfaulty processes always decide at the end of round 1.
- **Proof:**
 - Construct a chain of executions, each with ≤ 1 failure, such that:
 - First has decision value 0.
 - Last has decision value 1.
 - Any two consecutive executions in the chain are indistinguishable to some process i that is nonfaulty in both. So i must decide the same in both executions, and the two must have the same decision values.
 - So decision values in first and last executions must be the same.
 - Contradiction.


Special case: $f = 2$

- **Theorem 6:** Suppose $n \geq 4$. There is no n -process 2-fault stopping agreement algorithm in which nonfaulty processes always decide at the end of round 2.
- **Proof:**
 - Construct a chain of executions, each with ≤ 2 failures.
 - Start with α_0 : All processes have input 0, no failures
 - Work toward α_n , all 1's, no failures.
 - Each consecutive pair is indistinguishable to some nonfaulty process.
 - Use intermediate executions α_i in which:
 - Processes $1, \dots, i$ have initial value 1.
 - Processes $i+1, \dots, n$ have initial value 0.
 - No failures.



Special case: $f = 2$

- WLOG, show how to connect α_0 and α_1 , that is, change p1's initial value from 0 to 1.
- Start with α_0 , work toward killing p1 at the beginning, by removing messages.
- Change p1's initial value.
- Then replace messages, working back up to α_1 .
- Start by removing p1's round 2 messages, one by one.
- Can't continue by removing p1's round 1 messages, because consecutive executions would not look the same to anyone.
- E.g., removing $1 \rightarrow 2$ at round 1 allows p2 to tell everyone about the failure, at round 2.
- So, use many steps of the chain to remove the round 1 message from p1 to p2.
- In these steps, both p1 and p2 are faulty.



Removing p1's round 1 messages

- Start with execution where p1 sends to everyone at round 1, and to no one at round 2. Only p1 is faulty.
- Remove round 1 message $1 \rightarrow 2$:
 - p2 starts out nonfaulty, so sends all its round 2 messages.
 - Now fail p2, and remove its round 2 messages, one by one, until we reach an execution where $1 \rightarrow 2$ at round 1, but p2 sends no round 2 messages.
 - Now remove the round 1 message $1 \rightarrow 2$.
 - Executions look the same to everyone but p1 and p2.
 - Replace round 2 messages from p2, one by one, until p2 is no longer faulty.
- Repeat to remove p1's round 1 messages to p3, p4,...
- After removing all of p1's round 1 messages, change p1's initial value from 0 to 1.


General case: Any f

- **Theorem 7:** Suppose $n \geq f + 2$. There is no n -process f -fault stopping agreement algorithm in which nonfaulty processes always decide at the end of round f .
- **Proof:**
 - Same ideas, longer chain.
 - Must fail f processes in some executions in the chain, in order to remove all the required messages, at all rounds.
 - Construction in book, LTTR.
- **Alternative proof [Aguilera, Toueg]:**
 - Uses ideas from [Fischer, Lynch, Paterson] impossibility of consensus (which you will see later).
 - They assume strong validity, but their proof works for our weaker validity condition also.

[Aguilera,Toueg] lower bound proof

- By contradiction. Assume A solves stopping agreement for f failures and everyone decides after exactly f rounds.
- Consider only executions in which at most one process fails during each round.
- **Recall:** Failure at a round allows a process to send any subset of the messages, or to send all but halt before changing state.
- Regard vector of initial values as a 0-round execution.
- **Definitions** (adapted from [FLP]): α , an execution that completes some finite number (possibly 0) of rounds, is:
 - **0-valent**, if 0 is the only decision that can occur in any execution (of the kind we consider) that extends α .
 - **1-valent**, if 1 is the only decision that can occur in...
 - **Univalent**, if α is either 0-valent or 1-valent (essentially decided).
 - **Bivalent**, if both decisions occur in some extensions (undecided).

Univalence and Bivalence




Initial bivalence

- **Lemma 1:** There is some 0-round execution (vector of initial values) that is bivalent.
- **Proof** (derived from [FLP]):
 - Assume for contradiction that all 0-round executions are univalent.
 - 000...0 is 0-valent.
 - 111...1 is 1-valent.
 - So there must be two 0-round executions that differ in the value of just one process, i , such that one is 0-valent and the other is 1-valent.
 - But this is impossible, because if i fails at the start, no one else can distinguish the two 0-round executions.


Bivalence through $f-1$ rounds

- **Lemma 2:** For every k , $0 \leq k \leq f-1$, there is a bivalent k -round execution.
- **Proof:** By induction on k .
 - **Base ($k=0$):** Lemma 1.
 - **Inductive step:** Assume for k , show for $k+1$, where $k < f-1$.
 - Assume a bivalent k -round execution α .
 - Assume for contradiction that every 1-round extension of α (with at most one new failure) is univalent.
 - Let α^* be the 1-round extension of α in which no new failures occur in round $k+1$.
 - By assumption, this is univalent, say WLOG that it's 1-valent.
 - Since α is bivalent, there must be another 1-round extension of α , α^0 , that is 0-valent.



Bivalence through $f-1$ rounds

- In α^0 , some single process, say i , fails in round $k+1$, by not sending to some set of processes, say $J = \{j_1, j_2, \dots, j_m\}$.
- Define a chain of $(k+1)$ -round executions, $\alpha^0, \alpha^1, \alpha^2, \dots, \alpha^m$.
- Each α^l in this sequence is the same as α^0 except that i also sends messages to j_1, j_2, \dots, j_l .
 - Adding in messages from i , one at a time.
- Each α^l is univalent, by assumption.
- Since α^0 is 0-valent, either:
 - At least one of these is 1-valent, or
 - All are 0-valent.



Case 1: At least one α^l is 1-valent

- Then there must be some l such that α^{l-1} is 0-valent and α^l is 1-valent.
- But α^{l-1} and α^l differ after round $k+1$ only in the state of one process, j_l .
- We can extend both α^{l-1} and α^l by simply failing j_l at beginning of round $k+2$.
 - There is actually a round $k+2$ because we've assumed $k < f-1$, so $k+2 \leq f$.
- And no one left alive can tell the difference!
- Contradiction for Case 1.

Case 2: Every α^l is 0-valent


- Then compare:
 - α^m , in which i sends all its round $k+1$ messages and then fails, with
 - α^* , in which i sends all its round $k+1$ messages and does not fail.
- No other differences, since only i fails at round $k+1$ in α^m .
- α^m is 0-valent and α^* is 1-valent.
- Extend to full f -round executions:
 - α^m , by allowing no further failures,
 - α^* , by failing i right after round $k+1$ and then allowing no further failures.
- No one can tell the difference.
- Contradiction for Case 2.

Bivalence through $f-1$ rounds

- So we've proved, so far:
- **Lemma 2:** For every k , $0 \leq k \leq f-1$, there is a bivalent k -round execution.


Disagreement after f rounds

- **Lemma 3:** There is an f -round execution in which two nonfaulty processes decide differently.
- **Proof:**
 - Use Lemma 2 to get a bivalent $(f-1)$ -round execution α with $\leq f-1$ failures.
 - In every 1-round extension of α , everyone who hasn't failed must decide (and agree).
 - Let α^* be the 1-round extension of α in which no new failures occur in round f .
 - Everyone who is still alive decides after α^* , and they must decide the same thing. WLOG, say they decide 1.
 - Since α is bivalent, there must be another 1-round extension of α , say α^0 , in which some nonfaulty process (and so, all nonfaulty processes) decide 0.



Disagreement after f rounds

- In α^0 , some single process i fails in round f .
- Let j, k be two nonfaulty processes.
- Define a chain of three f -round executions, $\alpha^0, \alpha^1, \alpha^*$, where α^1 is identical to α^0 except that i sends to j in α^1 (it might not in α^0).
- Then $\alpha^1 \sim^k \alpha^0$.
- Since k decides 0 in α^0 , k also decides 0 in α^1 .
- Also, $\alpha^1 \sim^j \alpha^*$.
- Since j decides 1 in α^* , j also decides 1 in α^1 .
- Yields disagreement in α^1 , contradiction!



- So we've proved:
- **Lemma 3:** There is an f -round execution in which two nonfaulty processes decide differently.
- Which immediately yields the lower bound result.

Early-stopping agreement algorithms

- Tolerate f failures, but in executions with $f' < f$ failures, terminate correspondingly faster.
- [Dolev, Reischuk, Strong 90] :
- Stopping agreement algorithm in which all nonfaulty processes terminate within $\min(f' + 2, f + 1)$ rounds:
 - Always decide within $f + 1$ rounds.
 - If $f' + 2 \leq f$, decide “early”, within $f' + 2$ rounds.
- [Keidar, Rajsbaum 02]:
- Lower bound of $f' + 2$ for early-stopping agreement.
 - Not just $f' + 1$. Early stopping requires an extra round.
- **Theorem 1:** Assume $0 \leq f' \leq f - 2$ and $f < n$. Every early-stopping agreement algorithm tolerating f failures has an execution with f' failures in which some nonfaulty process doesn't decide by the end of round $f' + 1$.

Special case: $f' = 0$

- **Special case Theorem 2:** Assume $2 \leq f < n$. Every early-stopping agreement algorithm tolerating f failures has a **failure-free execution** in which some nonfaulty process does not decide by the end of round 1.
- **Definition:** Let α be an execution that completes some finite number (possibly 0) of rounds. Then $val(\alpha)$ is the unique decision value in the extension of α with no new failures.
 - Different from bivalence definitions from [Aguilera, Toueg] ---now consider value in just one extension.

Special case: $f' = 0$

- **Theorem 2:** Assume $2 \leq f < n$. Every early-stopping agreement algorithm tolerating f failures has a failure-free execution in which some nonfaulty process does not decide by the end of round 1.
- **Definition:** $val(\alpha)$ is the decision value in the extension of α with no new failures.
- **Proof of Theorem 2:**
 - Consider executions in which at most one process fails per round.
 - Identify 0-round executions with vectors of initial values.
 - Assume, for contradiction, that everyone decides by the end of round 1, in all failure-free executions.
 - $val(000 \dots 0) = 0, val(111 \dots 1) = 1$.
 - So there must be two 0-round executions α^0 and α^1 , that differ in the value of just one process i , such that $val(\alpha^0) = 0$ and $val(\alpha^1) = 1$.

Special case: $f' = 0$

- 0-round executions α^0 and α^1 , differing only in the initial value of process i , such that $\text{val}(\alpha^0) = 0$ and $\text{val}(\alpha^1) = 1$.
- In the failure-free extensions of α^0 and α^1 , all nonfaulty processes decide by the end of round 1.
- **Define:**
 - β^0 , 1-round extension of α^0 , in which process i fails, sends only to j .
 - β^1 , 1-round extension of α^1 , in which process i fails, sends only to j .
- **Then:**
 - β^0 looks to j like ff extension of α^0 , so j decides 0 in β^0 by round 1.
 - β^1 looks to j like ff extension of α^1 , so j decides 1 in β^1 by round 1.
- β^0 and β^1 are indistinguishable to all processes except i, j .
- **Define:**
 - γ^0 , infinite extension of β^0 , in which process j fails right after round 1.
 - γ^1 , infinite extension of β^1 , in which process j fails right after round 1.
- By agreement, all nonfaulty processes must decide 0 in γ^0 , 1 in γ^1 .
- But γ^0 and γ^1 are indistinguishable to all nonfaulty processes, so they can't decide differently, contradiction.

k -Agreement

k -agreement

- Also called k -set-agreement or k -set-consensus.
- Generalizes ordinary stopping agreement by allowing k different decisions instead of just one.
- Motivation:
 - Practical:
 - Allocating shared resources, e.g., agreeing on small number of radio frequencies to use for sending/receiving broadcasts.
 - Mathematical:
 - Natural generalization of ordinary 1-agreement.
 - Elegant theory: Nice topological structure, tight bounds.

The k -agreement problem

- Assume:
 - n -node complete undirected graph
 - Stopping failures only
 - Inputs, decisions in a finite totally-ordered set V (appear in state variables).
- Correctness conditions:
 - **Agreement:**
 - $\exists W \subseteq V, |W| = k$, all decision values in W .
 - That is, there are at most k different decision values.
 - **Validity:**
 - Any decision value is some process' initial value.
 - Like strong validity for 1-agreement.
 - **Termination:**
 - All nonfaulty processes eventually decide.

FloodMin k -agreement algorithm

- **Algorithm:**
 - Each process remembers the minimum value it has seen, initially its own value.
 - At each round, broadcasts its *min* value.
 - Decide after some generally-agreed-upon number of rounds, on current *min* value.
- **Q:** How many rounds are enough?
- **1-agreement:** $f + 1$ rounds
 - Argument like those for previous stopping agreement algorithms (LTTR).
- **k -agreement:** $\lfloor f/k \rfloor + 1$ rounds.
- Allowing k values **divides** the runtime by k .

FloodMin correctness

- **Theorem 1:** *FloodMin*, for $\lfloor f/k \rfloor + 1$ rounds, solves k -agreement.
- **Proof:**
- Define $M(r)$ = set of *min* values of active (not-yet-failed) processes after r rounds.
- This set can only decrease over time:
- **Lemma 1:** $M(r + 1) \subseteq M(r)$ for every r , $0 \leq r \leq \lfloor f/k \rfloor$.
- **Proof:** Any *min* value after round $r + 1$ is someone's *min* value after round r .

Proof of Theorem 1, cont'd

- **Lemma 2:** If at most $d - 1$ processes fail during round r , then $|M(r)| \leq d$.
- E.g., for $d = 1$: If no one fails during round r then all have the same *min* value after round r .
- **Proof:** Show the contrapositive.
 - Suppose that $|M(r)| > d$, show at least d processes must fail in round r .
 - Let $m = \max(M(r))$.
 - Let $m' < m$ be any other element of $M(r)$.
 - Then $m' \in M(r - 1)$ by Lemma 1.
 - Let i be a process that is active after $r - 1$ rounds and that has *min* = m' just after $r - 1$ rounds.
 - Claim i fails during round r :
 - If not, then everyone would receive m' in round r .
 - But then no one would choose $m > m'$ as its *min*, contradiction.
 - But this is true for every $m' < m$ in $M(r)$, so at least d processes fail in round r .

Proof of Theorem 1, cont'd

- **Validity:** Easy
- **Termination:** Obvious
- **Agreement:** By contradiction.
 - Assume an execution with $> k$ different decision values.
 - Then the number of *min* values for active processes after the full $\lfloor f/k \rfloor + 1$ rounds is $> k$.
 - That is, $|M(\lfloor f/k \rfloor + 1)| > k$.
 - Then by Lemma 1, $|M(r)| > k$ for every $r, 0 \leq r \leq \lfloor f/k \rfloor + 1$.
 - So by Lemma 2, at least k processes fail in each round.
 - That's at least $(\lfloor f/k \rfloor + 1) k$ total failures, which is $> f$ failures.
 - Contradiction!

Rounds for k -agreement


- **Theorem 1:** *FloodMin*, for $\lfloor f/k \rfloor + 1$ rounds, solves k -agreement.
- This is a tight bound!
- **Theorem 2:** Any algorithm for k -agreement requires $\geq \lfloor f/k \rfloor + 1$ rounds.

Lower Bound (sketch)

- **Theorem 2:** Any algorithm for k -agreement requires $\geq \lfloor f/k \rfloor + 1$ rounds.
- Recall old proof for $f + 1$ -round lower bound for 1-agreement.
 - Chain of executions for assumed algorithm:
$$\alpha_0 \text{ ----- } \alpha_1 \text{ ----- } \dots \text{ ----- } \alpha_j \text{ ----- } \alpha_{j+1} \text{ ----- } \dots \text{ ----- } \alpha_m$$
 - Each execution has a unique decision value.
 - Executions at ends of chain have specified decision values.
 - Two consecutive executions look the same to some nonfaulty process, who (therefore) decides the same in both.
- This argument doesn't extend immediately to k -agreement:
 - Can't assume a unique value in each execution.
 - Example: For 2-agreement, could have 3 different values in 2 consecutive executions without violating agreement.
- Instead, use a **k -dimensional generalized chain**.


Lower bound

- Assume, for contradiction:
 - An n -process k -agreement algorithm tolerating f failures.
 - $n \geq f + k + 1$ (so each execution we consider has $\geq k + 1$ nonfaulty processes)
 - All-to-all communication at all rounds.
 - $V = \{0, 1, \dots, k\}$, $k + 1$ values.
 - All processes decide just after round r , where $r \leq \lfloor f/k \rfloor$.
- Get a contradiction by finding an execution with $k + 1$ different decision values.
- Use a k -dimensional collection of executions rather than 1-dimensional.
 - $k = 2$: Triangle
 - $k = 3$: Tetrahedron, etc.



Labeling nodes with executions

- **Bermuda Triangle ($k = 2$):** Any algorithm must vanish somewhere in the interior.
- Label nodes with executions:
 - Corner: No failures, all have same initial value.
 - Boundary edge: Initial values chosen from those of the two endpoints
 - For $k > 2$, generalize to boundary faces.
 - Interior: Mixture of inputs
- Label so executions “morph gradually” in all directions:
- Difference between two adjacent executions along an edge:
 - Remove or add one message, to a process that fails immediately.
 - Fail or recover a process.
 - Change initial value of failed process.



Labeling nodes with process names

- Also label each node with the name of a process that is nonfaulty in the node's execution; indices chosen for the corners of any tiny triangle (simplex) are distinct.
- **Consistency property:** For every tiny triangle T , there is a single execution β , with at most f faults, that is “compatible” with the executions and processes labeling the corners of T :
 - All the corner-labeling processes are nonfaulty in β .
 - If (α, i) labels some corner of T , then α is indistinguishable from β by i .
- Formalizes the “gradual morphing” property.
- Proof by laborious, detailed construction.
- Can recast chain arguments for 1-agreement in this style:

$$\begin{array}{ccccccc}
 \alpha_0 & \text{-----} & \alpha_1 & \text{-----} & \dots & \text{-----} & \alpha_j & \overset{\beta}{\text{-----}} & \alpha_{j+1} & \text{-----} & \dots & \text{-----} & \alpha_m \\
 p_0 & & p_1 & & \dots & & p_j & & p_{j+1} & & & & p_m
 \end{array}$$


- β indistinguishable by p_j from α_j
- β indistinguishable by p_{j+1} from α_{j+1}

Bound on rounds

- This labeling construction uses the assumption that $r \leq \left\lfloor \frac{f}{k} \right\rfloor$, that is, $f \geq r k$.
- **How:**
 - We are essentially constructing chains simultaneously in k directions (2 directions, in the 2-dimensional case).
 - We use r failures (one per round) to construct the “chain” in each direction.
 - For k directions, that’s $r k$ total failures.
- Details LTTR (see book, or paper [Chaudhuri, Herlihy, Lynch, Tuttle])


Coloring the nodes

- Now color each node v with a “color” in $\{0, 1, \dots, k\}$:
 - If v is labeled with (α, i) then $color(v) = i$'s decision value in α .
- Properties:
 - Colors of the major corners are all different.
 - Color of each boundary edge node is the same as one of the endpoint corners.
 - For $k > 2$, generalize to boundary faces.
- Coloring properties follow from Validity, because of the way the initial values are assigned.



Sperner Colorings

- A coloring with the listed properties (suitably generalized to k dimensions) is called a **Sperner Coloring** (in algebraic topology).
- **Sperner's Lemma:** Any Sperner Coloring has some tiny triangle (simplex) whose $k + 1$ corners are colored by all $k + 1$ colors.
- Find one?



Applying Sperner's Lemma

- Apply Sperner's Lemma to the coloring we constructed.
- Yields a tiny triangle (simplex) T with $k + 1$ different colors on its corners.
- Which means $k + 1$ different decision values for the executions and processes labeling its corners.
- But recall that there must be a **single execution β** , with at most f faults, that is “compatible” with the executions and processes labeling the corners of T :
 - All the corner processes are nonfaulty in β .
 - If (α, i) labels some corner of T , then α is indistinguishable from β by i .
- So all the corner processes behave the same in β as they do in their own corner executions, and decide on the same values as in those executions.
- **That's $k + 1$ different decision values in one execution with at most f faults.**
- Contradicts k -agreement.

Approximate Agreement


Approximate Agreement problem

- Agreement on real number values, e.g.:
 - Readings of several altimeters on an aircraft.
 - Values of approximately-synchronized clocks.
- Consider Byzantine participants, e.g., faulty hardware.
- **Abstract approximate agreement problem:**
 - Inputs, outputs are reals
 - **Agreement:** Within ϵ .
 - **Validity:** Within range of initial values of nonfaulty processes.
 - **Termination:** Nonfaulty processes eventually decide.
- Assume: Complete n -node graph, $n > 3f$.
- Could solve by exact BA, using $f + 1$ rounds and lots of communication.
- But better algorithms exist:
 - Simpler, cheaper
 - Convergence strategy
 - Extend to asynchronous settings, whereas BA is unsolvable in asynchronous networks (as we will see).

Distributed Commit

Distributed Commit

- **Motivation:** Distributed database transaction processing
 - A database transaction performs work at several distributed sites.
 - Transaction manager (TM) at each site decides whether it would like to “commit” or “abort” the transaction.
 - Based on whether the transaction’s work has been successfully completed at that site, and results made stable.
 - All TMs must agree on whether to commit or abort.
- **Assume:**
 - Process stopping failures only.
 - n -node, complete, undirected graph.
- **Require:**
 - **Agreement:** No two processes decide differently (faulty or not, uniformity)
 - **Validity:**
 - If any process starts with 0 (abort) then 0 is the only allowed decision.
 - If all start with 1 (commit) and there are no faulty processes then 1 is the only allowed decision.




Correctness Conditions for Commit

- **Agreement:** No two processes decide differently.
- **Validity:**
 - If any process starts with 0 then 0 is the only allowed decision.
 - If all start with 1 and there are no faulty processes then 1 is the only allowed decision.
- Note the asymmetry: Guarantee abort (0) if **anyone** wants to abort; guarantee commit (1) if **everyone** wants to commit **and no one fails** (best case).
- **Termination:**
 - **Weak termination:** If there are no failures then all processes eventually decide.
 - **Strong termination (non-blocking condition):** (Even if there are failures), all nonfaulty processes eventually decide.



2-Phase Commit

- Traditional, blocking algorithm (guarantees weak termination only).
- Assumes distinguished process 1, acts as “coordinator” (leader).
- **Round 1:** All send initial values to process 1, who decides.
 - If it sees 0, or doesn’t hear from someone, it decides 0; otherwise it decides 1.
- **Round 2:** Process 1 sends the decision to everyone else.
- **Q:** When can the processes decide?
- Anyone with initial value 0 can decide at the beginning.
- Process 1 decides after receiving round 1 messages.
- Everyone else decides after round 2 (if there are no failures).






Correctness of 2-Phase Commit

- **Agreement:**
 - Because decision is centralized (and consistent with any individual initial decisions).
- **Validity:**
 - Because of how the coordinator decides.
- **Weak termination:**
 - If no one fails, everyone terminates by end of round 2.
- **Strong termination?**
 - No: If coordinator fails before sending its round 2 messages, then others with initial value 1 will never terminate.



Add a termination protocol?

- We might try to add a termination protocol: other processes try to detect failure of coordinator and finish agreeing on their own.
- But this can't always work:
 - If initial values are 0,1,1,1, then by validity, everyone is required to decide 0.
 - If initial values are 1,1,1,1 and process 1 fails just after deciding, and before sending out its round 2 messages, then:
 - Process 1 will decide 1.
 - By agreement, others must decide 1.
 - But the other processes can't distinguish these two situations.




Complexity of 2-phase commit

- Time:
 - 2 rounds
- Communication:
 - At most $2n$ messages


3-Phase Commit [Skeen]

- Yields strong termination.
- **Trick:** Introduce intermediate stage, before actually deciding.
- Process states are now classified into four categories:
 - *dec0*: Already decided 0.
 - *dec1*: Already decided 1.
 - *ready*: Ready to decide 1 but hasn't yet.
 - *uncertain*: Otherwise.
- Again, process 1 acts as “coordinator”.
- Communication pattern:



3-Phase Commit

- All processes are initially *uncertain*.
- **Round 1:**
 - All other processes send their initial values to p_1 .
 - All with initial value 0 **decide 0** (and enter *dec0* state)
 - If p_1 receives 1s from everyone and its own initial value is 1, p_1 becomes *ready*, but doesn't yet decide.
 - If p_1 sees 0 or doesn't hear from someone, p_1 **decides 0**.
- **Round 2:**
 - If p_1 has decided 0, it broadcasts "decide 0", else it broadcasts "ready".
 - Anyone else who receives "decide 0" **decides 0**.
 - Anyone else who receives "ready" becomes *ready*.
 - Now p_1 **decides 1** if it hasn't already decided.
- **Round 3:**
 - If p_1 has decided 1, it bcasts "decide 1".
 - Anyone else who receives "decide 1" **decides 1**.



3-Phase Commit

- Key invariants (after 0, 1, 2, or 3 rounds):
 - If any process is in *ready* or *dec1*, then all processes have initial value 1.
 - If any process is in *dec0* then:
 - No process is in *dec1*, and no non-failed process is *ready*.
 - If any process is in *dec1* then:
 - No process is in *dec0*, and no non-failed process is *uncertain*.
- **Proof:** LTTR.
 - Key step: Third condition is preserved when p_1 *decides 1* after round 2.
 - In this case, p_1 knows that:
 - Everyone's input is 1.
 - No one *decided 0* at the end of round 1.
 - Every other process has either become *ready* or has failed (without deciding).
 - Implies the third condition.
- **Note critical use of synchrony here:**
 - p_1 infers that non-failed processes are *ready* just because round 2 is completed.
 - Without synchrony, this would require explicit acknowledgments.

Correctness conditions (so far)

- Agreement and validity follow, for these three rounds.
- Weak termination holds
- Strong termination:
 - Doesn't hold yet---must add a termination protocol.
 - Allow process 2 to act as coordinator, then 3,...
 - “Rotating coordinator” strategy

3-Phase Commit

- **Round 4:**
 - All processes send current status (*dec0*, *uncertain*, *ready*, *dec1*) to p_2 .
 - If p_2 receives any *dec0*'s and hasn't already decided, then p_2 **decides 0**.
 - If p_2 receives any *dec1*'s and hasn't already decided, then p_2 **decides 1**.
 - If all received values, and its own value, are *uncertain*, then p_2 **decides 0**.
 - Otherwise (all values are *uncertain* or *ready* and at least one is *ready*), p_2 becomes *ready*, but doesn't decide yet.
- **Round 5** (analogous to round 2):
 - If p_2 has (ever) decided 0, broadcasts "decide 0", and similarly for 1.
 - Else broadcasts "ready".
 - Any undecided process who receives "decide()" decides accordingly.
 - Any process who receives "ready" becomes *ready*.
 - Now p_2 **decides 1** if it hasn't already decided.
- **Round 6** (analogous to round 3):
 - If p_2 has decided 1, broadcasts "decide 1".
 - Anyone else who receives "decide 1" **decides 1**.
- Continue with subsequent rounds for p_3, p_4, \dots

Correctness

- Key invariants still hold:
 - If any process is in *ready* or *dec1*, then all processes have initial value 1.
 - If any process is in *dec0* then:
 - No process is in *dec1*, and no non-failed process is *ready*.
 - If any process is in *dec1* then:
 - No process is in *dec0*, and no non-failed process is *uncertain*.
- Imply agreement, validity
- Strong termination:
 - Because eventually some coordinator will finish the job (unless everyone fails).

Complexity

- Time until everyone decides:
 - Normal case 3
 - Worst case $3n$
- Messages until everyone decides:
 - Normal case $O(n)$
 - Technicality: When can processes stop sending messages?
 - Worst case $O(n^2)$

Practical issues for 3-phase commit


- Depends on strong assumptions, which may be hard to guarantee in practice:
 - Synchronous model:
 - Could emulate with approximately-synchronized clocks, timeouts.
 - Reliable message delivery:
 - Could emulate with acks and retransmissions.
 - But if retransmissions add too much delay, then we can't emulate the synchronous model accurately.
 - Leads to unbounded delays, asynchronous model.
 - Accurate diagnosis of process failures:
 - Get this “for free” in the synchronous model.
 - E.g., 3-phase commit algorithm lets process that doesn't hear from another process i at a round conclude that i must have failed.
 - Very hard to guarantee in practice: In Internet, or even a LAN, how to reliably distinguish failure of a process from lost communication?
- Other consensus algorithms can be used for commit, including some that don't depend on such strong timing and reliability assumptions.

Paxos consensus algorithm [Lamport]



- A more robust consensus algorithm, can be used for commit.
- Tolerates process stopping and recovery, message losses and delays,...
- Runs in partially synchronous model.
- Similar to algorithm by [Dwork, Lynch, Stockmeyer].
- Algorithm idea:
 - Processes use an unreliable leader election subalgorithm to choose a coordinator, who tries to achieve consensus.
 - Coordinator decides based on active support from a majority of the processes.
 - Does not assume anything based on **not** receiving a message.
 - Subtleties arise when multiple coordinators are active---must ensure consistency.
- Practical difficulties with fault-tolerance in the synchronous model motivate moving on to study the asynchronous model (start this next time).

A Lower Bound for Commit

- How many messages are needed to solve the commit problem?
- **Theorem [Dwork, Skeen]:** Any algorithm that solves the commit problem, even with weak termination, uses at least $2n - 2$ messages in the failure-free execution α in which all inputs are 1.
- **Note:** That's what 2-phase commit uses, so 2-phase commit is “optimal”.
- Proof considers the communication pattern for α :





Information flow in a communication pattern



- i **affects** j in a pattern if there is a path in the pattern from i at time 0 to j at some time.
- In Pattern 1, all processes affect all processes.
- In Pattern 2, 4 does not affect 1.
- **Lemma:** In the failure-free, all-1-input run α , every i affects every j in the communication pattern of α .
- **Corollary:** The communication pattern of α has at least $2n - 2$ edges.



Proof of the Lemma

- **Lemma:** In the failure-free, all-1-input run α , every i affects every j in the communication pattern of α .
- **Proof:**
 - By contradiction. Suppose i does not affect j (for some particular i, j).
 - Then $i \neq j$.
 - Construct execution α' , which is the same as α except that:
 - i 's input is 0, and
 - Every process that is affected by process i in α fails just after it first gets affected by process i in α .
- **Example:** Process 4 does not affect process 1.



Proof of the Lemma

- **Lemma:** In the failure-free, all-1-input run α , every i affects every j in the communication pattern of α .
- **Proof, cont'd:**
 - Construct execution α' :
 - i 's input is 0, and
 - Every process that is affected by process i in α fails just after it first gets affected by process i in α .
 - In α , all processes eventually decide 1.
 - α' is indistinguishable from α to process j .
 - So process j decides 1 in α' , which contradicts the requirements.



Next time...

- Modeling asynchronous systems
- I/O automata
- **Reading:** Chapter 8