

6.852: Distributed Algorithms

Fall, 2015

Lecture 9


Today's plan

- Distributed commit
- Formal modeling of asynchronous systems:
 - I/O automata
 - Executions and traces
 - Operations: composition, hiding
 - Properties and proof methods:
 - Invariants
 - Simulation relations
- Reading: Section 7.3, Chapter 8
- Next:
 - Asynchronous network algorithms: Leader election, breadth-first search, shortest paths, spanning trees.
 - Reading: Chapters 14 and 15

Distributed Commit

Distributed Commit

- **Motivation:** Distributed database transaction processing
 - A database transaction performs work at several distributed sites.
 - Transaction manager (TM) at each site decides whether it would like to “commit” or “abort” the transaction.
 - Based on whether the transaction’s work has been successfully completed at that site, and results made stable.
 - All TMs must agree on whether to commit or abort.
- **Assume:**
 - Process stopping failures only.
 - n -node, complete, undirected graph.
- **Require:**
 - **Agreement:** No two processes decide differently (faulty or not, uniformity)
 - **Validity:**
 - If any process starts with 0 (abort) then 0 is the only allowed decision.
 - If all start with 1 (commit) and there are no faulty processes then 1 is the only allowed decision.




Correctness Conditions for Commit

- **Agreement:** No two processes decide differently.
- **Validity:**
 - If any process starts with 0 then 0 is the only allowed decision.
 - If all start with 1 and there are no faulty processes then 1 is the only allowed decision.
- Note asymmetry: Guarantee abort (0) if **anyone** wants to abort; guarantee commit (1) if **everyone** wants to commit **and no one fails** (best case).
- **Termination:**
 - **Weak termination:** If there are no failures then all processes eventually decide.
 - **Strong termination (non-blocking condition):** (Even if there are failures), all nonfaulty processes eventually decide.



2-Phase Commit

- Traditional, blocking algorithm (guarantees weak termination only).
- Assumes distinguished process 1, acts as “coordinator” (leader).
- **Round 1:** All send initial values to process 1, who decides.
 - If it sees 0, or doesn’t hear from someone, it decides 0; otherwise it decides 1.
- **Round 2:** Process 1 sends the decision to everyone else.
- **Q:** When can the processes decide?
- Anyone with initial value 0 can decide at the beginning.
- Process 1 decides after receiving round 1 messages.
- Everyone else decides after round 2 (if there are no failures).






Correctness of 2-Phase Commit

- **Agreement:**
 - Because decision is centralized (and consistent with any individual initial decisions).
- **Validity:**
 - Because of how the coordinator decides.
- **Weak termination:**
 - If no one fails, then everyone terminates by the end of round 2.
- **Strong termination?**
 - No: If the coordinator fails before sending its round 2 messages, then others with initial value 1 will never terminate.



Add a termination protocol?

- We might try to add a termination protocol: other processes try to detect failure of coordinator and finish agreeing on their own.
- But this can't always work:
 - If initial values are 0,1,1,1, then by validity, everyone is required to decide 0.
 - If initial values are 1,1,1,1 and process 1 fails just after deciding, and before sending out its round 2 messages, then:
 - Process 1 decides 1.
 - By agreement, others must decide 1.
 - But the other processes can't distinguish these two situations.




Complexity of 2-phase commit

- Time:
 - 2 rounds
- Communication:
 - At most $2n$ messages


3-Phase Commit [Skeen]

- Yields strong termination.
- **Trick:** Introduce intermediate stage, before actually deciding.
- Process states are now classified into four categories:
 - *dec0*: Already decided 0.
 - *dec1*: Already decided 1.
 - *ready*: Ready to decide 1 but hasn't yet.
 - *uncertain*: Otherwise.
- Again, process 1 acts as “coordinator”.
- Communication pattern:



3-Phase Commit

- All processes are initially *uncertain*.
- **Round 1:**
 - All other processes send their initial values to p_1 .
 - All with initial value 0 **decide 0** (and enter *dec0* state)
 - If p_1 receives 1s from everyone and its own initial value is 1, p_1 becomes *ready*, but doesn't yet decide.
 - If p_1 sees 0 or doesn't hear from someone, p_1 **decides 0**.
- **Round 2:**
 - If p_1 has decided 0, it broadcasts "decide 0", else it broadcasts "ready".
 - Anyone else who receives "decide 0" **decides 0**.
 - Anyone else who receives "ready" becomes *ready*.
 - Now p_1 **decides 1** if it hasn't already decided.
- **Round 3:**
 - If p_1 has decided 1, it bcasts "decide 1".
 - Anyone else who receives "decide 1" **decides 1**.



3-Phase Commit

- Key invariants (after 0, 1, 2, or 3 rounds):
 - If any process is in *ready* or *dec1*, then all processes have initial value 1.
 - If any process is in *dec0* then:
 - No process is in *dec1*, and no non-failed process is *ready*.
 - If any process is in *dec1* then:
 - No process is in *dec0*, and no non-failed process is *uncertain*.
- **Proof:** LTTR.
 - Key step: Third condition is preserved when p_1 *decides 1* after round 2.
 - In this case, p_1 knows that:
 - Everyone's input is 1.
 - No one *decided 0* at the end of round 1.
 - Every other process has either become *ready* or has failed (without deciding).
 - Implies the third condition.
- **Note critical use of synchrony here:**
 - p_1 infers that non-failed processes are *ready* just because round 2 is completed.
 - Without synchrony, this would require explicit acknowledgments.

Correctness conditions (so far)

- Agreement and validity follow, for these three rounds.
- Weak termination holds
- Strong termination:
 - Doesn't hold yet---must add a termination protocol.
 - Allow process 2 to act as coordinator, then 3,...
 - “Rotating coordinator” strategy

3-Phase Commit

- **Round 4:**
 - All processes send current status (*dec0*, *uncertain*, *ready*, *dec1*) to p_2 .
 - If p_2 receives any *dec0*'s and hasn't already decided, then p_2 **decides 0**.
 - If p_2 receives any *dec1*'s and hasn't already decided, then p_2 **decides 1**.
 - If all received values, and its own value, are *uncertain*, then p_2 **decides 0**.
 - Otherwise (all values are *uncertain* or *ready* and at least one is *ready*), p_2 becomes *ready*, but doesn't decide yet.
- **Round 5** (analogous to round 2):
 - If p_2 has (ever) decided 0, broadcasts "decide 0", and similarly for 1.
 - Else broadcasts "ready".
 - Any undecided process who receives "decide()" decides accordingly.
 - Any process who receives "ready" becomes *ready*.
 - Now p_2 **decides 1** if it hasn't already decided.
- **Round 6** (analogous to round 3):
 - If p_2 has decided 1, broadcasts "decide 1".
 - Anyone else who receives "decide 1" **decides 1**.
- Continue with subsequent rounds for p_3, p_4, \dots

Correctness

- Key invariants still hold:
 - If any process is in *ready* or *dec1*, then all processes have initial value 1.
 - If any process is in *dec0* then:
 - No process is in *dec1*, and no non-failed process is *ready*.
 - If any process is in *dec1* then:
 - No process is in *dec0*, and no non-failed process is *uncertain*.
- Imply agreement, validity
- Strong termination:
 - Because eventually some coordinator will finish the job (unless everyone fails).

Complexity

- Time until everyone decides:
 - Normal case 3
 - Worst case $3n$
- Messages until everyone decides:
 - Normal case $O(n)$
 - Technicality: When can processes stop sending messages?
 - Worst case $O(n^2)$

Practical issues for 3-phase commit


- Depends on strong assumptions, which may be hard to guarantee in practice:
 - Synchronous model:
 - Could emulate with approximately-synchronized clocks, timeouts.
 - Reliable message delivery:
 - Could emulate with acks and retransmissions.
 - But if retransmissions add too much delay, then we can't emulate the synchronous model accurately.
 - Leads to unbounded delays, asynchronous model.
 - Accurate diagnosis of process failures:
 - Get this “for free” in the synchronous model.
 - E.g., 3-phase commit algorithm lets process that doesn't hear from another process i at a round conclude that i must have failed.
 - Very hard to guarantee in practice: In Internet, or even a LAN, how to reliably distinguish failure of a process from lost communication?
- Other consensus algorithms can be used for commit, including some that don't depend on such strong timing and reliability assumptions.

Paxos consensus algorithm [Lamport]

- A more robust consensus algorithm, can be used for commit.
- Tolerates process stopping and recovery, message losses and delays,...
- Runs in partially synchronous model.
- Similar to algorithm by [Dwork, Lynch, Stockmeyer].
- Algorithm idea:
 - Processes use an unreliable leader election subalgorithm to choose a coordinator, who tries to achieve consensus.
 - Coordinator decides based on active support from a majority of the processes.
 - Does not assume anything based on **not** receiving a message.
 - Subtleties arise when multiple coordinators are active---must ensure consistency.
- Practical difficulties with fault-tolerance in the synchronous model motivate studying the asynchronous model (later today).



A Lower Bound for Commit

- How many messages are needed to solve the commit problem?
- **Theorem [Dwork, Skeen]:** Any algorithm that solves the commit problem, even with weak termination, uses at least $2n - 2$ messages in the failure-free execution α in which all inputs are 1.
- **Note:** That's what 2-phase commit uses, so 2-phase commit is “optimal”:



- Proof considers the communication pattern for α :



Information flow in a communication pattern



- i **affects** j in a pattern if there is a path in the pattern from i at time 0 to j at some later time.
- In Pattern 1, all processes affect all processes.
- In Pattern 2, 4 does not affect 1.
- **Lemma:** In the failure-free, all-1-input run α , every i affects every j in the communication pattern of α .
- **Corollary:** The communication pattern of α has at least $2n - 2$ edges.



Proof of the Lemma

- **Lemma:** In the failure-free, all-1-input run α , every i affects every j in the communication pattern of α .
- **Proof:**
 - By contradiction. Suppose i does not affect j (for some particular i, j).
 - Then $i \neq j$.
 - Construct execution α' , which is the same as α except that:
 - i 's input is 0, and
 - Every process that is affected by process i in α fails just after it first gets affected by process i in α .
- **Example:** Process 4 does not affect process 1.



Proof of the Lemma



- **Lemma:** In the failure-free, all-1-input run α , every i affects every j in the communication pattern of α .
- **Proof, cont'd:**
 - Construct execution α' :
 - i 's input is 0, and
 - Every process that is affected by process i in α fails just after it first gets affected by process i in α .
 - In α , all processes eventually decide 1.
 - α' is indistinguishable from α to process j .
 - So process j decides 1 in α' , which contradicts the requirements.



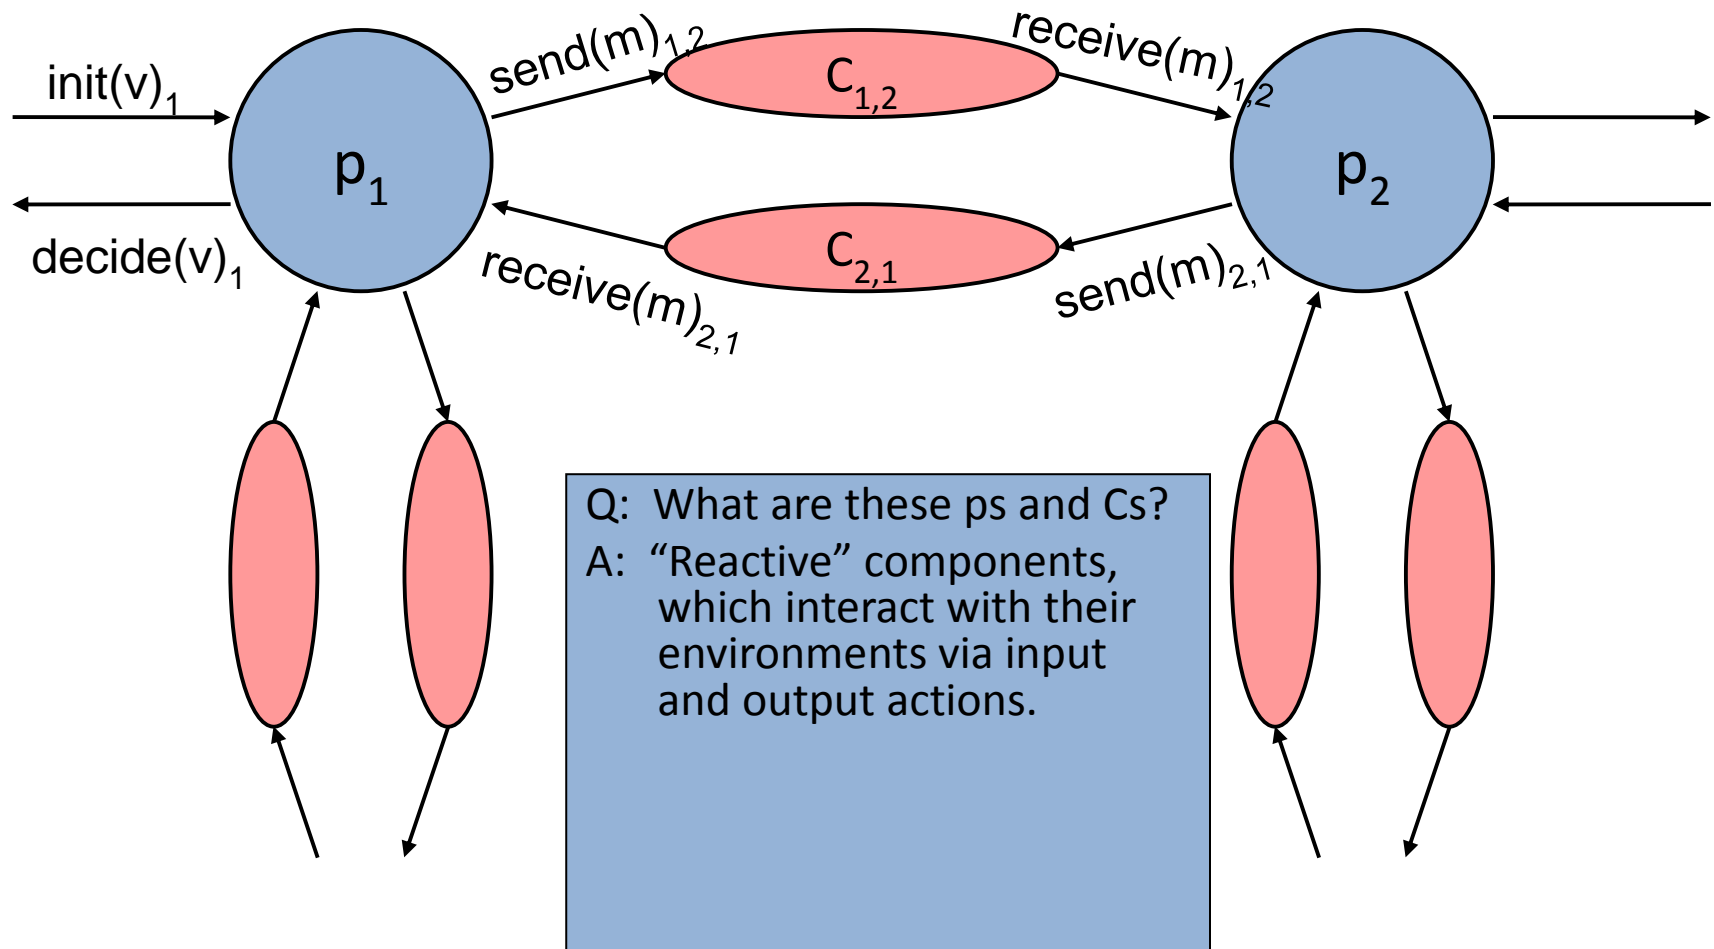
Asynchronous Systems

Asynchronous systems


- No timing assumptions
 - No rounds
- Two kinds of asynchronous models:
 - Asynchronous networks
 - Processes communicating via channels
 - Asynchronous shared-memory systems
 - Processes communicating via shared objects



Asynchronous network: Processes and channels



Asynchronous shared-memory system: Processes and objects



These processes and objects are also “reactive” components.

In both cases, we have reactive components.

We need a general model for reactive components.

Specifying problems and systems

- Processes, channels, and objects are **automata**
 - Perform **actions** while changing state.
 - Reactive
 - Interact with environment via input and output actions.
 - Not just functions from input values to output values; they may have more kinds of interactions.
- **Execution:**
 - Sequence of states and actions
 - Interleaving semantics
- **External behavior (trace):**
 - We observe **external** actions.
 - States and internal actions are hidden.
 - **Problems** are defined in terms of allowable traces.


Input/Output Automata

Input/Output Automata

- General **mathematical modeling framework** for reactive system components.
 - Little structure---must add special structure to specialize it for networks, shared-memory systems,...
- Designed for describing systems in a **modular** way:
 - Supports description of individual system components, and how they **compose** to yield a larger system.
 - Supports description of systems at different **levels of abstraction**, e.g.:
 - Detailed implementation vs. higher-level algorithm description.
 - Optimized algorithm vs. simpler, un-optimized version.
- Supports several standard **proof techniques**:
 - **Invariants**
 - **Simulation relations** (like running 2 algorithms side-by-side and relating their behavior step-by-step).
 - **Compositional reasoning** (prove properties of individual components; use compositional reasoning to infer properties for the overall system).

Input/Output Automaton

- State transition system
 - Transitions labeled by actions
- Actions classified as input, output, internal
 - Input, output are **external**.
 - Output, internal are **locally controlled**.




Input/Output Automaton, formally

- *sig* = (*in*, *out*, *int*)
 - input, output, internal actions (disjoint)
 - $in \cup out \cup int$
 - $ext = in \cup out$
 - $local = out \cup int$
- *states*: Not necessarily finite
- *start* $\subseteq states$
- *trans* $\subseteq states \times acts \times states$
 - Input-enabled: Any input “enabled” in any state.
- *tasks*, partition of locally controlled actions
 - Used to specify liveness.

Remarks


- A **step** of an automaton is an element of *trans*.
- Action π is **enabled** in a state s if *trans* contains a step (s, π, s') for some s' .
- I/O automata must be **input-enabled**.
 - Every input action is enabled in every state.
 - Captures the idea that an automaton cannot control its inputs.
 - If we want restrictions, model the environment as another automaton and express restrictions in terms of the environment.
 - Could allow a component to detect bad inputs and halt, or exhibit unconstrained behavior for bad inputs.
- Tasks correspond to “threads of control”.
 - Used to define fairness (give turns to all tasks).
 - Needed to guarantee liveness properties (e.g., the system keeps making progress, or eventually terminates).

Example: Channel automaton




- Reliable unidirectional FIFO channel between two processes.
 - Fix message alphabet M .
- signature
 - input actions: *send(m)*, $m \in M$
 - output actions: *receive(m)*, $m \in M$
 - No internal actions
- states
 - *queue*: FIFO queue of M , initially empty

Channel automaton



- trans
 - *send(m)*
 - effect: add m to (end of) *queue*
 - *receive(m)*
 - precondition: m is at head of *queue*
 - effect: remove head of *queue*
- tasks
 - All *receive* actions in one task.

Channel automaton



- trans

- *send(m)_{i,j}*

- effect: add m to (end of) *queue*


- *receive(m)_{i,j}*

- precondition: m is at head of *queue*
 - effect: remove head of *queue*

- tasks

- All *receive* actions in one task.

A process



- E.g., in a consensus protocol.
- See book, p. 205, for code details.
- Inputs arrive from the outside.
- Process sends/receives values, collects vector of values, one for each process.
- When vector is filled, outputs a decision obtained as a function of the vector.
- Can get new inputs, change values, send and output repeatedly.
- Tasks for:
 - Sending to each individual neighbor.
 - Outputting decisions.

Executions

- An I/O automaton executes as follows:
 - Start at some start state.
 - Repeatedly take step from current state to new state.
- Formally, an **execution** is a finite or infinite sequence:
 - $s_0 \pi_1 s_1 \pi_2 s_2 \pi_3 s_3 \pi_4 s_4 \pi_5 s_5 \dots$ (if finite, ends in state)
 - s_0 is a start state
 - $(s_i, \pi_{i+1}, s_{i+1})$ is a step (i.e., in trans)

$\lambda, send(a), a, send(b), ab, receive(a), b, receive(b), \lambda$

Execution fragments

- An I/O automaton executes as follows:
 - Start at some start state.
 - Repeatedly take step from current state to new state.
- Formally, an execution is a sequence:
 - $s_0 \pi_1 s_1 \pi_2 s_2 \pi_3 s_3 \pi_4 s_4 \pi_5 s_5 \dots$
 - s_0 is a start state
 - $(s_i, \pi_{i+1}, s_{i+1})$ is a step.

execution fragment

Invariants and reachable states

- A state is **reachable** if it appears in some execution.
 - Equivalently, at the end of some finite execution.
- An **invariant** is a predicate that is true for every reachable state.
 - Most important tool for proving properties of concurrent/distributed algorithms.
 - Typically proved by induction on length of execution.

Traces

- Traces allow us to focus on components' external behavior.
- Useful for defining correctness.
- A **trace** of an execution is the subsequence of external actions in the execution.
 - No states, no internal actions.
 - Denoted $trace(\alpha)$, where α is an execution.
 - Models “observable behavior”.

$\lambda, send(a), a, send(b), ab, receive(a), b, receive(b), \lambda$

$send(a), send(b), receive(a), receive(b)$

Operations on I/O Automata


Operations on I/O automata

- To describe how systems are built out of components, the model has operations for **composition, hiding, renaming**.
- **Composition:**
 - “Put multiple automata together.”
 - Output actions of one may be input actions of others.
 - All components having an action perform steps involving that action together (“synchronize on actions”).
- Composing finitely many (or countably many) automata $A_i, i \in I$:
- Need compatibility conditions:
 - Internal actions aren’t shared:
 - $int(A_i) \cap acts(A_j) = \emptyset$
 - Only one automaton controls each output:
 - $out(A_i) \cap out(A_j) = \emptyset$
 - But output of one automaton can be an input of one or more others.
 - No action is shared by infinitely many A_i s.

Composition of compatible automata

- For two automata A and B (see book for general case).
- $out(A \times B) = out(A) \cup out(B)$
- $int(A \times B) = int(A) \cup int(B)$
- $in(A \times B) = in(A) \cup in(B) - (out(A) \cup out(B))$
- $states(A \times B) = states(A) \times states(B)$
- $start(A \times B) = start(A) \times start(B)$
- $trans(A \times B)$: includes (s, π, s') iff
 - $(s_A, \pi, s'_A) \in trans(A)$ if $\pi \in acts(A)$; $s_A = s'_A$ otherwise.
 - $(s_B, \pi, s'_B) \in trans(B)$ if $\pi \in acts(B)$; $s_B = s'_B$ otherwise.
- $tasks(A \times B) = tasks(A) \cup tasks(B)$
- Notation: $\prod_{i \in I} A_i$, for composition of A_i , $i \in I$.

Composition of channels and consensus processes



Composition: Basic results

- Projection
 - Execution of composition “looks good” to each component.
- Pasting
 - If execution “looks good” to each component, it is good overall.
- Substitutivity
 - Can replace a component with one that implements it.

Composition: Basic results

Theorem 1: Projection

- If $\alpha \in \text{execs}(\Pi A_i)$ then $\alpha|A_i \in \text{execs}(A_i)$ for every i .
- If $\beta \in \text{traces}(\Pi A_i)$ then $\beta|A_i \in \text{traces}(A_i)$ for every i .

Composition: Basic results

Theorem 2: Pasting

Suppose β is a sequence of external actions of ΠA_i .

- If $\alpha_i \in \text{execs}(A_i)$ and $\beta|A_i = \text{trace}(\alpha_i)$ for every i , then there is an execution α of ΠA_i such that $\beta = \text{trace}(\alpha)$ and $\alpha_i = \alpha|A_i$ for every i .
- If $\beta|A_i \in \text{traces}(A_i)$ for every i then $\beta \in \text{traces}(\Pi A_i)$.

Composition: Basic results

Theorem 3: Substitutivity

- Suppose A_i and A'_i have the same external signature, and $traces(A_i) \subseteq traces(A'_i)$ for every i .
 - A kind of “implementation” relationship.
- Then $traces(\Pi A_i) \subseteq traces(\Pi A'_i)$ (assuming compatibility).

Proof:

- Follows from trace pasting and projection, Theorems 1 and 2.

Other operations on I/O automata

- Hiding
 - Reclassify some output actions as internal.
 - Hides internal communication among components of a system.
- Renaming
 - Change names of some actions.
 - Action names are important for specifying component interactions.
 - E.g., define a “generic” automaton, then rename actions to define many instances to use in a system.
 - As we did with channel automata.

Fairness

Fairness

- Task T (a set of actions) corresponds to a “thread of control”.
- Used to define “fair” executions: a task that is continuously enabled eventually takes a step.
- Tasks are used to state and prove liveness properties, e.g., that something eventually happens, like an algorithm terminating.
- Formally, an execution (or fragment) α of A is **fair** to task T if one of the following holds:
 - α is finite and T is not enabled in the final state of α .
 - α is infinite and contains infinitely many events in T .
 - α is infinite and contains infinitely many states in which T is not enabled.
- Execution of A is **fair** if it is fair to all tasks of A .
- Trace of A is **fair** if it is the trace of a fair execution of A .

Example

- Channel
 - Only one task (all receive actions).
 - A finite execution of Channel is fair iff *queue* is empty at the end.
 - **Q:** Is every infinite execution of Channel fair?
- Consensus process
 - Separate tasks for sending to each other process, and for output.
 - Means it “keeps trying” to do these forever.

Fairness and composition

- Fairness “behaves nicely” with respect to composition---results analogous to non-fair results:

- **Theorem 4: Projection**

- If $\alpha \in \text{fairexecs}(\Pi A_i)$ then $\alpha|A_i \in \text{fairexecs}(A_i)$ for every i .
- If $\beta \in \text{fairtraces}(\Pi A_i)$ then $\beta|A_i \in \text{fairtraces}(A_i)$ for every i .

- **Theorem 5: Pasting**

Suppose β is a sequence of external actions of ΠA_i .

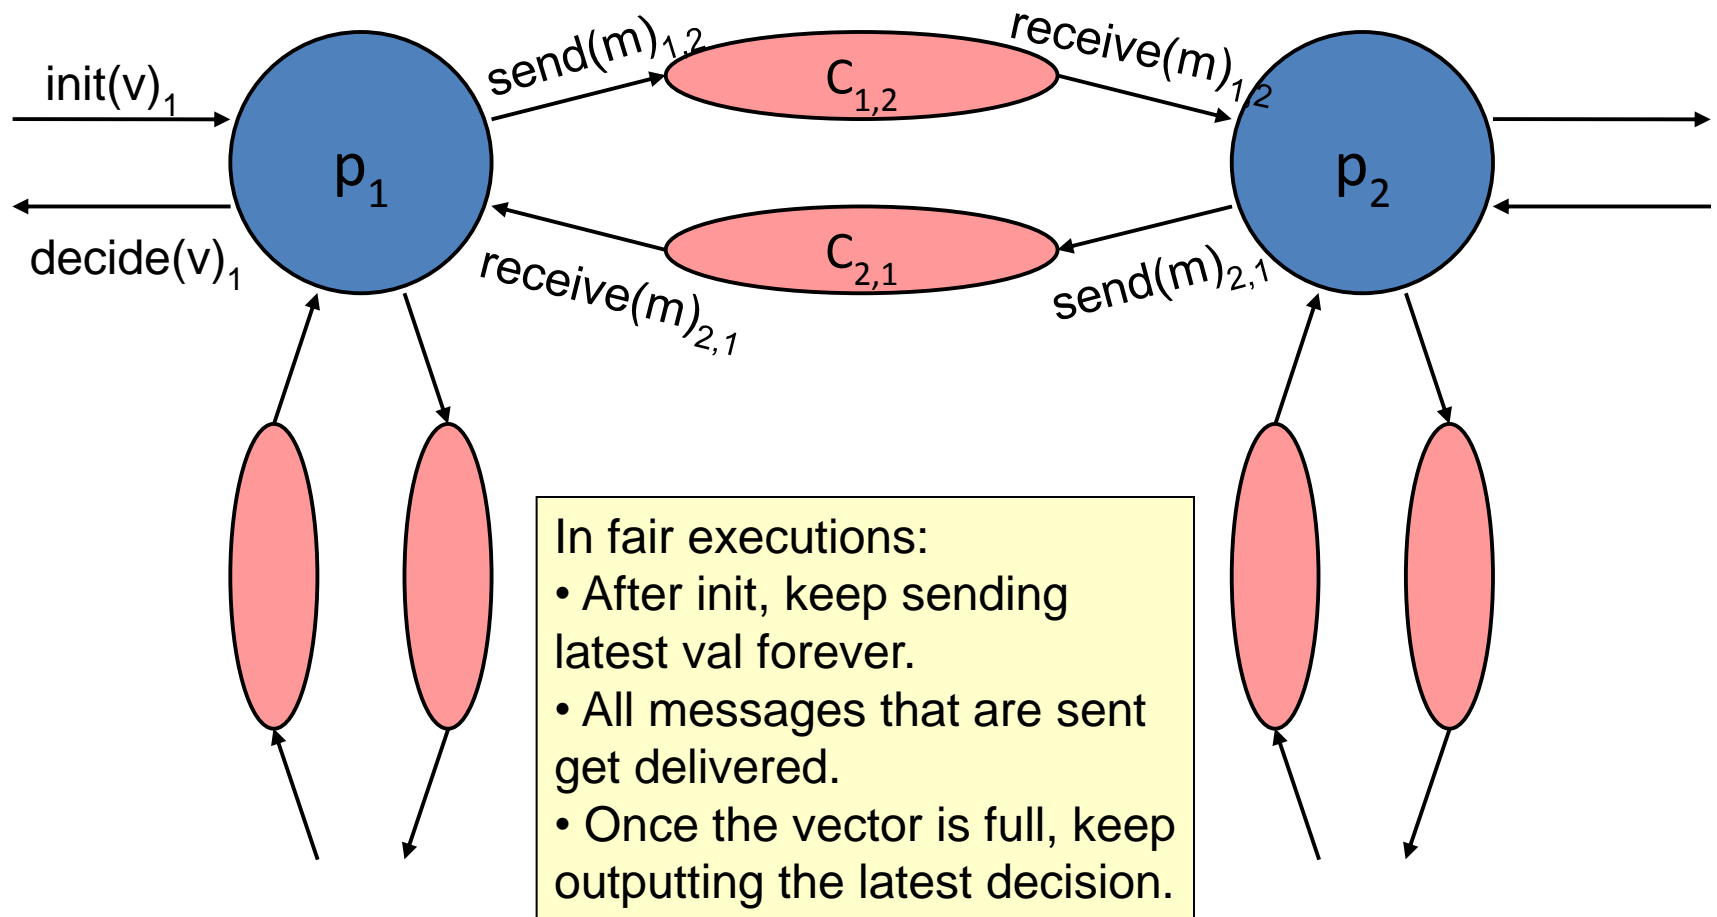
- If $\alpha_i \in \text{fairexecs}(A_i)$ and $\beta|A_i = \text{trace}(\alpha_i)$ for every i , then there is a **fair** execution α of ΠA_i such that $\beta = \text{trace}(\alpha)$ and $\alpha_i = \alpha|A_i$ for every i .
- If $\beta|A_i \in \text{fairtraces}(A_i)$ for every i then $\beta \in \text{fairtraces}(\Pi A_i)$.

Fairness and composition

- Theorem 6: Substitutivity

- Suppose A_i and A'_i have the same external signature, and $\text{fairtraces}(A_i) \subseteq \text{fairtraces}(A'_i)$ for every i .
 - Another kind of “implementation” relationship.
- Then $\text{fairtraces}(\Pi A_i) \subseteq \text{fairtraces}(\Pi A'_i)$.

Composition of channels and consensus processes



Properties and Proof Methods

- Compositional reasoning
- Invariants
- Trace properties
- Simulation relations

Compositional reasoning

- Use Theorems 1-6 to infer properties of a system from properties of its components.
- And vice versa.

Invariants

- A state is **reachable** if it appears in some execution (or, at the end of some finite execution).
- An **invariant** is a predicate that is true for every reachable state.
- Most important tool for proving properties of concurrent and distributed algorithms.
- Proving invariants:
 - Typically, by induction on length of execution.
 - Often prove batches of inter-dependent invariants together.
 - Step granularity is finer than round granularity, so proofs are more complicated and detailed than those for synchronous algorithms.

Example: Incrementing

- Two processes, P_1 and P_2 , communicating via channels C_{12} and C_{21} :
send(v)₁₂, receive(v)₁₂, send(v)₂₁, receive(v)₂₁.
- Each process has a local variable *val*.
- Initially $P_1.val = 1$, $P_2.val = 2$.
- Transitions:
 - *send(v)*, where $v = val$, at any time.
 - When *receive(v)*: $val := v + 1$.
- **Invariant 1:** $P_1.val$ is odd and $P_2.val$ is even
- **Proof:** By induction.
 - Base: Yes
 - Inductive step:
 - Cases based on various kinds of send/receive actions.
 - Strengthen invariant?
 - Add that any value in C_{12} is odd, and any value in C_{21} is even.

Example: Incrementing

- Initially $P_1.val = 1$, $P_2.val = 2$.
- Transitions:
 - $send(v)$, where $v = val$, at any time.
 - When $receive(v)$: $val := v + 1$.
- Invariant 1:** $P_1.val$ is odd and $P_2.val$ is even
- Invariant 2:** $|P_2.val - P_1.val| \leq 1$
- Proof:** By induction.
 - Base: Yes
 - Inductive step:
 - Cases based on various send/receive actions.
 - Strengthen invariant?
 - LTTR.

Trace properties

- A trace property is essentially a set of allowable external behavior sequences.
- Formally, a **trace property** P is a pair consisting of:
 - $sig(P)$: External signature (no internal actions).
 - $traces(P)$: Set of sequences of actions in $sig(P)$.
- Automaton A **satisfies** trace property P if (two different, alternative notions, depending on whether we want to consider fairness):
 - $extsig(A) = sig(P)$ and $traces(A) \subseteq traces(P)$
 - $extsig(A) = sig(P)$ and $fairtraces(A) \subseteq traces(P)$

Safety and liveness


- **Safety property:** “Bad” thing doesn't happen:
 - Nonempty (null trace is always safe).
 - Prefix-closed: Every prefix of a safe trace is safe.
 - Limit-closed: Limit of sequence of safe traces is safe.
- **Liveness property:** “Good” thing happens eventually:
 - Every finite sequence over $acts(P)$ can be extended to a sequence in $traces(P)$.
 - “It's never too late.”
- Define safety/liveness for executions similarly.

Automata as specifications

- Every I/O automaton specifies a trace property ($extsig(A), traces(A)$).
- So we can use an automaton as a problem specification.
- Automaton A “implements” automaton B if
 - $extsig(A) = extsig(B)$
 - $traces(A) \subseteq traces(B)$


Hierarchical proofs

- Important strategy for proving correctness of complex asynchronous distributed algorithms.
- Define a series of automata, each implementing the previous one (“successive refinement”).
- Highest-level automaton model captures the “real” problem specification.
- Next level is a high-level algorithm description.
- Successive levels represent more and more detailed versions of the algorithm.
- Lowest level is the full algorithm description.




Hierarchical proofs

- For example:
 - High levels centralized, lower levels distributed.
 - High levels inefficient but simple, lower levels optimized and more complex.
 - High levels with large granularity steps, lower levels with finer granularity steps.
- In all these cases, lower levels are harder to understand and reason about.
- So instead of reasoning about them directly, relate them to higher-level descriptions.
- Method similar to what we saw for synchronous algorithms.



Hierarchical proofs


- Recall, for synchronous algorithms:
 - Optimized algorithm runs side-by-side with unoptimized version, and “invariant” proved to relate the states of the two algorithms.
 - Prove using induction.
- For asynchronous systems, it’s harder:
 - Asynchronous model has **more nondeterminism** (in choice of new state, in order of steps).
 - So, it’s harder to determine which executions to compare.
- **One-way implementation relationship is enough:**
 - For each execution of the lower-level algorithm, there is a corresponding execution of the higher-level algorithm.
 - “Everything the algorithm does is allowed by the spec.”
 - Don’t need the other direction: it doesn’t matter if the algorithm does **everything** that is allowed.



Simulation relations

- Most common method of proving that one automaton implements another.
- Assume A and B have the same *extsig*, and R is a binary relation from $states(A)$ to $states(B)$.
- Then R is a **simulation relation** from A to B provided:
 - $s_A \in start(A)$ implies that there exists $s_B \in start(B)$ such that $s_A R s_B$.
 - If s_A, s_B are reachable states of A and B respectively, $s_A R s_B$ and (s_A, π, s'_A) is a step of A , then there is an execution fragment β of B , starting with s_B and ending with s'_B such that $s'_A R s'_B$ and $trace(\beta) = trace(\pi)$.


Simulation relations



- R is a **simulation relation** from A to B provided:
 - $s_A \in start(A)$ implies that there exists $s_B \in start(B)$ such that $s_A R s_B$.
 - If s_A, s_B are reachable states of A and B , $s_A R s_B$ and (s_A, π, s'_A) is a step, then there is an execution fragment β starting with s_B and ending with s'_B such that $s'_A R s'_B$ and $trace(\beta) = trace(\pi)$.


Simulation relations

- **Theorem:** If there is a simulation relation from A to B then $traces(A) \subseteq traces(B)$.
- All traces of A , not just finite traces.
- **Proof:** Fix a trace of A , arising from a (possibly infinite) execution of A .
- Create a corresponding execution of B , using an iterative construction.




Simulation relations

- **Theorem:** If there is a simulation relation from A to B then $traces(A) \subseteq traces(B)$.




Simulation relations

- **Theorem:** If there is a simulation relation from A to B then $traces(A) \subseteq traces(B)$.




Simulation relations

- **Theorem:** If there is a simulation relation from A to B then $traces(A) \subseteq traces(B)$.




Example: Channels

- Show two channels implement one.




- Rename some actions.
- Let $D = \text{hide}_{\{pass(m)\}} A \times B$.
- Show that $\text{traces}(D) \subseteq \text{traces}(C)$.

Recall: Channel automaton




- Reliable unidirectional FIFO channel.
- *sig*
 - Input actions: *send(m)*, $m \in M$
 - output actions: *receive(m)*, $m \in M$
 - No internal actions
- *states*
 - *queue*: FIFO queue of M , initially empty

Channel automaton




- *trans*
 - *send(m)*
 - effect: add m to *queue*
 - *receive(m)*
 - precondition: $m = \text{head}(\text{queue})$
 - effect: remove head of *queue*
- *tasks*
 - All *receive* actions in one task

Composing two channel automata



- Output of B is input of A
 - Rename *receive(m)* of B and *send(m)* of A to *pass(m)*.
- Claim $D = \text{hide}_{\{pass(m)\}} A \times B$ implements C .
- Define relation R :
 - For $s \in \text{states}(D)$ and $u \in \text{states}(C)$, define $s R u$ iff $u.queue$ is the concatenation of $s.A.queue$ and $s.B.queue$.
- Proof that R is a simulation relation:
 - Start condition: All queues are empty, so start states correspond.
 - Step condition: Define “step correspondence”:


Composing two channel automata



$s R u$ iff $u.queue$ is concatenation of $s.A.queue$ and $s.B.queue$

- Step correspondence:
 - For each step $(s, \pi, s') \in trans(D)$ and u such that $s R u$, define execution fragment β of C :
 - Starts with u , ends with u' such that $s' R u'$.
 - $trace(\beta) = trace(\pi)$
 - Here, actions in β depend only on π , and uniquely determine the states.
 - Same action if external, empty sequence if internal.

Composing two channel automata



$s \ R \ u$ iff $u.queue$ is concatenation of $s.A.queue$ and $s.B.queue$

- Step correspondence:
 - $\pi = \textit{send}(m)$ in D corresponds to $\textit{send}(m)$ in C
 - $\pi = \textit{receive}(m)$ in D corresponds to $\textit{receive}(m)$ in C
 - $\pi = \textit{pass}(m)$ in D corresponds to λ in C
- Verify that this works:
 - Same external actions (yes).
 - Actions of C are enabled.
 - Final states related by relation R .
- Routine case analysis:

Showing R is a simulation relation


$s R u$ iff $u.queue$ is concatenation of $s.A.queue$ and $s.B.queue$

- Case 1: $\pi = \text{send}(m)$
 - No enabling issues (input).
 - Must check that $s' R u'$.
 - Since $s R u$, $u.queue$ is the concatenation of $s.A.queue$ and $s.B.queue$.
 - Adding the same m to the end of $u.queue$ and $s.B.queue$ maintains the correspondence.
- Case 2: $\pi = \text{receive}(m)$
 - Enabling: Check that $\text{receive}(m)$, for the same m , is also enabled in u .
 - We know that m is first on $s.A.queue$.
 - Since $s R u$, m is also first on $u.queue$.
 - So $\text{receive}(m)$ is enabled in u .
 - $s' R u'$: Since m is removed from both $s.A.queue$ and $u.queue$.

Showing R is a simulation relation

$s R u$ iff $u.queue$ is concatenation of $s.A.queue$ and $s.B.queue$

- Case 3: $\pi = \text{pass}(m)$
 - No enabling issues (since no high-level steps are involved).
 - Must check $s' R u$:
 - Since $s R u$, $u.queue$ is the concatenation of $s.A.queue$ and $s.B.queue$.
 - The concatenation of the queues is unchanged as a result of this step, so also $u.queue$ is the concatenation of $s'.A.queue$ and $s'.B.queue$.



Next lecture

- A bit more on safety and liveness properties.
- Then, basic asynchronous network algorithms:
 - Leader election
 - Breadth-first search
 - Shortest paths
 - Spanning trees.
- Reading:
 - Chapters 14 and 15