

6.852: Distributed Algorithms

Fall, 2015

Lecture 10


Today's plan

- I/O Automata, cont'd
- Asynchronous network model
- Asynchronous network algorithms:
 - Leader election
 - Constructing a spanning tree
- Readings:
 - Chapter 8
 - Chapter 14
 - Section 15.1-15.3
- Next:
 - Breadth-first search
 - Shortest paths
 - Minimum spanning trees
 - Readings:
 - Section 15.3-15.5
 - [Gallager, Humblet, Spira]

Input/Output automaton

- *sig* = (*in*, *out*, *int*)
 - *act* = *in* \cup *out* \cup *int*
 - *ext* = *in* \cup *out*
 - *local* = *out* \cup *int*
- *states*
- *start* \subseteq *states*
- *trans* \subseteq *states* \times *acts* \times *states*
- *tasks*, partition of locally controlled actions
- Action π is **enabled** in a state *s* if *trans* contains a step (*s*, π , *s'*) for some *s'*.
- I/O automata are **input-enabled**.

Channel automaton



- signature
 - input actions: *send(m)*, $m \in M$
 - output actions: *receive(m)*, $m \in M$
- states
 - *queue*: FIFO queue of M , initially empty
- trans
 - *send(m)*
 - effect: add m to (end of) *queue*
 - *receive(m)*
 - precondition: m is at head of *queue*
 - effect: remove head of *queue*
- tasks
 - All *receive* actions in one task.

Executions

- An **execution** of an I/O automaton is a finite or infinite sequence:
 - $s_0 \pi_1 s_1 \pi_2 s_2 \pi_3 s_3 \pi_4 s_4 \pi_5 s_5 \dots$ (if finite, ends in state)
 - s_0 is a start state
 - $(s_i, \pi_{i+1}, s_{i+1})$ is a step (i.e., in trans)
- **Execution fragment:** Same, but might not begin in a start state.
- The **trace** of an execution is the subsequence of external actions in the execution.
- A **trace** of an I/O automaton is the trace of any execution of the automaton.

Properties and Proof Methods

- Compositional reasoning
- Invariants
- Trace properties
- Simulation relations

Compositional reasoning

- Use Theorems 1-6 (projection, pasting, substitutivity) to infer properties of a system from properties of its components.
- And vice versa.

Invariants

- A state is **reachable** if it appears in some execution (or, at the end of some finite execution).
- An **invariant** is a predicate that is true for every reachable state.
- Most important tool for proving properties of concurrent and distributed algorithms.
- Proving invariants:
 - Typically, by induction on length of execution.
 - Often prove batches of interdependent invariants together.
 - Step granularity is finer than round granularity, so proofs are more complicated and detailed than those for synchronous algorithms.

Example: Incrementing

- Two processes, P_1 and P_2 , communicating via channels C_{12} and C_{21} :
send(v)₁₂, receive(v)₁₂, send(v)₂₁, receive(v)₂₁.
- Each process has a local variable *val*.
- Initially $P_1.val = 1$, $P_2.val = 2$.
- Transitions:
 - *send(v)*, where $v = val$, at any time.
 - When *receive(v)*: $val := v + 1$.
- **Invariant 1:** $P_1.val$ is odd and $P_2.val$ is even
- **Proof:** By induction.
 - Base: Yes
 - Inductive step:
 - Cases based on various kinds of send/receive actions.
 - Strengthen invariant?
 - Add that any value in C_{12} is odd, and any value in C_{21} is even.

Example: Incrementing

- Initially $P_1.val = 1, P_2.val = 2$.
- Transitions:
 - $send(v)$, where $v = val$, at any time.
 - When $receive(v): val := v + 1$.
- Invariant 1:** $P_1.val$ is odd and $P_2.val$ is even
- Invariant 2:** $|P_2.val - P_1.val| \leq 1$
- Proof:** By induction.
 - Base: Yes
 - Inductive step:
 - Cases based on various send/receive actions.
 - Strengthen invariant?
 - LTTR.

Trace properties

- A trace property is essentially a set of allowable external behavior sequences.
- Formally, a **trace property** P is a pair consisting of:
 - $sig(P)$: External signature (no internal actions).
 - $traces(P)$: Set of sequences of actions in $sig(P)$.
- Automaton A **satisfies** trace property P if $extsig(A) = sig(P)$ and (two notions, depending on whether we consider fairness):
 - $traces(A) \subseteq traces(P)$
 - $fairtraces(A) \subseteq traces(P)$
- All problems we consider for asynchronous systems can be formulated as trace properties.
- When we care about liveness, we use the second def.

Safety and liveness properties

- **Safety property:** “Bad” thing doesn't happen:
 - Nonempty (null trace is always safe).
 - Prefix-closed: Every prefix of a safe trace is safe.
 - Limit-closed: Limit of sequence of safe traces is safe.
- **Liveness property:** “Good” thing happens eventually:
 - Every finite sequence over $acts(P)$ can be extended to a sequence in $traces(P)$.
 - “It's never too late.”
- Define safety/liveness for executions similarly.

Safety property S

- $\text{traces}(S)$ are nonempty, prefix-closed, limit-closed.
- **Examples:**
 - Consensus: Agreement, validity
 - Describe as set of sequences of init and decide actions in which we never disagree, or never violate validity.
 - Graph algorithms: Correct shortest paths, correct MSTs,...
 - Outputs do not yield any incorrect answers.
 - Mutual exclusion: No two grants without intervening return.

Proving a safety property

- That is, prove that all traces of A satisfy S.
- By limit-closure, it's enough to prove that all **finite** traces satisfy S.
- Use invariants:
 - Find an invariant corresponding to the trace safety property.
 - Example: Consensus
 - Record decisions in the state.
 - Express agreement and validity in terms of recorded decisions.
 - Then prove the invariant by induction.

Liveness property L

- Every finite sequence over $\text{sig}(L)$ has an extension in $\text{traces}(L)$.
- **Examples:**
 - Termination: No matter where we are, we could still terminate in the future.
 - Some event happens infinitely often.
- Proving liveness properties:
 - Measure progress toward goals, using progress functions.
 - Intermediate milestones.
 - Formal logical reasoning using temporal logic.
 - Methods are less agreed-upon than those for safety properties.

Safety and liveness


- **Theorem:** Every trace property can be expressed as the intersection of a safety property and a liveness property.
- So, to define a problem to be solved by an asynchronous system, it's enough to specify safety requirements and liveness requirements separately.
- This explains why typical specifications of problems for asynchronous systems consist of:
 - A list of safety properties.
 - A list of liveness properties.
 - And nothing else.

Automata as specifications

- Every I/O automaton specifies a trace property ($extsig(A), traces(A)$).
- So we can use an automaton as a problem specification.
- Automaton A “implements” automaton B if
 - $extsig(A) = extsig(B)$
 - $traces(A) \subseteq traces(B)$


Hierarchical proofs

- Important strategy for proving correctness of complex asynchronous distributed algorithms.
- Define a series of automata, each implementing the previous one (“successive refinement”).
- Highest-level automaton model captures the “real” problem specification.
- Next level is a high-level algorithm description.
- Successive levels represent more and more detailed versions of the algorithm.
- Lowest level is the full algorithm description.




Hierarchical proofs

- For example:
 - High levels centralized, lower levels distributed.
 - High levels inefficient but simple, lower levels optimized and more complex.
 - High levels with large granularity steps, lower levels with finer granularity steps.
- In all these cases, lower levels are harder to understand and reason about.
- So instead of reasoning about them directly, relate them to higher-level descriptions.
- Method similar to what we saw for synchronous algorithms.



Hierarchical proofs


- Recall, for synchronous algorithms:
 - Optimized algorithm runs side-by-side with unoptimized version, and “invariant” proved to relate the states of the two algorithms.
 - Prove using induction.
- For asynchronous systems, it’s harder:
 - Asynchronous model has **more nondeterminism** (in choice of new state, in order of steps).
 - So, it’s harder to determine which executions to compare.
- **One-way implementation relationship is enough:**
 - For each execution of the lower-level algorithm, there is a corresponding execution of the higher-level algorithm.
 - “Everything the algorithm does is allowed by the spec.”
 - Don’t need the other direction: it doesn’t matter if the algorithm does **everything** that is allowed.



Simulation relations

- Most common method of proving that one automaton implements another.
- Assume A and B have the same *extsig*, and R is a binary relation from $states(A)$ to $states(B)$.
- Then R is a **simulation relation** from A to B provided:
 - $s_A \in start(A)$ implies that there exists $s_B \in start(B)$ such that $s_A R s_B$.
 - If s_A, s_B are reachable states of A and B respectively, $s_A R s_B$ and (s_A, π, s'_A) is a step of A , then there is an execution fragment β of B , starting with s_B and ending with s'_B such that $s'_A R s'_B$ and $trace(\beta) = trace(\pi)$.


Simulation relations



- R is a **simulation relation** from A to B provided:
 - $s_A \in start(A)$ implies that there exists $s_B \in start(B)$ such that $s_A R s_B$.
 - If s_A, s_B are reachable states of A and B , $s_A R s_B$ and (s_A, π, s'_A) is a step, then there is an execution fragment β starting with s_B and ending with s'_B such that $s'_A R s'_B$ and $trace(\beta) = trace(\pi)$.


Simulation relations

- **Theorem:** If there is a simulation relation from A to B then $traces(A) \subseteq traces(B)$.
- All traces of A , not just finite traces.
- **Proof:** Fix a trace of A , arising from a (possibly infinite) execution of A .
- Create a corresponding execution of B , using an iterative construction.




Simulation relations

- **Theorem:** If there is a simulation relation from A to B then $traces(A) \subseteq traces(B)$.




Simulation relations

- **Theorem:** If there is a simulation relation from A to B then $traces(A) \subseteq traces(B)$.




Simulation relations

- **Theorem:** If there is a simulation relation from A to B then $traces(A) \subseteq traces(B)$.




Example: Channels

- Show two channels implement one.




- Rename *receive(m)* of *B* and *send(m)* of *A* to *pass(m)*.
- Let $D = \text{hide}_{\{pass(m)\}} (A \times B)$.
- Show that $\text{traces}(D) \subseteq \text{traces}(C)$.

Two channels implement one



- Let $D = \text{hide}_{\{pass(m)\}} (A \times B)$.
- Show that $\text{traces}(D) \subseteq \text{traces}(C)$.
- Define relation R : For $s \in \text{states}(D)$ and $u \in \text{states}(C)$, define:
 - $s R u$ iff $u.\text{queue}$ is the concatenation of $s.A.\text{queue}$ and $s.B.\text{queue}$.
- Prove that R is a simulation relation:
 - Start condition: All queues are empty, so start states correspond.
 - Step condition: Define “step correspondence”:


Two channels implement one



$s R u$ iff $u.queue$ is concatenation of $s.A.queue$ and $s.B.queue$

- Step correspondence:
 - For each step $(s, \pi, s') \in trans(D)$ and u such that $s R u$, define execution fragment β of C :
 - Starts with u , ends with u' such that $s' R u'$.
 - $trace(\beta) = trace(\pi)$
 - Here, actions in β depend only on π , and uniquely determine the states.
 - Same action if external, empty sequence if internal.

Two channels implement one



$s \ R \ u$ iff $u.queue$ is concatenation of $s.A.queue$ and $s.B.queue$

- Step correspondence:
 - $\pi = \textit{send}(m)$ in D corresponds to $\textit{send}(m)$ in C
 - $\pi = \textit{receive}(m)$ in D corresponds to $\textit{receive}(m)$ in C
 - $\pi = \textit{pass}(m)$ in D corresponds to λ in C
- Verify that this works:
 - Same external actions (yes).
 - Actions of C are enabled.
 - Final states related by relation R .
- Routine case analysis:

Showing R is a simulation relation


$s R u$ iff $u.queue$ is concatenation of $s.A.queue$ and $s.B.queue$

- Case 1: $\pi = \text{send}(m)$
 - No enabling issues (input).
 - Must check that $s' R u'$.
 - Since $s R u$, $u.queue$ is the concatenation of $s.A.queue$ and $s.B.queue$.
 - Adding the same m to the end of $u.queue$ and $s.B.queue$ maintains the correspondence.
- Case 2: $\pi = \text{receive}(m)$
 - Enabling: Check that $\text{receive}(m)$, for the same m , is also enabled in u .
 - We know that m is first on $s.A.queue$.
 - Since $s R u$, m is also first on $u.queue$.
 - So $\text{receive}(m)$ is enabled in u .
 - $s' R u'$: Since m is removed from both $s.A.queue$ and $u.queue$.

Showing R is a simulation relation

$s R u$ iff $u.queue$ is concatenation of $s.A.queue$ and $s.B.queue$

- Case 3: $\pi = \text{pass}(m)$
 - No enabling issues (since no high-level steps are involved).
 - Must check $s' R u$:
 - Since $s R u$, $u.queue$ is the concatenation of $s.A.queue$ and $s.B.queue$.
 - The concatenation of the queues is unchanged as a result of this step, so also $u.queue$ is the concatenation of $s'.A.queue$ and $s'.B.queue$.





Asynchronous network model

Send/Receive System


- Digraph $G = (V, E)$, with:
 - Process I/O automata associated with nodes, and
 - Channel I/O automata associated with directed edges.
- Compose all the automata to get a system automaton.

- Process:
 - User interface actions, e.g., invocations and responses
 - Send/receive actions for interaction with channels
- Problems specified in terms of allowable traces at the user interface.
 - Hide send/receive actions.
- Failure modeling:



- Having explicit *stop* actions in the external interface allows us to state requirements involving failures.

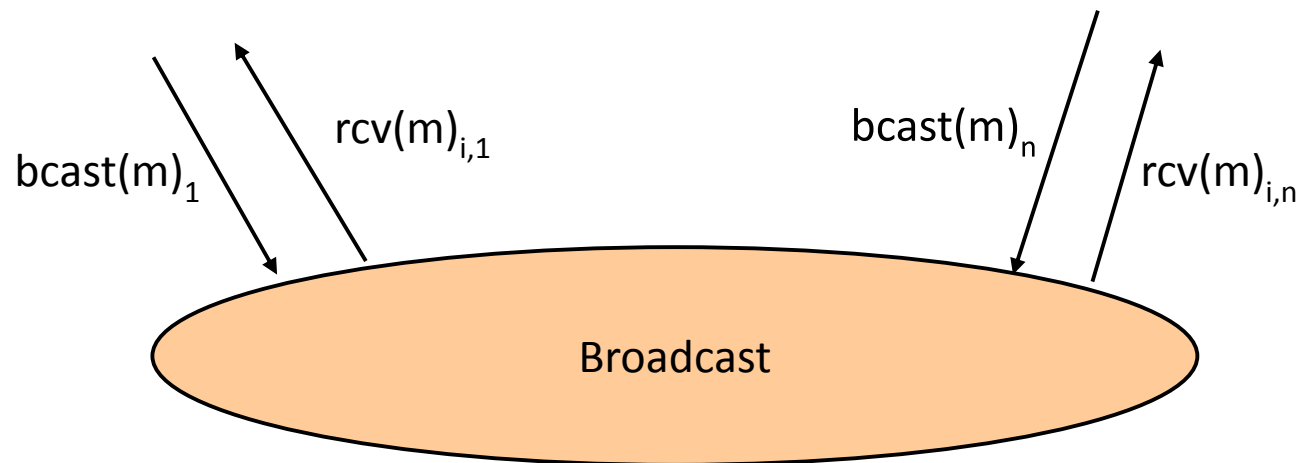
Channel automata




- Consider different kinds of channels with this interface:
 - Reliable FIFO, as before.
 - Weaker guarantees: Lossy, duplicating, reordering
- Can define channels by trace properties, using a *cause* function mapping receives to sends.
 - Integrity: The *cause* function preserves the message.
 - No loss: Function is onto (surjective).
 - No duplicates: Function is 1-1 (injective).
 - No reordering: Function is order-preserving.
- Reliable channel satisfies all of these conditions; weaker channels satisfy Integrity but may weaken some of the other properties.

Broadcast and multicast systems

- Broadcast
 - Reliable FIFO between each pair.
 - Different processes can receive messages from different senders in different orders.
 - Model abstractly using separate queues for each pair.
- Multicast: Processes designate recipients.




Asynchronous network algorithms



Asynchronous network algorithms

- Consider send/receive systems with reliable FIFO point-to-point channels
- Revisit problems we considered in synchronous networks:
 - Leader election, in a ring, and in general undirected networks.
 - Spanning tree
 - Breadth-first search
 - Shortest paths
 - Minimum spanning tree
- What results carry over?
- Where did we use the synchrony assumption?

Algorithms for Leader Election in a Ring



Leader election in a ring

- Assumptions:
 - G is a ring, with unidirectional or bidirectional communication
 - Local names for neighbors, UIDs
- *AsynchLCR* [LeLann] [Chang-Roberts]
 - Send UID clockwise around ring (unidirectional).
 - Discard UIDs smaller than your own.
 - Elect yourself if your UID comes back.
- Correctness: Basically the same as for synchronous version, with a few complications:
 - Finer granularity, must consider individual steps rather than entire rounds.
 - Must consider messages in channels.

AsynchLCR, process i

- Signature

- **in** $rcv(v)_{i-1,i}$, v a UID
- **out** $send(v)_{i,i+1}$, v a UID
- **out** $leader_i$

- State variables

- u , a UID, initially i 's UID
- $send$, a FIFO queue of UIDs, initially containing i 's UID
- $status$, unknown, chosen, or reported, initially unknown

- Tasks

- $\{ send(v)_{i,i+1} \mid v \text{ is a UID} \}$
- $\{ leader_i \}$

Transitions

- $send(v)_{i,i+1}$
pre: $v = \text{head}(send)$
eff: remove head of $send$
- $receive(v)_{i-1,i}$
eff:
if $v = u$ then $status := \text{chosen}$
if $v > u$ then add v to $send$
- $leader_i$
pre: $status = \text{chosen}$
eff: $status := \text{reported}$

AsynchLCR properties

- **Safety:** No process $i \neq i_{max}$ ever performs *leader_i*.
- **Liveness:** i_{max} eventually performs *leader_{i_{max}}*.

Safety proof

- **Safety:** No process $i \neq i_{max}$ ever performs *leader_i*.
- Recall the synchronous proof, based on showing an invariant of global states: After any number of **rounds**:
 - If $i \neq i_{max}$ and $j \in [i_{max}, i)$ then u_i not in *send_j*.
- We can use a similar invariant for the asynchronous algorithm:
 - If $i \neq i_{max}$ and $j \in [i_{max}, i)$ then u_i not in *send_j* or in *queue_{j,j+1}*.
- The main difference is that now the invariant must hold after any number of **steps**.
- Prove this by induction on number of steps.
 - Use cases based on the type of action.
 - Key case: *receive(v)*_{imax-1, imax}
 - Argue that if $v \neq u_{max}$ then v gets discarded.

Liveness proof

- **Liveness:** i_{max} eventually performs *leader* _{i_{max}} .
- Synchronous proof used an invariant saying exactly where the max is after r rounds.
- Now we don't have rounds, so we need a different proof.
- Can establish intermediate milestones, e.g.:
 - For $k \in [0, n - 1]$, u_{max} is eventually in *send* _{$i_{max}+k$}
 - Prove by induction on k ; use fairness for a process and a channel to prove the inductive step.

Complexity

- **Messages:** $O(n^2)$, as before.
- **Time:** $O(n(l + d))$
 - l is an upper bound on local step time for each process (that is, for each process task).
 - d is an upper bound on time to deliver the first message in each channel (that is, for each channel task).
 - Measuring real time here (not counting rounds).
 - Only upper bounds, so this does not restrict executions.
 - Bound still holds in spite of the possibility of “pileups” of messages in channels and send buffers.
 - Pileups can be interpreted as meaning that some tokens have sped up.
 - See analysis in book.

Reducing the message complexity

- Hirschberg-Sinclair:
 - Uses bidirectional communication.
 - Send in both directions, to successively doubled distances.
 - Extends immediately to the asynchronous model.
 - $O(n \log n)$ messages.
- Peterson:
 - Unidirectional communication
 - $O(n \log n)$ messages
 - Unknown ring size
 - Comparison-based

Peterson's algorithm

- Proceed in asynchronous “phases” (may execute concurrently).
- In each phase, each process is **active** or **passive**; passive processes just pass messages along.
- In each phase, at least half of the active processes become passive; so there are at most $\log n$ phases until election.
- **Phase 1:**
 - Send UID two processes clockwise; collect two UIDs from predecessors.
 - Remain active iff the middle UID is larger than the other two.
 - In this case, adopt the middle UID.
 - Some process remains active (assuming $n \geq 2$), but no more than half.
- **Later phases:**
 - Same, except that the passive processes just pass messages on.
 - No more than half of those active before the phase remain active.
- **Termination:**
 - If a process sees that its immediate predecessor's UID is the same as its own, it elects itself the leader (knows it's the only active process left).

PetersonLeader

- Signature

- **in** `receive(v)i-1,i`, v a UID
- **out** `send(v)i,i+1`, v a UID
- **out** `leaderi`
- **int** `get-second-uidi`
- **int** `get-third-uidi`
- **int** `advance-phasei`
- **int** `become-relayi`
- **int** `relayi`

- State variables

- **mode**: active or relay, initially active
- **status**: unknown, chosen, or reported, initially unknown
- **uid1**, initially i's UID
- **uid2**, initially null
- **uid3**, initially null
- **send**, FIFO queue of UIDs; initially contains just i's UID
- **receive**: FIFO queue of UIDs, initially empty

PetersonLeader

- get-second-uid_i
pre: $\text{mode} = \text{active}$
 receive is nonempty
 $\text{uid2} = \text{null}$
eff: $\text{uid2} := \text{head}(\text{receive})$
 remove head of receive
 add uid2 to send
 if $\text{uid2} = \text{uid1}$ then
 $\text{status} := \text{chosen}$
- get-third-uid_i
pre: $\text{mode} = \text{active}$
 receive is nonempty
 $\text{uid2} \neq \text{null}$
 $\text{uid3} = \text{null}$
eff: $\text{uid3} := \text{head}(\text{receive})$
 remove head of receive
- advance-phase_i
pre: $\text{mode} = \text{active}$
 $\text{uid3} \neq \text{null}$
 $\text{uid2} > \max(\text{uid1}, \text{uid3})$
eff: $\text{uid1} := \text{uid2}$
 $\text{uid2}, \text{uid3} := \text{null}$
 add uid1 to send
- become-relay_i
pre: $\text{mode} = \text{active}$
 $\text{uid3} \neq \text{null}$
 $\text{uid2} \leq \max(\text{uid1}, \text{uid3})$
eff: $\text{mode} := \text{relay}$
- relay_i
pre: $\text{mode} = \text{relay}$
 receive is nonempty
eff: move head(receive) to send


PetersonLeader

- Tasks:
 - $\{ \text{send}(v)_{i,i+1} \mid v \text{ is a UID} \}$
 - $\{ \text{get-second-uid}_i, \text{get-third-uid}_i, \text{advance-phase}_i, \text{become-relay}_i, \text{relay}_i \}$
 - $\{ \text{leader}_i \}$
- Number of phases is $O(\log n)$
- Complexity
 - Messages: $O(n \log n)$
 - Time: $O(n(l + d))$

Leader election in a ring


- **Q:** Can we do better than $O(n \log n)$ message complexity?
 - Not with comparison-based algorithms. (Why?)
 - **Not at all:**
 - Can prove another lower bound.
 - This one depends on asynchrony.

Lower Bound for Leader Election in a Ring





$\Omega(n \log n)$ lower bound

- Lower bound for leader election in asynchronous ring network.
- Assume:
 - Ring size n is unknown (algorithm must work in arbitrary size rings).
 - UIDS:
 - Chosen from some infinite set.
 - No restriction on allowable operations.
 - All processes are identical except for UIDs.
 - Bidirectional communication allowed.
- Consider combinations of processes to form:
 - Rings, as usual.
 - Lines, where nothing is connected to the ends and no input arrives at the ends.
 - Ring looks like a line if communication is delayed across a boundary.



$\Omega(n \log n)$ lower bound

- **Lemma 1:** There are infinitely many process automata, each of which can send at least one message without first receiving one (in some execution).
- **Proof:**
 - If not, then there are two processes i, j , neither of which sends a message without first receiving one.
 - Consider 1-node ring:
 - i must elect itself, with no messages sent or received.
 - Consider another 1-node ring:
 - j must elect itself, with no messages sent or received.
 - Now consider:
 - Both i and j elect themselves, contradiction.




$\Omega(n \log n)$ lower bound

- $C(L)$ = maximum (actually, supremum) of the number of messages that are sent in a single input-free execution of line L .
- **Lemma 2:** If L_1, L_2, L_3 are three line graphs of even length l such that each $C(L_i) \geq k$, then $C(L_i \text{ join } L_j) \geq 2k + l/2$ for some $i, j, i \neq j$.

- **Proof:**


- Suppose not.
- Construct an execution $\alpha_{1,2}$ of $L_1 \text{ join } L_2$:
- Let α_i be a finite execution of L_i with $\geq k$ messages, $i = 1, 2$.
- Run α_1 then α_2 then continue to a quiescent state (no more messages),
- May involve delivering delayed messages across the join boundary.
- By assumption, $< l/2$ additional messages are sent in $\alpha_{1,2}$.
- So, execution $\alpha_{1,2}$ quiesces before the effects of the new inputs can cross the middle edge of L_1 or the middle edge of L_2 .




$\Omega(n \log n)$ lower bound

- $C(L)$ = maximum (actually, supremum) of the number of messages that are sent in a single input-free execution of line L .
- **Lemma 2:** If L_1, L_2, L_3 are three line graphs of even length l such that each $C(L_i) \geq k$, then $C(L_i \text{ join } L_j) \geq 2k + l/2$ for some i, j .
- **Proof, cont'd:**

- Execution $\alpha_{1,2}$ quiesces before the effects of the new inputs can cross the middle edge of L_1 or L_2 .






- Similarly, construct $\alpha_{2,1}$, an execution of $L_2 \text{ join } L_1$.




- Execution $\alpha_{2,1}$ quiesces before the effects of the new inputs can cross the middle edge of L_1 or L_2 .

Proof of Lemma 2

- Now consider three rings:






Proof of Lemma 2




- Connect both ends of L_1 and L_2 .
 - Right neighbor in line is clockwise around ring.
- Run α_1 , then α_2 , then finish $\alpha_{1,2}$, then finish $\alpha_{2,1}$.
 - No interference between the last parts of $\alpha_{1,2}$ and $\alpha_{2,1}$.
 - Quiesces: Eventually no more messages are sent.
 - Must eventually elect a leader.
- Assume WLOG that elected leader is in the “bottom half”.



Proof of Lemma 2









- Same argument for ring constructed from L_2 and L_3 .
- Can the leader be in the bottom half?
- No!
- So it must be in top half.



Proof of Lemma 2



Proof of Lemma 2




Lower bound, cont'd

- Summarizing:
- **Lemma 1:** There are infinitely many process automata, each of which can send at least one message without first receiving one.
- **Lemma 2:** If L_1, L_2, L_3 are three line graphs of even length l such that each $C(L_i) \geq k$, then $C(L_i \text{ join } L_j) \geq 2k + l/2$ for some $i \neq j$.
- Combining, we get:
- **Lemma 3:** For any $r \geq 0$, there are infinitely many disjoint line graphs L of length 2^r such that $C(L) \geq r 2^{r-2}$.
- **Proof:** Induction on r .
 - Base ($r = 0$): Trivial claim.
 - Base ($r = 1$): Use Lemma 1
 - Just need length-2 lines sending at least one message.
 - Inductive step ($r \geq 2$):
 - Choose L_1, L_2, L_3 of length 2^{r-1} with $C(L_i) \geq (r-1) 2^{r-3}$.
 - By Lemma 2, for some i, j , $C(L_i \text{ join } L_j) \geq 2(r-1) 2^{r-3} + 2^{r-1}/2 = r 2^{r-2}$.

Lower bound, cont'd

- **Lemma 3:** For any $r \geq 0$, there are infinitely many disjoint line graphs L of length 2^r such that $C(L) \geq r 2^{r-2}$.
- **Theorem:** For any $r \geq 0$, there is a ring R of size $n = 2^r$ such that $C(R) = \Omega(n \log n)$.
 - Choose L of length 2^r such that $C(L) \geq r 2^{r-2}$.
 - Connect ends, but delay communication across boundary.
- Theorem can be extended to non-powers of 2. LTTR.

Leader Election in General Networks



Leader election in general networks

- Consider undirected graphs.
- We can get an asynchronous version of the synchronous *FloodMax* algorithm:
 - Simulate rounds with local counters.
 - Need to know the diameter for termination.
- We'll see several better asynchronous algorithms later:
 - Don't need to know diameter.
 - Lower message complexity.
- Depend on techniques such as:
 - Breadth-first search
 - Convergecast using a spanning tree
 - Synchronizers to simulate synchronous algorithms
 - Consistent global snapshots to detect termination

Spanning Trees and Searching


Spanning trees and searching

- Spanning trees are used for communication, e.g., bcast/ccast
- Start with the simple task of setting up some (arbitrary) spanning tree with a (given) root i_0 .
- **Assume:**
 - Undirected, connected graph (i.e., bidirectional communication).
 - Root i_0
 - Size and diameter unknown.
 - UUIDs, with comparisons for equality.
 - Can identify in- and out-edges to same neighbor.
- **Require:** Each process should output its parent in tree, with a *parent* output action.
- Starting point: *SynchBFS* algorithm:
 - i_0 floods a *search* message; parent of a node is the first node from which it receives a *search* message.
 - Q: What if we try to run the same algorithm in an asynchronous network?
 - Still yields a spanning tree, but not necessarily a breadth-first tree.


AsynchSpanningTree, Process i

- Signature
 - **in** $\text{receive}(\text{search})_{j,i}$, $j \in \text{nbrs}$
 - **out** $\text{send}(\text{search})_{i,j}$, $j \in \text{nbrs}$
 - **out** $\text{parent}(j)_i$, $j \in \text{nbrs}$
- State
 - **parent**: $\text{nbrs} \cup \{\perp\}$, init \perp
 - **reported**: Boolean, init false
 - for each $j \in \text{nbrs}$:
 - $\text{send}(j) \in \{\text{search}, \perp\}$,
init search if $i = i_0$, else \perp
- $\text{send}(\text{search})_{i,j}$
pre: $\text{send}(j) = \text{search}$
eff: $\text{send}(j) := \perp$
- $\text{receive}(\text{search})_{j,i}$
eff: if $i \neq i_0$ and $\text{parent} = \perp$ then
 $\text{parent} := j$
 for $k \in \text{nbrs} - \{j\}$ do
 $\text{send}(k) := \text{search}$
- $\text{parent}(j)_i$
pre: $\text{parent} = j$
 $\text{reported} = \text{false}$
eff: $\text{reported} := \text{true}$


AsynchSpanningTree




AsynchSpanningTree




AsynchSpanningTree




AsynchSpanningTree




AsynchSpanningTree




AsynchSpanningTree




AsynchSpanningTree



AsynchSpanningTree




AsynchSpanningTree



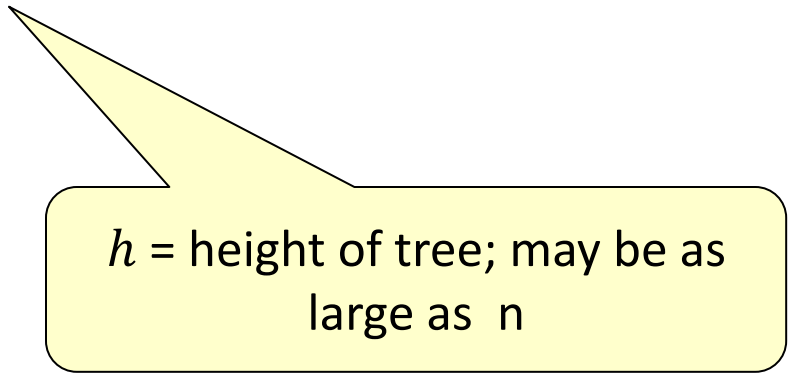
AsynchSpanningTree

- Complexity
 - Messages: $O(|E|)$
 - Time: $diam(l + d) + l$
- Anomaly: Paths may be longer than the diameter!
 - Messages may travel faster along longer paths, in asynchronous networks.



Applications of AsynchSpanningTree

- Similar to synchronous BFS
- Message broadcast: Piggyback on *search* message.
- Child pointers: Add responses to *search* messages, easy because of bidirectional communication.
- Use precomputed tree for bcast/convergecast
 - Now the timing anomaly becomes significant.
 - $O(h(l + d))$ time complexity.
 - $O(|E|)$ message complexity.
 - See book for details.



h = height of tree; may be as large as n

More applications

- Asynchronous broadcast/convergecast:
 - Can also construct spanning tree while using it to broadcast message and also to collect responses.
 - E.g., to tell the root when the bcast is done, or to collect aggregated data.
 - See book, p. 499-500, *AsynchBcastAck*.
 - Complexity:
 - $O(|E|)$ message complexity.
 - $O(n(l + d))$ time complexity, timing anomaly.
 - See book for details.
- Elect leader when nodes have no info about the network (no knowledge of n , $diam$, etc.; no root, no spanning tree):
 - All independently initiate *AsynchBcastAck*, use it to determine max, max elects itself.

Next lecture

- More asynchronous network algorithms
 - Breadth-first search
 - Shortest paths
 - Minimum spanning tree (GHS)
- Readings:
 - Sections 15.3-15.5
 - [Gallager, Humblet, Spira]