

6.852: Distributed Algorithms


Fall, 2015

Lecture 11

Today's plan

- Basic asynchronous network algorithms, general networks:
 - Leader election
 - (Arbitrary) spanning trees
 - Breadth-first spanning trees
 - Shortest-paths spanning trees
 - Minimum Spanning Trees (MSTs)
- Readings:
 - Chapter 15
 - [Gallager, Humblet, Spira]
- Next time:
 - Synchronizers
 - Reading: Chapter 16.

Leader Election in General Networks



Leader election in general networks

- Consider undirected graphs.
- We can get an asynchronous version of the synchronous *FloodMax* algorithm:
 - Simulate rounds with local counters.
 - Need to know the diameter for termination.
- We'll see several better asynchronous algorithms later:
 - Don't need to know diameter.
 - In some cases, better message complexity.
- Depend on techniques such as:
 - Breadth-first search
 - Convergecast using a spanning tree
 - Synchronizers to simulate synchronous algorithms
 - Consistent global snapshots

Spanning Trees and Searching


Spanning trees and searching

- Start with the simple task of setting up some (arbitrary) spanning tree with a (given) root i_0 .
- **Assume:**
 - Undirected, connected graph (i.e., bidirectional communication).
 - Root i_0
 - Size and diameter unknown.
 - UUIDs, with comparisons for equality.
 - Can recognize when in-edges and out-edges connect to the same neighbor.
- **Require:** Each process should output its parent in tree, with a *parent* output action.
- Starting point: *SynchBFS* algorithm:
 - i_0 floods a *search* message; parent of a node is the first neighbor from which it receives a *search* message.
- If we try to run the same algorithm in an asynchronous network, then we still get a spanning tree, but not necessarily a breadth-first tree.


AsynchSpanningTree, Process i

- Signature
 - **in** $\text{receive}(\text{search})_{j,i}$, $j \in \text{nbrs}$
 - **out** $\text{send}(\text{search})_{i,j}$, $j \in \text{nbrs}$
 - **out** $\text{parent}(j)_i$, $j \in \text{nbrs}$
- State
 - **parent**: $\text{nbrs} \cup \{\perp\}$, init \perp
 - **reported**: Boolean, init false
 - for each $j \in \text{nbrs}$:
 - $\text{send}(j) \in \{\text{search}, \perp\}$,
init search if $i = i_0$, else \perp
- $\text{send}(\text{search})_{i,j}$
pre: $\text{send}(j) = \text{search}$
eff: $\text{send}(j) := \perp$
- $\text{receive}(\text{search})_{j,i}$
eff: if $i \neq i_0$ and $\text{parent} = \perp$ then
 $\text{parent} := j$
 for $k \in \text{nbrs} - \{j\}$ do
 $\text{send}(k) := \text{search}$
- $\text{parent}(j)_i$
pre: $\text{parent} = j$
 $\text{reported} = \text{false}$
eff: $\text{reported} := \text{true}$


AsynchSpanningTree




AsynchSpanningTree




AsynchSpanningTree




AsynchSpanningTree




AsynchSpanningTree




AsynchSpanningTree




AsynchSpanningTree



AsynchSpanningTree




AsynchSpanningTree



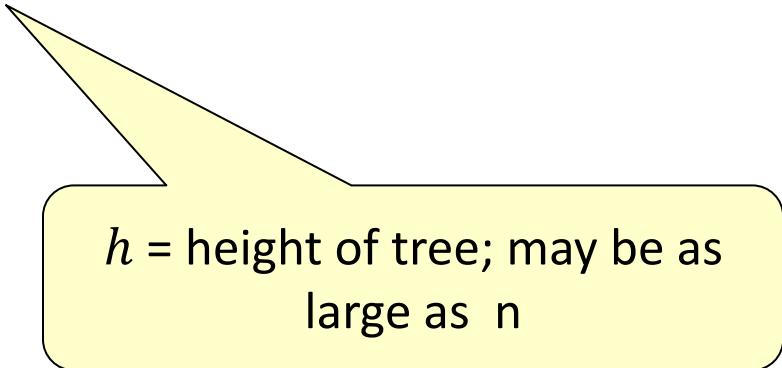
AsynchSpanningTree

- Complexity
 - Messages: $O(|E|)$
 - Time: $diam(l + d) + l$
- Anomaly: Paths may be longer than the diameter!
 - Messages may travel faster along longer paths, in asynchronous networks.



Applications of AsyncSpanningTree

- Similar to those for synchronous BFS
- Message broadcast: Piggyback on *search* message.
- Child pointers: Add responses to *search* messages, easy because of bidirectional communication.
- Use precomputed tree for broadcast/convergecast
 - Convergecast works as in the synchronous setting.
 - Now the timing anomaly becomes significant.
 - $O(h(l + d))$ time complexity.
 - $O(n)$ message complexity.
 - See book for details.



h = height of tree; may be as large as n

More applications

- Asynchronous broadcast/convergecast:
 - Can also construct spanning tree while using it to broadcast a message and also to collect responses.
 - E.g., to tell the root when the bcast is done, or to collect aggregated data.
 - See book, p. 499-500, *AsynchBcastAck*.
 - Complexity:
 - $O(|E|)$ message complexity.
 - $O(n(l + d))$ time complexity, timing anomaly.
 - See book for details.
- Elect leader when nodes have no info about the network (no knowledge of n , $diam$, etc.; no root, no spanning tree):
 - All independently initiate *AsynchBcastAck*, use it to determine max, max elects itself.

Breadth-First Spanning Trees

Breadth-first spanning trees

- Assume (same as above):
 - Undirected, connected graph (i.e., bidirectional communication).
 - Root i_0 .
 - Size and diameter unknown.
 - UIDs, with comparisons.
- Require: Each process should output its parent in a breadth-first spanning tree.
- *AsynchSpanningTree* does not guarantee that the spanning tree constructed is breadth-first.
 - Long paths may be traversed faster than short ones.
- Now modify each process to keep track of distance, change parent when it hears of a shorter path.
 - Relaxation algorithm (like Bellman-Ford).
 - Must inform neighbors of changes.
 - Eventually, tree stabilizes to a breadth-first spanning tree.

AsynchBFS

- Signature

- **in** $\text{receive}(m)_{j,i}$, $m \in \mathbf{N}$, $j \in \text{nbrs}$
- **out** $\text{send}(m)_{i,j}$, $m \in \mathbf{N}$, $j \in \text{nbrs}$

- State


- **dist**: $\mathbf{N} \cup \{\infty\}$, initially 0 if $i = i_0$, else ∞
- **parent**: $\text{nbrs} \cup \{\perp\}$, init \perp
- for each $j \in \text{nbrs}$:
 - **send(j)**: FIFO queue of \mathbf{N} , initially (0) if $i = i_0$, else empty

- $\text{send}(m)_{i,j}$
pre: $m = \text{head}(\text{send}(j))$
eff: remove head of $\text{send}(j)$


- $\text{receive}(m)_{j,i}$
eff: if $m+1 < \text{dist}$ then
 dist := $m+1$
 parent := j
 for $k \in \text{nbrs} - \{j\}$ do
 add **dist** to $\text{send}(k)$

Note: No parent output actions---no one knows when the algorithm is done


AsynchBFS




AsynchBFS




AsynchBFS




AsynchBFS




AsynchBFS




AsynchBFS




AsynchBFS




AsynchBFS




AsynchBFS




AsynchBFS




AsynchBFS




AsynchBFS




AsynchBFS



AsynchBFS



AsynchBFS



AsynchBFS

- Complexity:
 - **Messages:** $O(n |E|)$
 - May send $O(n)$ messages on each link (one for each distance estimate).
 - **Time:** $O(\text{diam } n (l + d))$ (taking pileups into account).
 - We can reduce complexity if we know an upper bound D on diameter:
 - Allow only distance estimates $\leq D$.
 - Messages: $O(D |E|)$; Time: $O(\text{diam } D (l + d))$
- Termination:
 - No one knows when this is done, so they can't produce *parent* outputs.
 - Can augment with *acks* for search messages, convergecast back to i_0 .
 - i_0 learns when the tree has stabilized, tells everyone else.
 - A bit tricky:
 - Tree grows and shrinks.
 - Some processes may participate many times, as they learn improvements.
 - Bookkeeping needed.
 - Complexity?

Layered BFS

- Asynchrony leads to many corrections, which lead to lots of communication.
- **Idea:** Slow down communication, grow the tree in synchronized phases.
 - In phase k , incorporate all nodes at distance k from i_0 .
 - i_0 synchronizes between incorporating nodes at distance k and $k + 1$.
- **Phase 1:**
 - i_0 sends *search* messages to neighbors.
 - Neighbors set *dist* := 1, send *acks* to i_0 .
- **Phase $k + 1$:**
 - Assume phases 1, ..., k are completed: each node at distance $\leq k$ knows its parent, and each node at distance $\leq k - 1$ also knows its children.
 - i_0 broadcasts *newphase* message along tree edges, to distance- k processes.
 - Each of these sends *search* message to all neighbors except its parent.
 - When any non- i_0 process receives its first *search* message, it sets *parent* := sender and sends *ack*; sends *nacks* for subsequent *search* messages.
 - When distance- k process receives *acks/nacks* for all its *search* messages, it designates nodes that sent *acks* as its children.
 - Distance- k processes convergecast back to i_0 along the depth k tree to say that they're done; include a bit saying whether any new nodes were found.

Layered BFS

- **Terminates:** When i_0 learns, in some phase, that no new nodes were found.
- Obviously produces BFS tree, in *diam* phases.
- **Complexity:**
 - **Messages:** $O(|E| + n \text{ diam})$

Each edge is explored at most once in each direction by search/ack.

Each tree edge is traversed at most once in each phase by newphase/convergecast.

– Time:

- Simplified analysis:
 - Neglect local computation time l
 - Assume every message in a channel is delivered in time d (ignore congestion delays).
- $O(\text{diam}^2 d)$

LayeredBFS vs AsynchBFS

- **Message complexity:**

- *AsynchBFS*: $O(\text{diam } |E|)$, assuming diameter is known, $O(n |E|)$ if not
- *LayeredBFS*: $O(|E| + n \text{diam})$

- **Time complexity:**

- *AsynchBFS*: $O(\text{diam } d)$
- *LayeredBFS*: $O(\text{diam}^2 d)$

- Can also define “hybrid” algorithm (in book)

- Add m layers in each phase instead of just one.
- Within each phase, layers get constructed asynchronously.
- Intermediate performance.

Shortest-Paths Spanning Trees

Shortest paths

- Assumptions:

- Same as for BFS, plus edge weights.
- *weight(i, j)*, nonnegative real, same in both directions.

- Require:

- Output shortest distance and parent in shortest-paths tree.

- Use Bellman-Ford asynchronously

- Used to establish routes in ARPANET 1969-1980.
- Can augment with convergecast as for BFS, for termination.
- But worst-case complexity is **very, very bad**...

AsynchBellmanFord

- Signature

- **in** $\text{receive}(w)_{j,i}$, $w \in \mathbf{R}^{\geq 0}$, $j \in \text{nbrs}$
- **out** $\text{send}(w)_{i,j}$, $w \in \mathbf{R}^{\geq 0}$, $j \in \text{nbrs}$

- State

- **dist**: $\mathbf{R}^{\geq 0} \cup \{ \infty \}$, initially 0 if $i = i_0$, else ∞
- **parent**: $\text{nbrs} \cup \{ \perp \}$, init \perp
- for each $j \in \text{nbrs}$:
 - **send(j)**: FIFO queue of $\mathbf{R}^{\geq 0}$; init (0) if $i = i_0$, else empty

- Transitions

- $\text{send}(w)_{i,j}$
pre: $w = \text{head}(\text{send}(j))$
eff: remove head of $\text{send}(j)$
- $\text{receive}(w)_{j,i}$
eff: if $w + \text{weight}(j,i) < \text{dist}$
then
 $\text{dist} := w + \text{weight}(j,i)$
 $\text{parent} := j$
 for $k \in \text{nbrs} - \{ j \}$ do
 add dist to $\text{send}(k)$



AsynchBellmanFord

- **Termination:**

- Use convergecast (as for AsynchBFS).



- **Complexity:**

- $O(n!)$ simple paths from i_0 to any other node, which is $O(n^n)$.
- So the number of messages sent on any channel is $O(n^n)$.
- So message complexity = $O(n^n |E|)$, time complexity = $O(n^n n (l + d))$.
- **Q:** Are the message and time complexity really exponential in n ?
- **A:** Yes: In some execution of the network below, i_k sends 2^k messages to i_{k+1} , so message complexity is $\Omega(2^{n/2})$ and time complexity is $\Omega(2^{n/2} d)$.



Exponential time/message complexity

- In some execution, i_k sends 2^k messages to i_{k+1} , so message complexity is $\Omega(2^{n/2})$ and time complexity is $\Omega(2^{n/2} d)$.
- Possible distance estimates for i_k are $2^k - 1, 2^k - 2, \dots, 0$.
- Moreover, i_k can take on all these estimates in sequence:
 - First, messages traverse upper links, $2^k - 1$.
 - Then last lower message arrives at i_k , $2^k - 2$.
 - Then lower message $i_{k-2} \rightarrow i_{k-1}$ arrives, reduces i_{k-1} 's estimate by 2, message $i_{k-1} \rightarrow i_k$ arrives on upper links, $2^k - 3$.
 - Etc. Count down in binary.
 - If this happens quickly, get pileup of 2^k search messages in $C_{k,k+1}$.



Shortest Paths

- Moral: Unrestrained asynchrony can cause problems.
- Return to this problem after we have better synchronization methods.
- Now, another good illustration of the problems introduced by asynchrony:

Minimum Spanning Tree

Minimum spanning tree

- Assumptions:

- $G = (V, E)$ connected, undirected.
- Weighted edges, weights known to endpoint processes, weights distinct.
- UIDs
- Processes don't know $n, diam$.
- Can identify in-edge and out-edge connecting to the same neighbor.
- Input: *wakeup* actions, occurring at any time at one or more nodes.
- Process wakes up when it first receives either a *wakeup* input or a protocol message.

- Requires:

- Produce MST, where each process knows which of its incident edges belong to the tree.
- Guaranteed to be unique, because of unique weights.

- [Gallager-Humblet-Spira]: Recommended reading!

Recall synchronous algorithm

- Proceeds in **phases (levels)**.
- After each phase, we have a **spanning forest**, in which each component tree has a leader.
- In each phase, each component finds **min weight outgoing edge (*MWOE*)**, then components merge using all *MWOE*s to get components for next phase.
- **In more detail:**
 - Each node is initially in component by itself (level 0 components).
 - **Phase 1 (produces level 1 components):**
 - Each node uses its min weight edge as the component *MWOE*.
 - Send ***connect*** message across *MWOE*.
 - There is a unique edge that is the *MWOE* of two components.
 - Leader of new component is higher-id endpoint of this unique edge.
 - **Phase $k + 1$ (produces level $k + 1$ components):**

Synchronous algorithm

- Phase $k + 1$ (produces level $k + 1$ components):
 - Leader of each component initiates search for $MWOE$ (broadcast *initiate* on tree edges).
 - Each node finds its *mwoe*:
 - Send *test* on potential edges, wait for *accept* (different component) or *reject* (same component).
 - Test edges one at a time in order of weight.
 - Report to leader (convergecast *report*); remember direction of best edge.
 - Leader picks $MWOE$ for component.
 - Send *changeroot* message to $MWOE$'s endpoint, using remembered best edges.
 - Send *connect* message across $MWOE$.
 - There is a unique edge that is the $MWOE$ of two components.
 - Leader of new component is higher-id endpoint of this unique edge.
 - Wait sufficient time for phase to end.

Synchronous algorithm

- Complexity is good:
 - Messages: $O(n \log n + |E|)$
 - Time (rounds): $O(n \log n)$
- Low message complexity depends on the way nodes test their incident edges, in order of weight, not retesting the same edge once it's rejected.
- Q: How to run this algorithm asynchronously?

Running the algorithm asynchronously

- Problems arise:


- Inaccurate information about outgoing edges:

- In the synchronous algorithm, when a node tests its edges, it knows that its neighbors are already up to the same level, and have up-to-date information about their component.
 - In asynchronous version, neighbors could lag behind; they might be in same component but not yet know this.

- Less “balanced” combination of components:

- In synchronous algorithm, level k components have $\geq 2^k$ nodes, and level $k + 1$ components are constructed from at least two level k components.
 - In asynchronous version, components at different levels could be combined.
 - Can lead to more messages overall.


- **Example:** One component might keep merging with level 0 single-node components. After each merge, the number of messages sent in the tree is proportional to the component’s size. Leads to $\Omega(n^2)$ messages overall.



Running the algorithm asynchronously

- Problems arise:

- Inaccurate information about outgoing edges.
- Less “balanced” combination of components:




- Concurrent overlapping searches/convergecasts:
 - When nodes are out of synch, concurrent searches for MWOEs could interfere with each other (we’ll see this).
- These problems result from nodes being out-of-synch, at different levels.
- We could try to synchronize levels, but carefully, so as not to hurt the time and message complexity too much.

GHS algorithm (asynchronous)

- Same basic ideas as before:
 - Form components, combine along *MWOEs*.
 - Within any component, processes cooperate to find component *MWOE*.
 - Broadcast from leader, convergecast, etc.
- Introduce **synchronization** to prevent nodes from getting too far ahead of their neighbors.
 - Associate a **level** with each component, as before.
 - Number of nodes in a level k component $\geq 2^k$, as before.
 - Now, each level $k + 1$ component will be (initially) formed from **exactly two** level k components.
 - Level numbers are used for synchronization, and for determining who is in the same component.
- Complexity:
 - **Messages:** $O(|E| + n \log n)$
 - **Time:** $O(n \log n (d + l))$


GHS algorithm

- Combine pairs of components in two ways, **merging** and **absorbing**.
- **Merging:**



- C and C' have same level k , and have a common $MWOE$.
- Result is a new merged component C'' , with level $k + 1$.

GHS algorithm



- **Absorbing:**
 - $level(C) < level(C')$, and C 's $MWOE$ leads to C' .
 - Result is to absorb C into C' .
 - Not creating a new component---just adding C to existing C' .
 - C “catches up” with the more advanced C' .
 - Absorbing is cheap, local.
- Merging and absorbing ensure that the number of nodes in any level k component $\geq 2^k$.
- Merging and absorbing are both allowable operations in computing the MST, because they are allowed by the general theory for MSTs.

Liveness

- **Q:** Why are merging and absorbing sufficient to ensure that the construction is eventually completed?
- **Lemma:** After any allowable finite sequence of merges and absorbs, either the forest consists of one tree (so we're done), or some merge or absorb is enabled.
- **Proof:**
 - Consider the current “component digraph”:
 - Nodes = components
 - Directed edges correspond to *MWOEs*
 - Then there must be some C, C' whose *MWOEs* point to each other. (Why?)
 - These *MWOEs* must be the same edge. (Why?)
 - Can combine, using either merge or absorb: If same level, merge, else absorb.
- So, merging and absorbing are enough.
- Now, how to implement them with a distributed algorithm?

Component names and leaders

- For every component with $level > 1$, define the **core edge** of the component's tree.
 - Defined in terms of the merge and absorb operations used to construct the component:
 - After merge: Use the common *MWOE*.
 - After absorb: Keep the old core edge of the higher-level component.
 - “The edge along which the most recent merge occurred.”
-
- **Component name:** $(core, level)$
 - **Leader:** Endpoint of core edge with higher id.

Determining whether an edge is outgoing


- Suppose i wants to know whether the edge (i, j) is outgoing from i 's current component.
- At that point, i 's component name info is up-to-date:
 - Component is in “search mode”.
 - i has received an *initiate* message from the leader, which included the component name.
- So i sends j a *test* message.
- **Three cases:**
 - If j 's current $(core, level)$ is the same as i 's, then j knows that it is in the same component as i .
 - If j 's $(core, level)$ is different from i 's and j 's level is $\geq i$'s, then j knows that j is in a different component from i .
 - Each component has only one core per level.
 - No one in the same component currently has a higher level than i does, since the component is still searching for its *MWOE*.
 - If j 's level is $< i$'s, then j doesn't know if it is in the same or a different component. So it doesn't yet respond---it waits to catch up to i 's level.

Liveness, again

- **Q:** Can the extra delays imposed here affect the progress argument?
- **No:**
 - We can redo the progress argument, this time considering only those components with the lowest current level k .
 - All processes in these components must succeed in determining their *mwoes*, so these components succeed in determining the component *MWOE*.
 - If any of these level k components' *MWOE*s leads to a higher level, then we can absorb.
 - If not then all lead to other level k components, so as before, we must have two components that point to each other; so we can merge.

Interference between *MWOE* searches


- Suppose C gets absorbed into C' via an edge from i to j , while C' is working on determining its *MWOE*.



- Two cases:
 - When the absorb occurs, j has not yet reported its local *mwoe*.
 - Then it's **not too late** for C' to include C in its *MWOE* search. So j passes the *initiate* message into C .
 - j has already reported its local *mwoe*.
 - Then it's **too late** to include C in the search.
 - But it doesn't matter: the *MWOE* for the combined component can't be outgoing from a node in C anyhow!
 - Why not?

Interference between *MWOE* searches

- If j has already reported its local *mwoe*, then the *MWOE* for the combined component is not outgoing from a node in C .



- Claim 1:** j 's reported *mwoe* is not the edge (i, j) .

- Proof:**

- j 's *mwoe* must lead to a node with $level \geq level(C')$.
- But i 's $level < level(C')$ when the *absorb* occurs.
- So j 's *mwoe* must be a different edge, with weight $< weight(i, j)$.

- Claim 2:** *MWOE* for combined component is not outgoing from a node in C .

- Proof:**

- The weight of the *MWOE* of the combined component is \leq the weight of j 's *mwoe*, so is $< weight(i, j)$.
- Since (i, j) is the *MWOE* of C , there are no edges outgoing from C with weight $< weight(i, j)$.
- So *MWOE* of combined component isn't outgoing from C .

A few details

- Specific messages:
 - *initiate*: Broadcast from leader to find *MWOE*; piggyback the component name on the message.
 - *report*: Convergecast the responses back to the leader.
 - *test*: Asks whether an edge is outgoing from the component.
 - *accept/reject*: Answers.
 - *changeroot*: Sent from leader to endpoint of *MWOE*.
 - *connect*: Sent across the *MWOE*, to connect components.
 - We say *merge* occurs when a *connect* message has been sent both ways on the edge (2 nodes must have same level).
 - We say *absorb* occurs when a *connect* message has been sent on the edge from a lower-level to a higher-level node.

Test-Accept-Reject Protocol

- Bookkeeping: Each process i keeps a list of incident edges in order of weight, classified as:
 - *branch* (in the MST),
 - *rejected* (leads to same component, not in the MST), or
 - *unknown* (not yet classified).
- Process i tests only *unknown* edges, in order of weight:
 - Sends *test* message, with $(core, level)$; recipient j compares.
 - If same $(core, level)$, j sends *reject* (same component), and i reclassifies edge as *rejected*.
 - If $(core, level)$ pairs are unequal and $level(j) \geq level(i)$ then j sends *accept* (different component). i does not reclassify the edge.
 - If $level(j) < level(i)$ then j delays responding, until $level(j) \geq level(i)$.
- Retesting is possible, for accepted edges.
- Reclassify edge as *branch* as a result of *changeroot* message.

Complexity

- As for synchronous version.
- **Messages:** $O(|E| + n \log n)$
 - $4|E|$ for *test* + *reject* messages (one pair for each direction of every edge)
 - n *initiate* messages per level (broadcast: only sent on tree edges)
 - n *report* messages per level (convergecast)
 - $2n$ *test* + *accept* message pairs per level (one pair per node)
 - n *changeroot* + *connect* messages per level (leader to *MWOE* path)
 - $\log n$ levels
 - Total: $4|E| + 5n \log n$
- **Time:** $O(n \log n (l + d))$

Proving Correctness

- **GHS** MST is hard to prove, because it's complicated.
- **GHS** paper includes informal arguments.
 - Pretty convincing, but not formal.
 - Also simulated the algorithm extensively.
- Many successful attempts to formalize, all complicated
 - Many invariants because many variables and actions.
 - Some use simulation relations.
 - Recent proof by **Moses and Shimony**.

Minimum spanning tree

- Application to leader election:
 - Convergecast from leaves until messages meet at node or edge.
 - Works with any spanning tree, not just MST.
 - E.g., in asynchronous ring, this yields $O(n \log n)$ messages for leader election.
- Lower bounds on message complexity for MST:
 - $\Omega(n \log n)$, from leader election lower bound and the reduction above.

Next time

- Synchronizers
- Reading: Chapter 16