

6.852: Distributed Algorithms

Fall, 2015

Lecture 12


Last Time

- Asynchronous spanning tree algorithms:
 - Asynch Spanning Tree (not necessarily breadth-first)
 - Asynchronous BFS
 - Asynchronous Shortest Paths
- Important observation: In a distributed algorithm, fast execution of portions of the algorithm don't necessarily result in fastest execution overall.
- Different from sequential algorithms.
- Also GHS Minimum Spanning Tree algorithm.
- Questions?

Today's plan

- Simulating synchronous algorithms in asynchronous networks.
- Synchronizers
- Lower bound for global synchronization
- **Reading:** Chapter 16
- **Next:**
 - Logical time, state machine emulation, vector timestamps
 - Readings:
 - Chapter 18
 - [Lamport] Time, Clocks, and the Ordering of Events in a Distributed System
 - [Mattern] Vector timestamps

Simulating Synchronous Algorithms in Asynchronous Networks (Synchronizers)



Minimum spanning tree, revisited


- In GHS, complications arise because different processes can be at very different levels at the same time.
- Alternative, simpler, more synchronized approach:
 - Keep levels of nearby nodes close, by restricting the asynchrony.
 - Each process uses a **level** variable to keep track of the level of its current component (according to its local knowledge).
 - Each process at level k delays all “interesting” processing until it hears **that all its neighbors have reached $level \geq k$** .
 - Looks (to each process) like global synchronization, but easier to achieve.
 - Each node inform its neighbors whenever it changes level.
- Resulting algorithm is simpler than GHS.
- **Complexity:**
 - Time: $O(n \log n)$, like GHS.
 - Messages: $O(|E| \log n)$, worse than GHS.

A strategy for designing asynchronous distributed algorithms

- Assume undirected graph $G = (V, E)$.
- Design a synchronous algorithm for G , then transform it into an asynchronous algorithm using local synchronization.
- Synchronize at every round (not every “level” as above).
- Method works only for non-fault-tolerant algorithms.
 - In fact, no general transformation can work for fault-tolerant algorithms.
 - E.g., simple fault-tolerant stopping agreement is solvable in synchronous networks, but unsolvable in asynchronous networks [FLP].
- Present a general strategy, and some special implementations.
 - Describe in terms of sub-algorithms, modeled as abstract services.
 - [Raynal book], [Awerbuch papers]
- Then a lower bound on the time for global synchronization.
 - Larger than upper bounds for local synchronization.


Synchronous model, reformulated in terms of I/O automata

- Global synchronizer automaton
- User process automata:
 - Processes of an algorithm that uses the synchronizer.
 - May have other inputs/outputs, for interacting with other programs.
- Interactions between user process i and synchronizer:
 - $usersend(T, r)_i$
 - T = set of (message, destination) pairs, destinations are neighbors of i .
 - T = empty set \emptyset , if no messages are sent by i at round r .
 - r = round number
 - $userrcv(T, r)_i$
 - T = set of (message, source) pairs, where source is a neighbor of i .
 - r = round number




Behavior of *GlobSynch*

- Manages global synchronization of rounds:
 - Users send packages of all their round 1 messages, using *usersend*($T, 1$) actions.
 - GlobSynch waits for all round 1 messages, sorts them, then delivers to users, using *userrcv*($T, 1$) actions.
 - Users send round 2 messages, etc.
- Not exactly the same as the synchronous model:
 - *GlobSynch* can receive round 2 messages from a user before it finishes delivering all the round 1 messages.
 - But it doesn't do anything with these until it's finished round 1 deliveries.
 - So, essentially the same.
- *GlobSynch* synchronizes globally between each pair of rounds.




Requirements on each U_i

- Well-formed:
 - U_i sends the right kinds of messages, in the right order, at the right times.
- Liveness:
 - After receiving the messages for any round r , U_i eventually submits messages for round $r + 1$.
- Code for *GlobSynch* in [book, p. 534].
 - State consists of:
 - A **tray** of messages for each (destination, round).
 - Some Boolean flags to keep track of which sends and rcvs have happened.
 - Transitions obvious.
 - Liveness expressed by tasks, one for each (destination, round).




The Synchronizer Problem

- Design an automaton A that “implements” *GlobSynch* in the sense that it “looks the same” to each U_i :
 - Has the right interface.
 - Exhibits the right behavior:
 - For every fair execution α of the U_i s and A ,
 - There exists a fair execution α' of the U_i s and *GlobSynch*, such that
 - For every i , α is indistinguishable by U_i from α' , written as $\alpha \sim_{U_i} \alpha'$.
- A “behaves like” *GlobSynch*, as far as any individual U_i can tell.
- Allows global reordering of events at different U_i .



Local Synchronizer, *LocSynch*

- Enforces local (not global) synchronization, still looks the same locally.
- Only difference from *GlobSynch*: the precondition for $usrrcv(T, r)_i$.
 - To deliver round r messages to user i , *LocSynch* checks only that i 's neighbors have sent round r messages.
 - Doesn't wait for all nodes.




- **Lemma 1:** For every fair execution α of the U_i s and *LocSynch*, there is a fair execution α' of the U_i s and *GlobSynch*, such that for each i , $\alpha \sim_{U_i} \alpha'$.
- **Proof:**
 - Can't use a simulation relation, since the global order of external events need not be the same, and simulation relations preserve external order.
 - So consider a partial order of events and dependencies:

Proof sketch for Lemma 1

- Consider partial order of events and dependencies:
 - Each U_i event depends on previous U_i events.
 - $userrcv(*, r)_i$ depends on $usersend(*, r)_j$ for every neighbor j of i .
 - Take transitive closure.
- **Claim:** If we start with a (fair) execution of the *LocSynch* system and reorder events while preserving these dependencies, the result is still a (fair) execution of the *LocSynch* system.
- So, obtain α' by reordering the events of α so that:
 - These dependencies are preserved, and
 - The rounds are globally aligned: events associated with any round r precede those of round $r+1$.
- OK because round r events don't depend on round $r + 1$ events.
- This reordering preserves the view of each U_i .
- Also satisfies the extra *userrcv* precondition needed by *GlobSynch*.

Trivial distributed algorithm to implement *LocSynch*

- Processes, point-to-point channels.
- SimpleSynch* algorithm, process i :
 - After *usersend*(T, r) _{i} , send a message to each neighbor j containing round number r and any algorithm messages i has for j .
 - Send (\emptyset, r) message if i has no basic algorithm messages for j .
 - Wait to receive round r messages from all neighbors.
 - Output *userrcv*(T', r).
- Lemma 2:** For every fair execution α of U_i s and *SimpleSynch*, there is a fair execution α' of U_i s and *LocSynch*, such that for each i , $\alpha \sim_{U_i} \alpha'$.
- In fact, indistinguishable by all the U_i s together---preserves external order.




SimpleSynch, cont'd

- **Proof of Lemma 2:**
 - No reordering needed, preserves order of external events.
 - Can use a simulation relation (for the safety part).
- **Corollary:** For every fair execution α of U_i s and *SimpleSynch*, there is a fair execution α' of U_i s and *GlobSynch*, such that for each i , $\alpha \sim_{U_i} \alpha'$.
- **Proof:** Combine Lemmas 1 and 2.
- **Complexity:**
 - Messages: $\leq 2 |E|$ per simulated round.
 - Time:
 - Assume user always sends ASAP.
 - l , upper bound on time for each task of each process.
 - d , upper bound on time for first message in channel to be delivered
 - Then r rounds completed within time $r (d + O(l))$.


Reducing the communication

- General **Safe Synchronizer** strategy [Awerbuch].
 - If there's no message from U_i to U_j at round r of the underlying synchronous algorithm, try to avoid sending such messages in the simulating asynchronous algorithm.
 - Can't just omit them, since each process must determine, for each round r , when it has received all of its round r messages.
 - **Key idea:** Separate the functions of:
 - Sending the actual messages, and
 - Determining when the round is over.
 - Algorithm decomposes into:
 - Front Ends + channels + *SafeSynch*
- For the actual messages For deciding when round is finished



Safe Synchronizers

- Front End:
 - Sends, receives algorithm messages for each round r .
 - Sends *acks* for received messages.
 - Waits to receive *acks* for its own messages.
- Notes:
 - Sends only actual algorithm messages, no dummies.
 - *acks* double the messages, but can still be a win.
- Front End, cont'd:
 - When FE receives *acks* for all its round r messages, it's **safe**: it knows that all its messages have been received by its neighbors.
 - Then sends **OK** for round r to *SafeSynch*.
 - Before responding to user, FE must know that it has received all its neighbors' messages for round r .
 - Suffices to know that **all its neighbors are safe**, that is, that they know that their messages have been received.
- *SafeSynch*:
 - Tells each FE when its neighbors are safe.
 - After it has received **OK** from i and all its neighbors, sends **GO** to i .




Correctness of SafeSynch


- **Lemma 3:** For every fair execution α of the *SafeSynch* system, there is a fair execution α' of the *LocSynch* system, such that for each i , $\alpha \sim U_i \alpha'$.
- (Actually, indistinguishable to all the U_i s together.)
- **Corollary:** For every fair execution α of the *SafeSynch* system, there is a fair execution α' of the *GlobSynch* system, such that for each i , $\alpha \sim U_i \alpha'$.
- We must still implement *SafeSynch* with a distributed algorithm...
- Three *SafeSynch* implementations: Synchronizers A, B, and Γ [Awerbuch].
- All implement *SafeSynch*, in the sense that the resulting systems are indistinguishable to each U_i (in fact, to all the U_i s together).

SafeSynch Implementations

- *SafeSynch's* job: After receiving *OK* for round r at location i and all its neighbors, send *GO* for round r at location i .
- **Synchronizer A:**
 - When process i receives OK_i , sends to neighbors.
 - When process i hears that it and all its neighbors have received OK s, outputs GO_i .
- Obviously implements *SafeSynch*.
- **Complexity:** To emulate r rounds:
 - Messages: $\leq 2m + 2r|E|$, if synchronous algorithm sends m messages in r rounds.



- Time: $\leq r(3d + O(l))$



Comparisons

- To emulate r rounds:
 - *SafeSynch* system with Synchronizer A
 - Messages: $2m + 2r |E|$
 - Time: $r(3d + O(l))$
 - *SimpleSynch*
 - Messages: $2r |E|$
 - Time: $r(d + O(l))$
- So Synchronizer A hasn't improved anything.
- Next, Synchronizer B, with lower message complexity, but higher time complexity.
- Then Synchronizer Γ , does well in terms of both messages and time, in an important subclass of networks (those with a “cluster” structure).

Synchronizer B

- Assumes rooted spanning tree of the graph, height h .
- Algorithm:**
 - All processes convergecast *OK* to the root, using spanning tree edges.
 - Root then broadcasts permission to *GO*, again using the spanning tree.
- Obviously implements *SafeSynch* (overkill).
- Complexity:** To emulate r rounds, in which synchronous algorithm sends m messages:

– Messages: $2m + 2rn$

Messages and acks by *FEs*

Messages within *B*

– Beats A: $2m + 2r|E|$


– Time: $\leq r(2d + O(l)) + 2h(d + O(l))$

FEs

B, convergecast and broadcast

Synchronizer Γ


- Hybrid of A and B.
- In “clustered” (almost partitionable) graphs, can get performance advantages of both:
 - Time like A, communication like B.
- Assume spanning forest of rooted trees, each tree spanning a “cluster” of nodes.
- Example:
 - Clusters = triangles
 - All edges between adjacent triangles in the line.
 - Spanning forest:



- Use B within each cluster, A between clusters.

Decomposition of Γ

- *ClusterSynch*:
 - After receiving *OK*s from everyone in the cluster, sends *clusterOK* to *ForestSynch*.
 - After receiving *clusterGO* from *ForestSynch*, sends *GO* to everyone in the cluster.
 - Similar to B.
- *ForestSynch*:
 - Essentially, a safe synchronizer for the “Cluster Graph” G' :
 - Nodes of G' are the clusters.
 - Edge between two clusters if and only if they contain nodes that are adjacent in G .
 - Send *clusterGO* to a cluster after receiving *clusterOK* from that cluster and all its neighboring clusters.
- **Lemma:** Automaton Γ Implements *SafeSynch*
- **Proof idea:**
 - Must show: If $GO(r)_i$ occurs, then there must be a previous $OK(r)_i$, and also a previous $OK(r)_j$ for every neighbor j of i .





Γ Implements SafeSynch

- **Show:** If $GO(r)_i$ occurs, then there must be a previous $OK(r)_i$, and also previous $OK(r)_j$ for every neighbor j of i .
- Must be a previous $OK(r)_i$:
 - $GO(r)_i$ preceded by $clusterGO(r)$ for i 's cluster (*ClusterSynch*),
 - Which is preceded by $clusterOK(r)$ for i 's cluster (*ForestSynch*),
 - Which is preceded by $OK(r)_i$ (*ClusterSynch*).
- Must be a previous $OK(r)_j$ for any neighbor j in the same cluster as i .
 - $GO(r)_i$ preceded by $clusterGO(r)$ for i 's cluster (*ClusterSynch*),
 - Which is preceded by $clusterOK(r)$ for i 's cluster (*ForestSynch*),
 - Which is preceded by $OK(r)_j$ (*ClusterSynch*).
- Must be a previous $OK(r)_j$ for any neighbor j in a different cluster.
 - Then the two clusters are **neighboring clusters in the cluster graph G'** , because i and j are neighbors in G .
 - $GO(r)_i$ preceded by $clusterGO(r)$ for i 's cluster (*ClusterSynch*),
 - Which is preceded by $clusterOK(r)$ for j 's cluster (*ForestSynch*),
 - Which is preceded by $OK(r)_j$ (*ClusterSynch*).

Implementing *ClusterSynch*, *ForestSynch*

- *ClusterSynch*:
 - Use variant of Synchronizer B on cluster tree:
 - Convergecast *OKs* to root on the cluster tree,
 - root outputs *clusterOK*, receives *clusterGO*,
 - root broadcasts *GO* on the cluster tree.
- *ForestSynch*:
 - Clusters run Synchronizer A.
 - But clusters can't actually run anything...
 - So *cluster roots* run A.
 - Simulate communication channels between neighboring clusters by indirect communication paths between the roots.
 - These paths must exist: Run through the trees and across edges that join the clusters.
- *clusterOK* and *clusterGO* are internal actions of the cluster root processes.



Putting the pieces together

- In Γ , process i emulates $FrontEnd_i$, process i in $ClusterSynch$ algorithm, and process i in $ForestSynch$ algorithm.
 - Formally, it's the composition of three I/O automata.
- Real channel $C_{i,j}$ emulates channel from $FrontEnd_i$ to $FrontEnd_j$, channel from i to j in the $ClusterSynch$ algorithm, and channel from i to j in the $ForestSynch$ algorithm.
- Orthogonal decompositions of Γ :
 - Physical: Nodes and channels.
 - Logical: FEs , $ClusterSynch$, and $ForestSynch$
 - Same system, two views.
 - Works because composition of I/O automata is associative, commutative.
- Such decompositions are common for complex distributed algorithms:
 - Each node runs pieces of algorithms at several layers.
- **Theorem 1:** For every fair execution α of the Γ (or A , or B) system, there is a fair execution α' of the $GlobSynch$ system, such that for each i , $\alpha \sim_{U_i} \alpha'$.

Complexity of Γ

- Consider r rounds, in which the synchronous algorithm sends m messages.
- Let:
 - h = max height of a cluster tree
 - e' = total number of edges on shortest paths between roots of neighboring clusters.

- Messages: $2m + O(r(n + e'))$

Messages and acks by FEs

Messages between roots,
In *ForestSynch* algorithm

Messages in cluster trees,
In *ClusterSynch* algorithm

- Time: $O(r h (d + l))$
- If $n + e' \ll |E|$, then Γ 's message complexity is much better than A 's.
- If $h \ll$ height of spanning tree of entire network, then Γ 's time complexity is much better than B 's.
- Both of these are true for “nicely clustered” networks.


Comparison of Costs

- r rounds
- m messages sent by synchronous algorithm
- d , message delay
- Ignore local processing time l .
- e' = total length of paths between roots of neighboring clusters
- h = height of global spanning tree
- h' = max height of cluster tree

	Messages	Time
A	$2m + 2r E $	$O(rd)$
B	$2m + 2rn$	$O(rhd)$
Γ	$2m + O(r(n + e'))$	$O(rh'd)$

Example

- $p \times p$ grid of complete k -node graphs, with all nodes of neighboring k -node graphs connected.
- Clusters = k -node graphs
- $h = O(p)$
- $h' = O(1)$
- $e' = O(p^2)$



	Messages	Time
A	$2 m + O(r p^2 k^2)$	$O(r d)$
B	$2 m + O(r p^2 k)$	$O(r p d)$
Γ	$2 m + O(r p^2 k)$	$O(r d)$

Synchronizer Applications

Application 1: Breadth-First Search

- **Recall:**
 - *SynchBFS*:
 - Constructs BFS tree.
 - $O(|E|)$ messages, $O(\text{diam})$ rounds
 - When run in asynchronous network:
 - Constructs a spanning tree, but not necessarily a BFS tree.
 - Modified version, with corrections:
 - Constructs BFS tree.
 - $O(n|E|)$ messages, $O(\text{diam} \cdot n \cdot d)$ time (counting pileups)
- **BFS using Synchronizers:**
 - Runs more like *SynchBFS*, avoids corrections, pileups
 - With Synchronizer A:
 - $O(\text{diam} |E|)$ messages, $O(\text{diam} \cdot d)$ time
 - With Synchronizer B :
 - Better communication, but worse time.
 - With Synchronizer Γ :
 - Better overall, in clustered graphs.

Application 2: Broadcast/Ack

- Assume known leader, but no spanning tree.
- Recall:
 - Synchronous Bcast/Ack:
 - Constructs spanning tree while broadcasting
 - $O(|E|)$ messages, $O(diam)$ rounds
 - Asynchronous Bcast/Ack:
 - Timing anomaly: Construct non-minimum-hop paths, on which acks travel.
 - $O(|E|)$ messages, $O(nd)$ time
- Bcast/Ack using Synchronizers:
 - Using (e.g.) Synchronizer A:
 - Avoids timing anomaly.
 - Bcast travels on min-hop paths, so Acks follow min-hop paths.
 - $O(diam |E|)$ messages, $O(diam d)$ time


Application 3: Shortest paths

- Assume weights on edges.
- Without termination detection.
- Recall:
 - Synchronous Bellman-Ford:
 - Makes some corrections, due to low-cost high-hop-count paths.
 - $O(n |E|)$ messages, $O(n)$ rounds
 - Asynchronous Bellman-Ford
 - Many more corrections possible (exponential), due to message delays.
 - (Worst-case) message complexity is exponential in n .
 - Time complexity also exponential in n , counting message pileups.
- Using (e.g.) Synchronizer A:
 - Behaves like Synchronous Bellman-Ford.
 - Avoids corrections due to message delays.
 - Still has corrections due to low-cost high-hop-count paths.
 - $O(n |E|)$ messages, $O(n d)$ time
 - Big improvement.

Further reading

- To read more:
 - See Awerbuch's extensive work on
 - Applications of synchronizers.
 - Distributed algorithms for clustered networks.
 - Also work by Peleg.
 - [Awerbuch, Peleg] Dijkstra Prize paper on clustered networks.
- **Q:** This work used a strategy of purposely slowing down portions of a system in order to improve overall performance. In which situations is this strategy a win?

Lower Bound on Time for Synchronization




Lower bound on time for synchronization

- A, B, Γ emulate synchronous algorithms only in a local sense:
 - Looks the same to individual users,
 - Not to the combination of all users---can reorder events at different users.
- Good enough for many applications (e.g., data management).
- Not for others (e.g., embedded systems).
- Now a theoretical result showing that **global synchronization is inherently more costly than local synchronization**, in terms of time complexity.
- **Approach:**
 - Define a toy global synchronization problem, the **k -Session Problem**.
 - Show that this problem has a fast synchronous algorithm, and thus, a fast algorithm using *GlobSynch*.
 - Time $O(k d)$, assuming *GlobSynch* takes steps ASAP.
 - Prove that all asynchronous distributed algorithms for this problem are slow.
 - Time $\Omega(k \text{ diam } d)$.
 - Implies *GlobSynch* has no fast distributed implementation.
- Contrast:
 - Synchronizer A, *SimpleSynch* are fast distributed impls of *LocSynch*.


k -Session Problem

- Session:
 - Any sequence of *flash* events containing at least one *flash_i* event for each location i .
- k -Session problem:
 - In every fair execution, perform at least k separate sessions, and eventually halt.
- Original motivation:
 - Synchronization of this kind is useful for performing parallel matrix computations that require enough interleaving of process steps, but tolerate extra steps.




Application: Boolean matrix computation

- $n = m^3$ processes cooperate to compute the transitive closure of an $m \times m$ Boolean matrix M .
- $p_{i,j,k}$ repeatedly does:
 - read $M(i, k)$, read $M(k, j)$
 - If both are 1 then write 1 in $M(i, j)$
- Each $\text{flash}_{i,j,k}$ in the abstract session problem represents a chance for $p_{i,j,k}$ to read or write a matrix entry.
- With enough interleaving ($O(\log n)$ sessions), this is guaranteed to compute the transitive closure.




Synchronous solution

- Fast algorithm using *GlobSynch*:
 - Just flash once at every round.
 - k sessions done in time $O(k d)$, assuming *GlobSynch* takes steps ASAP.



Asynchronous lower bound

- Consider a distributed algorithm A that solves the k -session problem.
- Consists of process automata and FIFO send/receive channel automata.



- Assume:
 - d = upper bound on time to deliver any message (don't count pileups)
 - l = local processing time, $l \ll d$
- Define time measure $T(A)$:
 - Timed execution α : Fair execution with times labeling events, subject to upper bound of d on message delay, l for local processing.
 - $T(\alpha)$ = time of last flash in α .
 - $T(A)$ = supremum, over all timed executions α , of $T(\alpha)$.

Lower bound

- **Theorem 2:** If A solves the k -session problem then
$$T(A) \geq (k - 1) \text{diam } d.$$
- Factor of diam worse than the synchronous algorithm.
- **Definition: Slow timed execution:** All message deliveries take exactly the upper bound time d .
- **Proof:** By contradiction.
 - Suppose $T(A) < (k - 1) \text{diam } d$.
 - Consider α , any slow timed execution of A .
 - α contains at least k sessions.
 - α contains no flash event at a time $\geq (k - 1) \text{diam } d$.
 - So we can decompose $\alpha = \underbrace{\alpha_1 \alpha_2 \dots \alpha_{k-1}}_{\alpha'} \alpha''$, where:
 - Time of last event in α' is $< (k - 1) \text{diam } d$.
 - No flash events occur in α'' .
 - The difference between the times of the first and last events in each α_r is strictly less than $\text{diam } d$.

Lower bound, cont'd

- Now reorder events in α , while preserving dependencies:
 - Events of same process.
 - Send and corresponding receive.
- Reordered execution will have strictly fewer than k sessions, which will yield a contradiction.
- Fix processes, j_0 and j_1 , with $\text{dist}(j_0, j_1) = \text{diam}$ (maximum distance apart).
- Reorder within each α_r separately:
 - For α_1 : Reorder to $\beta_1 = \gamma_1 \delta_1$, where:
 - γ_1 contains no event of j_0 , and
 - δ_1 contains no event of j_1 .
 - For α_2 : Reorder to $\beta_2 = \gamma_2 \delta_2$, where:
 - γ_2 contains no event of j_1 , and
 - δ_2 contains no event of j_0 .
 - And alternate thereafter.

Lower bound, cont'd

- If the reordering yields a fair execution of A (can ignore timing), then we get a contradiction, because it contains at most $k - 1$ sessions:
 - No session entirely within γ_1 , (no event of j_0).
 - No session entirely within $\delta_1 \gamma_2$ (no event of j_1).
 - No session entirely within $\delta_2 \gamma_3$ (no event of j_0).
 - ...
 - Thus, every session must span some $\gamma_r - \delta_r$ boundary.
 - But, there are only $k - 1$ such boundaries.
- It remains only to construct the reordering...

Constructing the reordering

- For example, consider α_r for r odd (analogous construction for r even).
- Need $\beta_r = \gamma_r \delta_r$, where γ_r contains no event of j_0 , δ_r no event of j_1 .
- If α_r contains no event of j_0 then don't reorder, define $\gamma_r = \alpha_r$, $\delta_r = \lambda$.
- If α_r contains no event of j_1 then don't reorder, define $\gamma_r = \lambda$, $\delta_r = \alpha_r$.
- Now assume α_r contains at least one event of each.
- Let π be the first event of j_0 , φ the last event of j_1 in α_r .
- **Claim:** φ does not depend on π .
- **Why:** There is insufficient time for messages to travel from j_0 to j_1 :
 - Execution α is slow (message deliveries take time d).
 - Time between π and φ is $< diam$.
 - j_0 and j_1 are $diam$ hops apart.
- Then, we can reorder α_r to β_r , in which π comes after φ .
- Consequently, in β_r , all events of j_1 precede all events of j_0 .
- Define γ_r to be the part ending with φ , δ_r the rest.

Next time...

- Time, clocks, and the ordering of events in a distributed system.
- State-machine emulation.
- Vector timestamps.
- Reading:
 - Chapter 18
 - [Lamport] Time, Clocks...paper
 - [Mattern] paper