

6.852: Distributed Algorithms

Fall, 2015

Lecture 14, Part 1

Weak Logical Time and Vector Timestamps

Weak Logical Time

- Logical time imposes a **total ordering** on events, assigning them values from a totally-ordered set T .
- Sometimes we don't need to order all events---it may be enough to **order just the ones that are causally dependent**.
- **Mattern** (also **Fidge**) developed an alternative notion of logical time based on a **partial ordering** of events, assigning them values from a partially-ordered set P .
- Function **ltime** from events in α to partially-ordered set P is a **weak logical time assignment** if:
 1. **ltime**s are distinct: $\text{ltime}(e_1) \neq \text{ltime}(e_2)$ if $e_1 \neq e_2$.
 2. **ltime**s of events at each process are monotonically increasing.
 3. $\text{ltime}(\text{send}) < \text{ltime}(\text{receive})$ for the same message.
 4. For any t , the number of events e with $\text{ltime}(e) < t$ is finite.
- Same as for logical time, but using partial order.

Weak Logical Time

- In fact, Mattern's partially-ordered set P represents causality exactly.
- Logical times of two events are ordered in P if and only if the two events are causally related (related by the causality ordering).
- Might be useful in distributed debugging: A log of local executions with weak logical times could be observed after the fact, used to infer causality relationships among events.

Algorithm for weak logical time

- Based on **vector timestamps**: vectors of nonnegative integers indexed by processes.
- **Algorithm:**
 - Each process maintains a local **vector clock**, called *vclock*.
 - When a non-receive event occurs at process i , it increments its own component of its *vclock*, which is $vclock(i)$, and assigns the new *vclock* to be the vector timestamp of the event.
 - Whenever process i **sends a message**, it attaches the vector timestamp of the send event.
 - When i **receives a message**, it first increases its *vclock* to the component-wise maximum of its current *vclock* and the incoming vector timestamp. Then it increments its $vclock(i)$ as before, and assigns the new *vclock* to the **receive** event.
- A process' *vclock* represents the latest known “tick values” for all processes.
- **Partially ordered set P :**
 - The vector timestamps, ordered based on \leq in all components.
 - $V \leq V'$ if and only if $V(i) \leq V'(i)$ for all i .

Key theorems about vector clocks

- **Theorem 1:** The vector clock assignment is a weak logical time assignment.
- **Lemma 1:** If event π causally precedes event π' , then the logical times are ordered, in the same order.
- **Proof:**
 - True for direct causality.
 - Use induction on the number of direct causality relationships.
- Claim this assignment **exactly captures causality:**
- **Lemma 2:** If the vector timestamp V of event π is (component-wise) \leq the vector timestamp V' of event $\pi' \neq \pi$, then π causally precedes π' .
- **Proof:** Prove the contrapositive: Assume π does not causally precede π' and show that V is not $\leq V'$.

Proof of Lemma 2

- **Lemma 2:** If the vector timestamp V of event π is (component-wise) \leq the vector timestamp V' of event $\pi' \neq \pi$, then π causally precedes π' .
- **Proof:**
 - Prove the contrapositive: Assume π does not causally precede π' and show that V is not $\leq V'$.
 - **Case 1:** π and π' are events of the same process i .
 - Then since π does not causally precede π' , it must be that π' precedes π in time.
 - Then $V'(i) < V(i)$.
 - So V is not $\leq V'$.
 - **Case 2:** π is an event of process i and π' an event of another process $j \neq i$.

Proof of Lemma 2


- **Lemma 2:** If the vector timestamp V of event π is (component-wise) \leq the vector timestamp V' of event $\pi' \neq \pi$, then π causally precedes π' .
- **Proof:**
 - Prove the contrapositive: Assume π does not causally precede π' and show that V is not $\leq V'$.
 - **Case 2:** π is an event of process i and π' an event of process $j \neq i$.
 - i increases its $vclock(i)$ for event π , say to value t .
 - Without causality, there is no way for this value t for i to propagate to j before π' occurs.
 - So, when π' occurs at process j , j 's $vclock(i) < t$.
 - So V is not $\leq V'$.

Back to Theorem 1

- **Theorem 1:** The vector clock assignment is a weak logical time assignment.
- **Lemma 1:** If event π causally precedes event π' , then the logical times are ordered, in the same order.
- **Lemma 2:** If the vector timestamp V of event π is (component-wise) \leq the vector timestamp V' of event $\pi' \neq \pi$, then π causally precedes π' .
- **Proof of Theorem 1:**
 - The ordering is a partial order.
 - Lemma 1 yields Properties 2 and 3.
 - Lemma 2 yields Property 1 (uniqueness).
 - Property 4 (non-Zeno) LTTR.


Another important theorem about vector timestamps [Mattern]

- Relates timestamps to **consistent cuts** of the causality graph.
- **Cut:** A point between events at each process.
 - Specify a cut by a vector giving the number of preceding steps at each location.
- **Consistent cut:** “Closed under causality”: If event π causally precedes event π' and π' is before the cut, then so is π .
- **Example:**



The theorem

- Consider any particular cut.
- Let V_i be the vector clock of process i exactly at i 's cut-point.
- Then $V = \max(V_1, V_2, \dots, V_n)$ gives the maximum information obtainable by combining everyone knowledge at the cut-points.
 - Component-wise max.
- **Theorem 2:** The cut is consistent iff, for every i , $V(i) = V_i(i)$.
- That is, the maximum information about i that anyone knows at its cut point is the same as what i knows about itself at its cut point.
- “No one else knows more about i than i itself knows.”
- Rules out j receiving a message before its cut point that i sent after its cut point; in that case, j would have more information about i than i had about itself.



The theorem

- Let V_i be the vector clock of process i exactly at i 's cut-point.
- $V = \max(V_1, V_2, \dots, V_n)$.
- **Theorem 2:** The cut is consistent iff, for every i , $V(i) = V_i(i)$.
- Stated slightly differently:
- **Theorem 2:** The cut is consistent iff, for every i and j , $V_j(i) \leq V_i(i)$.
- **Proof:** LTTR (see Mattern's paper).

- **Q:** What is this good for?

Application: Debugging

- **Theorem 2:** The cut is consistent iff, for every i and j , $V_j(i) \leq V_i(i)$.
- **Example: Debugging**
 - Each node keeps a log of its local execution, with vector timestamps for all events.
 - Collect information, find a cut for which $V_j(i) \leq V_i(i)$ for every i and j . (**Mattern** gives an algorithm to do this.)
 - By Theorem 2, this is a consistent cut.
 - Such a cut yields:
 - States for all processes at the cut, and
 - Information about messages sent before the cut and not received until after the cut, i.e., messages “in transit” at the cut.
 - Put this together, get a “consistent” global state (we will study this next).
 - Use this to check correctness properties for the execution, e.g., invariants.